

ADVAIT M. NARVEKAR

narvekaradvait03@gmail.com • (315) 832 0340 • [Linkedin](#) • New York, USA

SUMMARY

Data Engineer with **5 years of experience** building and operating production pipelines across **Azure Databricks, Snowflake, and Azure Data Factory** in healthcare and financial services. Strong in **SQL and Python** for transformation, validation, and performance tuning. Known for stabilizing unreliable pipelines, improving data quality, and delivering clean datasets for analytics and regulated reporting.

EDUCATION

Master of Science (MS), Data Science | Syracuse University, Syracuse, New York.

Aug 2023 - May 2025

Bachelor of Engineering (BE), Computer Engineering | University of Mumbai, India.

Aug 2016 - May 2020

TECHNICAL SKILLS

- **Core:** SQL, Python, PySpark, Spark SQL, R, MongoDB
- **Cloud & Data Platforms:** Azure Databricks, Snowflake, Azure Synapse Analytics, ADLS Gen2, Azure SQL, Azure Storage Explorer
- **Processing & Frameworks:** Apache Spark (PySpark, Spark SQL), Spark Structured Streaming, Delta Lake
- **Analytics Engineering:** SQL-based ELT, DBT based transformations (models, tests, incremental builds), analytics data marts
- **Data Modeling & Architecture:** Medallion Architecture (Bronze / Silver / Gold), Data Vault 2.0, dimensional modeling
- **Data Quality & Governance:** Data validation rules, schema enforcement, RBAC, secrets management (Azure Key Vault), CDC
- **Monitoring & Reliability:** Azure Monitor, Log Analytics, Git, Azure DevOps, CI/CD pipelines
- **Visualization & Analysis:** Power BI, Streamlit, Excel (Power Query, PivotTables)
- **Orchestration & Scheduling:** Apache Airflow (DAG-driven dependency management alongside Azure Data Factory)

PROFESSIONAL EXPERIENCE

Data Analytics Engineer, Syracuse University, United States.

Aug 2023 - Present

- Built and maintained ingestion and transformation pipelines using **PySpark and SQL across Azure Databricks, Snowflake, and Azure Data Factory**, processing more than 15 million healthcare claims and EHR records per month.
- Developed analytics-ready tables and marts in **Snowflake using DBT-based incremental models**, enabling consistent metric definitions for reporting and self-service analysis.
- Implemented validation checks using **DBT tests and SQL** reconciliation queries (row counts, null thresholds, referential integrity), improving reporting accuracy from 83% to 95%.
- Tuned **Spark and Snowflake workloads** through **SQL refactoring, join strategy improvements, and caching**, reducing end-to-end runtime by 65% while maintaining stable compute usage.
- Orchestrated end-to-end pipeline execution using **Apache Airflow** to coordinate ingestion, transformation, and downstream **Snowflake refreshes alongside Azure Data Factory**.
- Productionized pipelines using **Git, Azure DevOps CI/CD, and Airflow scheduling**, enabling controlled releases, repeatable deployments, and safe historical backfills across environments.

Data Integration Engineer, Loopback Analytics – Contract.

Oct 2025 - Dec 2025

- Joined as a contract backfill to stabilize **Databricks and Snowflake** ingestion workflows after mid-cycle attrition, supporting a time-sensitive healthcare production release.
- Supported ingestion from **more than 10 upstream healthcare systems**, reducing release-blocking defects and downstream data quality incidents during production runs.
- Built an **XML-based schema mapping** and validation layer to catch data-type and structural issues before execution, reducing ETL reruns and saving 8 to 12 engineering hours per release cycle.
- Converted ad-hoc ingestion into **repeatable Databricks Jobs and DBT-driven Snowflake ELT steps**, with documentation and operational handoff to internal teams.
- Diagnosed failures using **SQL checks and Spark job logs**, then added validation gates and retry logic to reduce on-call escalations during release windows.
- Ensured secure handling of **HIPAA-sensitive** data using restricted access policies, environment isolation, and secrets management.

Data Engineer, NatWest (Accenture), India.

May 2020 - May 2023

- Built ELT pipelines using **PySpark, SQL, and DBT-style** transformations across **Azure Databricks and Snowflake**, processing more than 5 terabytes of trading and regulatory data per cycle.
- Reduced **Basel III reporting turnaround** from over 12 hours to under 5 hours by optimizing Spark transformations, Snowflake query patterns, and pipeline execution flow.
- Implemented near-real-time ingestion using **Spark Structured Streaming** for market and risk feeds, enabling analytics with data latency under 10 minutes.
- Modeled more than 200 financial-domain tables using **Data Vault 2.0 and dimensional schemas**, improving traceability and audit readiness for regulated reporting.
- Created feature engineering datasets for risk and analytics teams, ensuring reproducible training snapshots across historical re-runs and regulatory scenarios.
- Set up pipeline monitoring and alerting using **Azure Monitor and Log Analytics**, reducing recurring SLA breaches and manual recovery.
- Worked with compliance and security teams to enforce **RBAC**, encryption, and secrets management using **Azure Key Vault**, supporting **GDPR and SOX** controls.