

Yashi Agarwal

CA, USA | +1 (925) 524-8625 | yashiagarwal762@gmail.com | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Versatile Data Science and AI professional with 3+ years of hands-on experience solving analytics and machine learning challenges across business and product domains. Adept at building ML models, scalable ETL pipelines, and interactive dashboards. Skilled in statistical analysis, cloud platforms, data storytelling, and responsible AI, blending technical innovation with real-world business support.

SKILLS

Python | SQL | R | Advanced Excel | Power Query | Git | Docker | Tableau | Power BI | MLflow | TensorFlow | PyTorch | LangChain | Hugging Face | Pandas | NumPy | Scikit-learn | Deep Learning | NLP | Prompt Engineering | Feature Engineering | Data Storytelling | Predictive Modeling | A/B Testing | Statistics | Responsible AI (Fairlearn, AIF360) | Communication | Apache Airflow | FastAPI | AutoML | Google AutoML | H2O.ai | DataRobot | Data Cleaning | BigQuery | Snowflake | Redshift | Spark | Hadoop | PostgreSQL | MySQL | MongoDB | Vector Databases | REST API | Version Control | Data Governance | D3.js | Plotly

EXPERIENCE

AI Engineer **Feb 2025 – Present**
Kaiyros Inc **Los Angeles, CA**

- Built NLP and text classification pipelines using Python, spaCy, LangChain, and Hugging Face Transformers for automated content tagging of social and product data.
- Led integration and deployment of ML models with FastAPI and Docker for scalable APIs, managed CI/CD using GitHub Actions for version control.
- Created interactive Power BI dashboards and custom D3.js analytics tools for campaign reporting and business insights.
- Automated ETL workflows with Python, Power Query, and Airflow; utilized AWS Lambda for real-time updates and storage in BigQuery.
- Collaborated cross-functionally to ensure responsible AI and fairness in recommendations using Fairlearn, documented processes for compliance.

Data Scientist **Jun 2024 – Aug 2024**
Roundglass **Bellevue, WA**

- Designed recommendation systems using Sentence Transformers, pgvector with PostgreSQL, and optimized query performance for large datasets.
- Built and validated ML models using Scikit-learn, AutoML (H2O.ai), and MLflow for experiment and metric tracking.
- Supported RAG chatbot deployment with LangChain, prompt engineering, and OpenAI APIs; delivered contextual analytics as REST endpoints.
- Created Tableau dashboards, conducted data storytelling sessions, and documented workflow using Jupyter Notebooks and Git.
- Integrated data from Snowflake, performed quality checks, and maintained robust data governance using standardized controls.

Data Scientist **Jan 2021 – Dec 2022**
ThoughtWorks **Gurugram, India**

- Developed supervised and unsupervised models using Pandas, Scikit-learn, and Spark for large-scale customer analytics.
- Built automated ETL pipelines employing SQL, Power Query, and Hadoop to prepare structured model training datasets.
- Containerized and deployed model APIs using Docker and FastAPI, tracked performance and drift using MLflow.
- Delivered interactive and actionable reports through Tableau, Power BI, and custom Plotly dashboards for executive stakeholders.
- Drove critical thinking and collaborative business analysis, translating model results into clear recommendations for product and strategy teams.

Data Scientist **Aug 2020 – Dec 2020**
DataOne Innovation Labs **Remote, India**

- Fine-tuned domain-specific LLMs and transformer models using PyTorch Lightning, managed data experiments via MLflow.
- Built feature engineering and hyperparameter tuning frameworks leveraging Google AutoML and integrated workflow orchestration with Airflow.
- Developed ETL and data validation pipelines in Python, Excel/Power Query, and maintained structured logs in MongoDB and PostgreSQL.
- Created interactive dashboards and conducted data storytelling sessions to translate technical findings into actionable business intelligence.
- Supported responsible AI initiatives, implemented model fairness checks using Fairlearn and documented compliance strategies.

EDUCATION

Master's in Statistics – Data Science, California State University, East Bay, CA | **Jan 2023 – Dec 2024**
B.Tech in Electronics & Communication Engineering, Guru Gobind Singh Indraprastha University, Delhi, India | **Aug 2018 – May 2022**

PROJECTS

Customer Churn Prediction Platform
Developed a production-ready ML pipeline with Random Forest and Scikit-learn, tracked via MLflow and visualized in Tableau.
Tech Stack: Python, SQL, Pandas, Scikit-learn, MLflow, Tableau

Sales Forecasting Dashboard
Created time-series forecasts with LSTM models, automated via Airflow and GCP, and reporting powered by Power BI and Plotly.
Tech Stack: Python, TensorFlow, Airflow, GCP, Power BI, Plotly, SQL

Sentiment Analysis Engine
Built and deployed a FastAPI service using BERT, Hugging Face, and Docker for review analytics and prompt engineering tasks.
Tech Stack: Python, BERT, Hugging Face, FastAPI, Docker, PostgreSQL, Git