

SAI AKHIL SAKHAMURI

✉ akhilsakhamurijob@gmail.com | 🌐 akhil-sakhamuri | ☎ +1 (945)-233-9203 | Dallas, Tx.

SUMMARY

Senior Software Engineer (5+ years) specializing in distributed systems, cloud-native services, and mission-critical APIs across AWS and Azure. Proven track record designing, building, and operating high-availability, low-latency services with Kubernetes, Docker, and Terraform improving p95 latency by 20–40% and cutting release lead time by ~30%. Experienced owning features end-to-end: requirements, architecture, implementation, observability, and on-call support. Strong in Python and C#/.NET with deep exposure to data-intensive systems (PostgreSQL, Kafka, Spark) and modern SDLC practices (CI/CD, SLOs, runbooks). Known for simplifying complex systems, driving cross-team alignment, and mentoring engineers in clean architecture and operational excellence.

EDUCATION

University of Texas, Dallas

Master of Science in Information Technology & management

Dallas, Texas

Aug 2022 - May 2024

Dean's Excellence Scholar, Coursework: Software Development & Networking, Data Visualization, AI & ML for everyone, Advance

Statistics, Business Analytics with Python and R.

Karunya University

Coimbatore, India

Bachelor of Technology, Mechanical Engineering

June 2016 - June 2020

Coursework: Python for beginners, C, Database foundations, Supply chain strategy, Industrial Engineering, Computer Engineering.

EXPERIENCE

Adora AI

July 2025 - Present

Senior Software Engineer, Cloud & AI Platform

Dallas, Texas

- Designed and scaled multi-tenant, containerized services on Azure Kubernetes Service (AKS) and Docker, improving measured uptime and horizontal scalability by ~30% while keeping infrastructure costs predictable.
- Built Azure ML/Databricks release pipelines with automated tests, canary deployments, and promotion gates, reducing end-to-end model/service release lead time by ~30%.
- Hardened production readiness with Python integration tests, infra config checks, and health probes; reduced post-deployment defects by ~40% and improved on-call stability.
- Optimized API-first request paths (routing + Postgres/Supabase I/O) for retrieval-heavy workloads; improved p95 latency by ~20% for key customer journeys.
- Automated multi-cloud infrastructure delivery using Terraform and container pipelines, standardizing environments across AWS and Azure and significantly reducing configuration drift.
- Integrated with .NET microservices via OpenAPI contracts; contributed to API design reviews and auth/observability standards

Technologies: Python, FastAPI, AWS (SageMaker, Lambda, API GW, S3, IAM, CloudWatch), LangChain/LangGraph, MLflow, Terraform.

Verizon:

December 2023 - July 2025

Senior Software Engineer

Dallas, Texas

- Partnered with CX and operations stakeholders to map real-world associate workflows; delivered internal analytics and insights services that automated manual steps and cut time-to-answer by ~35% for support agents.
- Led an experiment-driven improvement cycle using evaluation harnesses (quality + latency), dashboards, and golden datasets boosting exact-match accuracy from 61% → 74% on targeted customer journeys.
- Shipped FastAPI-based microservices on AWS (Lambda, ECR, SageMaker) with clearly defined SLOs, structured logging, and error budgets, supporting mission-critical, tier-one internal experiences.
- Instituted CI/CD with GitHub Actions/Bitbucket (versioning, tests, image build & deploy) to raise deployment frequency while meeting change-failure targets.
- Automated MLOps with GitHub Actions/Bitbucket (versioning, tests, image build, deploy); improved deployment frequency while maintaining change failure rate targets.
- Implemented evaluation & tracing (faithfulness, EM, latency) and golden-set regression checks to block quality regressions prior to release.

Technologies: Python, PyTorch, FastAPI, Pinecone, Postgres, OpenAI embeddings, LangChain/LangGraph, GitHub Actions, MLflow.

Genpact

July 2019 - July 2022

Python Developer

Hyderabad, India

- Built RESTful backends and workflow automation services for operations teams, introducing modular service design, code reviews, and CI tests that strengthened SDLC quality and predictability.
- Containerized 12+ services and rolled out CI/CD pipelines; reduced lead time for changes from weeks to a few days, doubled deployment frequency, and lowered change-fail rate by ~6 percentage points via automated smoke tests.
- Tuned SQL/Postgres schemas & queries to improve reliability and sprint predictability (~15%) across releases.
- Introduced prompt templates, retrieval filters, rerankers, and eval harnesses that produced double-digit accuracy gains and 35% faster TTA across knowledge flows.

Technologies: Python, Flask, HTML, CSS, Bootstrap, MySQL, PostgreSQL. Django, REST APIs, MySQL, Agile scrum.

Languages & Frameworks: Python (Pandas, NumPy, scikit-learn, TensorFlow), FastAPI, Flask, Django, .NET, Salesforce.

AI/ML Frameworks: PyTorch, TensorFlow, Hugging Face, LangChain/LangGraph, CrewAI, MLflow, LlamaIndex.

Data & RAG Infra: Pinecone/FAISS, PostgreSQL/Supabase; robust chunking, metadata, retrieval filters

MLOps: AWS (Lambda, ECR, SageMaker, S3, EC2, CloudWatch), Azure (Functions, AKS, ML, Blob), Docker, Kubernetes, GitHub Actions, Bitbucket, terraform, Jenkins, Canary deployments.

APIs: FastAPI, REST/OpenAPI; secure auth, audit logs, RBAC

Data/ETL: SQL/PostgreSQL, Databricks, Spark, Supabase, Spark, Airflow, Kafka, Docker, Terraform, Power BI, Tableau.