ARCHITA SEVAK| Data Analyst

+1 (864) 283-4482| architasevak97@gmail.com | [LinkedIn]

SUMMARY

Data Engineer with 6+ years of experience designing and optimizing real-time and batch data pipelines using Apache Spark, PySpark, Kafka, and Airflow. Expert in AWS and Azure cloud platforms, Databricks, Snowflake, and integrating ML models via SageMaker and Scikit-learn. Successfully increased data throughput by 40%, reduced latency from 15 minutes to under 2 minutes, and improved data accuracy to 99.5%. Skilled in dashboarding (Tableau, Power BI) and delivering actionable insights in healthcare and financial analytics.

TECHNICAL SKILLS

Programming: Python, Scala, C++, Golang, TypeScript, JavaScript, SQL, R, HTML, MATLAB
Frameworks & Libraries: Flask, Go-Kit, Apache Flink, Apache Spark, Scikit-learn, Gradient Boosting (GBM), PySpark, SQLAlchemy
Cloud & DevOps: AWS (Lambda, S3, Athena, Kinesis, DynamoDB, ECS, Glue, Data Pipeline, CloudWatch), Azure (Data Factory, Synapse, Functions, App Services, Event Hubs, Azure DevOps), GitHub Actions, SageMaker, CI/CD, Terraform, Argo, Jenkins
Big Data & ETL: Kafka, Airflow, Snowflake, MongoDB, Redis, Databricks, Cassandra, Hadoop
Testing & Monitoring: JMeter, Unit Testing, Data Validation, Alerting Systems, Datadog, Prometheus
Databases: SQL Server, Postgres, Oracle, OLAP/OLTP Data Modeling, Vector Databases
Data Formats: JSON, XML
Tools & Platforms: GenAI, ML DevOps, SQL Server, Tableau (assumed for dashboards), Flink Watermarks, Power BI, SAS, SPSS
Data Engineering: Real-Time Data Pipelines, Data Lake Architecture, Schema Partitioning, Data Cleansing, Data Governance, Data Integration, Kimball Methodology, Async Event-Driven Processing
Collaboration: Agile, Cross-Functional Teaming, Dashboarding, Model Presentation, Healthcare Analytics (Patient Care & Predictive Insights), KPI Reporting, Financial Data Analytics / Risk Modeling

PROFESSIONAL EXPERIENCE

***Tenet Health* – USA** **Jan 2025 – May 2025**
Data Engineer

- Assisted in developing and optimizing ETL pipelines using Apache Spark, PySpark, Kafka, and Airflow, contributing to a 40% increase in data throughput and 25% fewer pipeline failures.
- Supported the implementation of real-time data streaming systems with AWS Kinesis, Lambda, and DynamoDB, helping reduce latency from 15 minutes to under 2 minutes and enabling faster analytics.
- Contributed to the design of a data lake on AWS S3 with schema partitioning and Flink watermarks, improving query performance by 35% while lowering storage costs by 18%.
- Helped automate batch and streaming workflows on Databricks and Snowflake, achieving 99.9% uptime and reducing operational overhead by 20%.
- Assisted in implementing data validation, unit testing, and alerting systems with Datadog and Prometheus, increasing data accuracy to 99.5% and reducing incident resolution time by 30%.
- Collaborated with data science teams to build dashboards in Tableau and Power BI, supporting KPI reporting and improving healthcare analytics outcomes by 15%.
- Helped integrate ML models via SageMaker and Scikit-learn into production pipelines, enabling predictive analytics that reduced patient readmission rates by 10%.

IBM PVT LTD – India **April 2017 – July 2023**
Data Engineer

- Accelerated API performance by 30% by reengineering backend services with Python and Flask, leveraging async event-driven processing to handle 50K+ daily active user requests with lower latency.
- Streamlined event processing for 5M+ daily records by developing a real-time data pipeline using Apache Flink, Kafka, and Redis, enabling near-instant updates to dashboards and alerting systems.
- Reduced manual deployment time by 85% through CI/CD automation with GitHub Actions, integrated with unit tests and data validation, supporting faster iteration across environments.
- Cut ETL job failures by 60% by restructuring Airflow DAGs, applying schema validation, and optimizing Snowflake partitioning logic, resulting in more stable ingestion of financial data.
- Boosted anomaly detection coverage by 40% by building a Python-based test suite that combined unit testing, data validation, and real-time monitoring hooks for ML outputs.
- Enhanced model performance by integrating Scikit-learn and Gradient Boosting (GBM) into pipelines, delivering 12% higher predictive accuracy for operational risk models.
- Optimized distributed data transformations with PySpark and Databricks, reducing large-scale processing runtimes by 35% and supporting scalable analytics workloads.

EDUCATION

Master of Computer Science (Data Analyst) **June 2026**
Westcliff University
Master of Science (Electronics) **Mar 2019**
Sri Indu College of Engineering and Technology
Bachelor of Science (Instrumentation) **Mar 2017**
Sri Indu College of Engineering and Technology