

CHANDRA SEKHAR THOLLA

Cloud DevOps Engineer | Data Engineer | AI Infrastructure Specialist
Oklahoma City, OK | +1 (316) 665-5540 | Chsekhui@gmail.com
LinkedIn: <https://www.linkedin.com/in/chandra9989/> | Portfolio: <https://chandracloud.net>

PROFESSIONAL SUMMARY

Result-oriented Cloud & AI Data Engineer with 5+ years of experience designing scalable infrastructure and data pipelines on GCP, AWS, and Azure. Deep expertise in Unified Data Analytics (Databricks, Snowflake), Container Orchestration (Kubernetes/Docker), and Generative AI integration. Proven track record of handling the full SDLC—from gathering requirements to production deployment—ensuring 99.9% uptime and strict security compliance. Adept at bridging the gap between Development, Operations, and Data teams to deliver cost-effective, high-performance solutions.

TECHNICAL SKILLS

Cloud Platforms:	GCP: AWS:	• Compute Engine, Cloud Run, GKE, Cloud Functions, BigQuery, Cloud Build, IAM, VPC. • EC2, S3, Lambda, RDS, DynamoDB, EKS, VPC Peering, CloudWatch, Route53.
Data Engineering & Analytics:	Platforms: ETL/Orchestration: Visualization:	• Databricks (PySpark/Delta Lake), Snowflake, Google BigQuery, AWS Glue. • Apache Airflow, Alteryx, dbt (Data Build Tool), Kafka (Streaming). • Tableau, Power BI, Oracle OTBI/BIP, Looker Studio.
DevOps & Infrastructure as Code (IaC):	Tools: Containerization:	• Terraform, Ansible, Jenkins, GitHub Actions, GitLab CI, ArgoCD. • Docker, Kubernetes (K8s), Helm Charts.
AI & Automation:	GenAI: Scripting:	• OpenAI API (GPT-4), Gemini, LangChain, RAG Pipelines, Vector Databases (Pinecone/ChromaDB). • Python (Boto3, Pandas), Bash/Shell, SQL, Node.js.
Monitoring & Tools:	Observability: Management:	• Splunk, Prometheus, Grafana, Datadog, ELK Stack. • Jira (Agile/Scrum), Confluence, ServiceNow (Incident Management), Git.

PROFESSIONAL EXPERIENCE

Cloud AI Architect (SaaS Product) | <https://chandrahere.github.io/chandracloud> | Remote | Oct 2025 – Present

Designing a comprehensive SaaS ecosystem that integrates Cloud, Data, and AI services.

- End-to-End Architecture: Built a microservices-based architecture on Google Kubernetes Engine (GKE), utilizing Docker containers for modular deployment of web apps and AI agents.

- **Databricks & Data Lake:** Implemented a Lakehouse architecture prototype to process user activity logs. Used PySpark on Databricks to perform transformations and stored processed data in BigQuery for real-time analytics.
- **DevOps Automation:** Configured Cloud Build triggers to automate testing and deployment pipelines in GCP. Implemented "GitOps" workflows where every code commit automatically updates the infrastructure via Terraform.
- **AI Agent Integration:** Developed automated customer support agents using Gemini Pro and Python, capable of handling natural language queries to deploy infrastructure (e.g., "Launch a VM for me").
- **Security & Compliance:** Enforced Organization Policies, setup VPC Service Controls, and managed secrets using Google Secret Manager to ensure zero-trust security.

Cloud DevOps & Data Engineer

I2U Systems Inc | Dallas, TX | Feb 2023 – Present

Managed cloud infrastructure and data pipelines for enterprise clients, ensuring high availability and scalability.

- **Data Migration & Modernization:** Led the migration of on-premise SQL databases to Snowflake and AWS Redshift. Utilized AWS DMS (Database Migration Service) for minimal downtime.
- **Big Data Processing:** Engineered scalable ETL pipelines using Databricks (Apache Spark) to ingest and process Terabytes of streaming data from IoT devices. Optimized Delta Lake tables for faster query performance.
- **Infrastructure Optimization:** Managed AWS EKS (Kubernetes) clusters. Wrote Terraform modules to provision VPCs, Subnets, and Security Groups, ensuring Infrastructure-as-Code (IaC) best practices.
- **CI/CD Implementation:** Designed Jenkins and GitHub Actions pipelines to build, test, and deploy React/Node.js applications. Integrated SonarQube for code quality checks and Trivy for container vulnerability scanning.
- **Production Support:** Handled L3 support tickets via Jira/ServiceNow. Used Splunk and CloudWatch logs to troubleshoot critical production issues, reducing Mean Time To Resolution (MTTR) by 40%.

Data Engineer & BI Developer

Dummadu Games Pvt Ltd | Bengaluru, India | May 2019 – Dec 2021

Collaborated with cross-functional teams to deliver business intelligence solutions.

- **ETL Development:** Developed complex workflows in Alteryx and Python to cleanse and blend data from multiple sources (Oracle ERP, Excel, CSVs). Automated daily data loads into the Data Warehouse.
- **Reporting & Visualization:** Designed interactive executive dashboards in Tableau and Oracle OTBI, providing visibility into Procurement (P2P), HR, and Finance metrics.
- **Database Management:** Optimized complex PL/SQL queries and Stored Procedures in Oracle Fusion Cloud to improve report generation speed by 60%.
- **Team Collaboration:** Participated in daily Agile/Scrum standups, sprint planning, and retrospectives to align technical deliverables with business goals.

KEY PROJECTS

Real-Time Log Analysis Platform (ELK Stack + Kafka)

- Built a centralized logging system using Kafka for log ingestion and Elasticsearch for indexing. Created Kibana dashboards to visualize application error rates in real-time.

AI-Powered IT Support Bot (RAG Pipeline)

- Developed an internal tool using Python, LangChain, and Vector DBs to ingest internal documentation. The bot answers employee technical queries instantly, reducing manual ticket volume.

EDUCATION & CERTIFICATIONS

- **Master of Science in Computer Science** | Wichita State University, KS | 3.68/4 GPA | May 2023
- Plan and prepare to develop Microsoft Azure AI
- Oracle Fusion Technical Consultant
- K8S Certified, HashiCorp Certified: Terraform Associate (In Progress)
- Databricks Certified Data Engineer Associate (Self-Study)