Chandra Mouli Cherukuri Visweswararao

Data Analyst

chandramouli858304@gmail.com | +1 (781) 956-0628 | Boston, USA | LinkedIn

Summary

Data Analyst with over 4.5+ years of experience building scalable data pipelines, predictive models, and real-time dashboards using Python, SQL, and cloud technologies (AWS, Azure, GCP). Skilled in statistical analysis, A/B testing, and machine learning, including regression and time series forecasting. Proven ability to partner with cross-functional teams to deliver actionable insights that drive strategic decisions. Holds an M.S. in Computer Science from UMass Boston with strong expertise in data governance, data quality, and stakeholder collaboration.

Skills

Languages: Python (Pandas, NumPy, Seaborn, Plotly, Scikit-learn), SQL (PostgreSQL, Snowflake, BigQuery), R, Bash
Analytics & Visualization: Tableau, Power BI, Looker, Streamlit, Matplotlib
ETL & Orchestration: Apache Airflow, dbt, Fivetran, Informatica, Talend, PySpark, Kafka
Cloud Platforms: AWS (Redshift, S3, Lambda, SageMaker), Azure (Synapse, Data Factory, Databricks), GCP (BigQuery, Dataflow)
Data Management: Data Modeling, Data Cleaning, Alation, Collibra, Data Quality Audits
Statistical & ML Methods: A/B Testing, Hypothesis Testing, Linear/Logistic Regression, Time Series (ARIMA), Market Basket Analysis
Other Tools: Git, JIRA, Agile/Scrum, Docker (basic), Excel (Advanced Pivoting/Power Query)

Professional Experience

Data Analyst - Remote — 03/2024 – present | Massachusetts, USA
Dell Technologies
- Spearheaded optimization of complex ETL pipelines using Apache Airflow and Python, improving data throughput by 42% across 20+ workflows that ingested over 1 million Redshift records per month.
- Built and deployed AI server demand forecasting models using Scikit-learn and ARIMA, improving capacity planning accuracy by 18% and reducing overprovisioning costs by 10%.
- Created interactive Tableau and Power BI dashboards for key initiatives like Clara AI and NRF 2025, integrating GeoPandas for location-based visual insights; dashboards auto-refresh every 4 hours for real-time decision-making.
- Conducted cohort analysis on Dell Pro AI Studio adoption using Snowflake SQL and Dask, uncovering usage trends by segment (e.g., by region, role, frequency) that shaped go-to-market strategy.
- Delivered strategic recommendations using A/B testing results, supported by exploratory data analysis (EDA) with Plotly and Seaborn, helping senior leadership prioritize features in upcoming quarterly releases.

Tools: Python, Apache Airflow, Redshift, Snowflake, GeoPandas, Tableau, Power BI, ARIMA, Dask, Plotly

Data Analyst — 01/2022 – 12/2023 | Massachusetts, USA
University of Massachusetts Boston
- Designed and implemented PySpark pipelines that processed over 500GB of student engagement and performance data daily, reducing data load and transformation time by 35% and enabling near real-time analytics.
- Tuned complex PostgreSQL and MySQL queries, resulting in a 40% improvement in dashboard response time for academic performance reports viewed by 200+ stakeholders, including deans and department heads.

- Developed Power BI dashboards integrated with Azure Synapse to track enrollment trends, dropout rates, and funding distribution across 50+ programs, improving strategic academic planning.
- Implemented real-time anomaly detection in financial aid and student billing systems using Kafka and Spark Streaming, enabling alerts for suspicious transactions with sub-second latency.
- Oversaw AWS-based infrastructure (EC2, Lambda, Redshift) to ensure cost-effective, scalable, and secure compute for long-running analytics jobs and ad hoc query support.

Tools: PySpark, Azure Synapse, Power BI, PostgreSQL, MySQL, AWS Lambda, EC2, Redshift, Kafka, Spark Streaming

Associate Data Analyst 08/2020 – 12/2021 | Ahmedabad, India
MoneyGram

- Conducted detailed analysis of over 8 million customer support logs using Python and SQL, identifying root causes of delays and reducing average ticket resolution time from 3.2 to 2.4 days (25% improvement).
- Built Looker dashboards and dbt models to monitor SLA compliance and ticket escalation trends across 15+ support datasets, helping support leadership identify underperforming regions.
- Mapped and visualized 30,000+ geocoded tickets using GeoPandas, revealing technical issue clusters in high-volume corridors and guiding infrastructure enhancements in Asia-Pacific markets.
- Applied graph analysis with NetworkX to identify frequent sequences in 5K+ chatbot queries, leading to a 30% boost in intelligent routing accuracy and resolution rate.
- Created Kafka-driven AWS Lambda functions to automate alerts for operational triage, reducing response lag to high-severity events by 40%.

Tools: Python, SQL, Looker, dbt, GeoPandas, AWS Lambda, Kafka, NetworkX

Education

M.S. in Computer Science 01/2022 – 12/2023 | USA
University of Massachusetts Boston

B.S. in Computer Science Engineering 05/2017 – 08/2021 | India
REVA University

Certificates

Data Analysis with Python