

SRI CHARAN CHITTURI

Location: FL | Email: charanchitturi399@gmail.com | Phone: (813) 603-0353

SUMMARY

- Experienced **Data Analyst** with over 2 years of expertise in building and optimizing scalable data pipelines across healthcare and financial domains.
- Specialized in **ETL/ELT development** using **Databricks**, **Apache Spark**, and **Airflow** to process structured, semi-structured, flatfiles, and CSV data exceeding **25M records**.
- Skilled in managing **cloud-native platforms** including **AWS (Glue, S3, Lambda, RDS)**, **Azure Synapse**, and **Google BigQuery**.
- Proficient in designing high-quality data workflows with **schema evolution**, **data validation**, **deduplication**, and **anomaly detection**.
- Strong knowledge of **data warehousing**, dimensional modeling, **star/snowflake schemas**, and analytics optimization practices.
- Hands-on experience with **MLOps** and **CI/CD pipelines** using **GitHub Actions**, **Terraform**, **Jenkins**, and **Docker** for automated and secure deployments.
- Integrated and tested **REST**, **GraphQL**, and **Webhook APIs** to ingest third-party data from **ERP**, **CRM**, and marketing platforms.
- Applied **DevSecOps** practices including **RBAC**, encryption, logging, and remediation workflows to meet **HIPAA** and **GDPR** compliance standards.
- Strong in **OLTP/OLAP systems**, **ODBC/SQL**, **XML**, **JSON**, **Flatfiles**, and **Linux/Unix scripting** for full-cycle data engineering support.

TECHNICAL SKILLS

Programming Languages:	Python, SQL, R, Scala, Java, TypeScript
Databases & Data Warehousing:	MySQL, PostgreSQL, Oracle, SQL Server, MongoDB, Cassandra, Snowflake, Delta Lake, Amazon Neptune, DynamoDB, HBase, Hive, Cloudera, OLTP, OLAP, RDBMS, HDFS, Flat Files, CSV
Cloud Platforms & DevOps Tools:	AWS (S3, Glue, Lambda, Kinesis, RDS, Athena, Firehose, EventBridge, SNS, SQS, CloudWatch), Azure (ADLS, Synapse), Google BigQuery, Docker, Kubernetes, Jenkins, Terraform, Ansible, PowerShell, ZooKeeper
ETL & Data Integration:	Talend, Informatica, SSIS, Apache NiFi, Matillion, SAP, REST APIs, GraphQL, Webhooks, Data Quality Checks, Schema Validation, Data Profiling, Matching Logic, Schema Evolution, Ad-Hoc Queries
Data Visualization & Reporting:	Power BI, Tableau, MS Excel (Macros, VBA), SSAS
Big Data & Streaming Frameworks:	Apache Spark, Spark Streaming, Apache Flink, Kafka, MapReduce, Databricks, Teradata
Testing, Compliance & Deployment:	Unit Testing, Data Validation Frameworks, CI/CD, DevSecOps, HIPAA, GDPR, Code Reviews, Vulnerability Assessments.
Methodologies & Architecture:	SDLC, Agile, Waterfall, MLOps, Data Mesh, Data Fabric
Additional Skills:	CRM Systems Integration, Graph Algorithms, Data Architecture, Business Intelligence, Effective Communication

PROFESSIONAL EXPERIENCE

KPMG	India
Data Analyst	Aug 2022 – Jul 2023
- Engineered 12+ ETL/ELT pipelines using Databricks, PySpark, and SQL, processing 25M+ financial records per batch from flatfiles, APIs, and CSV sources. - Constructed 8+ Delta Lake tables with schema evolution and ACID compliance, supporting audit analytics with 4-year data traceability and full CDC capability. - Orchestrated data ingestion via Azure Data Factory, integrating sources from PostgreSQL, MongoDB, and Snowflake into layered architecture across 5 projects. - Automated 40+ Airflow DAGs daily, boosting pipeline SLA compliance to 98% and cutting manual interventions by 14+ hours/week. - Integrated 6 REST API connectors for ERP, CRM, and banking systems, eliminating 100% of prior manual file imports across 3 data teams. - Formulated graph-based fraud models in Amazon Neptune, enabling detection of 12+ anomaly patterns from live banking logs processed via Spark Streaming. - Provisioned CI/CD infrastructure with Jenkins, Terraform, and Azure DevOps, accelerating release cycles by 3x with rollback-safe IaC. - Delivered 6 Tableau dashboards with KPI-based alerts, supporting real-time monitoring for 20+ stakeholders across finance and audit divisions. - Enforced HIPAA/GDPR compliance by embedding RBAC and encrypting sensitive data with Azure Key Vault, impacting 100% of production pipelines. - Implemented serverless JSON ingestion into Snowflake using AWS Lambda and API Gateway, improving data arrival time by 70% for external sources.	

- Executed scalable analytics with **Spark Core, Spark SQL, and Spark Streaming**; profiled data using **Pandas, NumPy**, and **Matplotlib**, covering 10+ audit projects.
- Collaborated with **Splunk** engineers to monitor ETL events and establish **360° pipeline observability**, decreasing resolution time by **4 hours per issue**.

Cipla
Data Analyst

India
Aug 2021 - Jul 2022

- Architected analytical layers in **Amazon Redshift** and **Snowflake** to monitor **200+ product KPIs**, improving regulatory audit readiness for CDSCO and WHO.
- Ingested 10M+ nested JSON and flatfiles from **AWS S3** using **Glue + Lambda triggers**, integrating sales feeds from 15+ hospitals and vendor systems.
- Optimized prescription log processing using **AWS Step Functions** and **PySpark**, slashing reporting latency from **7 days to under 24 hours** across 3 units.
- Engineered Spark Streaming jobs to detect batch recall events in real-time, improving traceability and response for **30+ manufacturing sites**.
- Structured OLTP schemas and OLAP cubes for on-demand analytics, reducing query time by **60%** for forecasts across **200 SKUs**.
- Automated 8 Power BI dashboards with region-wise metrics for respiratory, oncology, and specialty therapies, covering **15+ zones** and enabling national sales insights.
- Executed deduplication and anomaly detection for **10M+ batch records**, increasing batch quality assurance effectiveness by **40%**.
- Enabled secure ingestion via **API Gateway + Lambda**, reducing manual interventions and ensuring 100% API traceability with audit logs.
- Configured Jenkins pipelines on **Linux servers** with rollback logic triggered by **Splunk alerts**, enhancing deployment safety for all supply chain tools.
- Instituted HIPAA-compliant encryption protocols and **RBAC** policies with **AWS IAM + KMS**, protecting sensitive data for 5+ operational environments.
- Coordinated CI/CD through **Azure DevOps (VSTS)** for environment-specific deployments across 4 pharmaceutical data applications.

EDUCATION

The University of Tampa Masters in Information and Technology Management	Tampa, FL, USA Aug 2023 - May 2025
Koneru Lakshmaiah Education Foundation Bachelor of Technology in Electronics and Communication Engineering	Andhra Pradesh, India Jul 2018 - May 2022

CERTIFICATIONS

- AWS Certified Data Engineer Associate
- AWS AI Practitioner
- Databricks Certified Data Engineer Associate
- Oracle Database SQL Certified Professional

PROJECTS

Generative AI Research Engineer | The University of Tampa | Tampa, Florida

- Designed and deployed scalable ETL pipelines in Databricks to clean and structure 1M+ clinical records for downstream NLP processing.
- Integrated BERT, T5, and GPT models using Hugging Face Transformers to perform summarization and question generation on patient data.
- Developed RESTful endpoints for real-time inference via AWS SageMaker, Azure ML, and Bedrock.
- Implemented schema validation and data profiling logic to ensure compliance and accuracy of model input.
- Created testing frameworks to validate model predictions and deployed them with Docker and GitHub CI/CD workflows.
- Applied MLOps techniques to track model lineage, automate deployments, and enable rollback with model version control.
- Leveraged Kubernetes to containerize models for scalable access during high-volume medical queries.
- Partnered with healthcare analysts to design dashboards visualizing model accuracy and data distribution.

Data Quality and Observability Platform

- Engineered a modular observability platform using Python and Airflow to detect schema changes, missing values, and data duplication.
- Automated lineage capture and profiling across 20+ Delta Lake and Snowflake tables using custom DAGs.
- Set up Grafana dashboards and Slack integrations to provide real-time anomaly alerts for Airflow tasks.
- Implemented SQL-based validation rules to benchmark column-level data quality against thresholds.
- Developed PowerShell and Bash scripts to monitor AWS Glue job logs and automate error remediation workflows.
- Reduced pipeline downtime by creating recovery scripts and health checks with metadata validation triggers.
- Collaborated with DevOps engineers to integrate observability alerts into Terraform-managed Airflow environments.