

Anirudh Korukanti

USA | [linkedin.com/in/anirudhkorukanti](https://www.linkedin.com/in/anirudhkorukanti) | +1 (850) 405-7140 | anirudhkorkanti27@gmail.com | [Credly](#)

SUMMARY

Data Engineer with 4+ YOE building scalable real-time and batch data pipelines using streaming, cloud, and ETL tools to drive insights in healthcare, education, and telecom.

EDUCATION

Florida State University, Tallahassee, FL, USA Aug 2022 – May 2024
Master’s in computer science.
Course Work: Artificial Intelligence, Advanced Data Science, Data and Computer Communication, Concurrent Parallel and Distributed Programming, Network Security, Active and Passive Defenses.

PROFESSIONAL EXPERIENCE

Florida Department of Health | Data Engineer Sep 2024 – Present

- Engineered an end-to-end **Apache Kafka pipeline** to ingest Avro-encoded telemetry from 50+ medical devices across 10 hospitals.
- Developed **Spark Structured Streaming** jobs on **AWS EMR** to process and aggregate event windows (1-minute tumbling) and persisted results to **Redshift** for near-real-time dashboards.
- Reduced end-to-end latency from device to visualization from 15 minutes to under 2 minutes, meeting SLA > 99.5%.
- Automated nightly **ETL using Airflow**: 20+ DAGs orchestrating **Python/SQL** scripts to extract EHR and insurance claims from on-prem Oracle, transform (data quality checks, deduplication, schema validation), and load into S3 in Parquet format.
- Designed **Datadog** and **CloudWatch** dashboards to track pipeline throughput, job durations, memory utilization, and data freshness.

Florida State University | Data Engineer Sep 2022 – May 2024

- Structured and published canonical data models for student performance, enrollment, and semester metrics.
- Created an automated metadata registry in **PostgreSQL**, generating JSON schemas for incoming CSV feeds and validating against Avro definitions.
- Developed **Python-based ETL pipelines** to process 100K+ daily records from campus systems (Banner, Blackboard) into **PostgreSQL and Oracle**.
- Containerized and deployed Java/Spring Boot microservices on Kubernetes** to expose RESTful endpoints for dashboarding teams; reduced manual data requests by 60%.
- Designed and launched **Tableau** and **Power BI** dashboards on top of star-schema models, enabling drill-through by department, demographic, and grant program.

Verizon Data Services | Data Engineer I Jan 2022 – Aug 2022

- Re-architected legacy shell-scripts into Dockerized Spring Boot services running on Kubernetes; processed 10 M+ CDR (call detail records) daily.
- Leveraged **Apache Kafka** for event-driven ingestion of network KPIs into HDFS; implemented custom serializers/deserializers for Avro payloads.
- Designed partition strategies for **Oracle OLAP tables**, reducing query times by 50% for 1-year historical trend analysis.
- Automated statistics collection and indexing maintenance via **PL/SQL jobs**, improving overall system throughput.
- Authored detailed data pipeline design documents, API contracts, and runbooks for cross-team collaboration and onboarding.

Cognizant Technology Solutions | Software Engineer – Data Integration Jan 2021 – Jan 2022

- Developed a distributed data pipeline using **Kafka** and **Spark Streaming** to process 200K+ financial transactions per day.
- Enriched incoming records with reference data from Redis and applied anomaly detection logic using **PySpark UDFs**.
- Designed checkpointing and windowed aggregations to ensure fault tolerance and accurate real-time analytics.
- Built nightly Spark batch jobs that summarized transactional data and saved them in Parquet format to S3, partitioned by region and transaction date.
- Integrated **AWS Athena** for ad-hoc querying by the analytics team, reducing report generation from hours to minutes.
- Created modular Jenkins pipelines with stages for build, test, deploy, and rollback across multiple environments.
- Wrote reusable Terraform modules to provision **AWS Lambda, EC2 instances, IAM roles, and RDS** for internal tooling.
- Maintained **Git** branching strategy and enforced PR reviews, ensuring stable and auditable infrastructure changes.
- Built a log aggregation pipeline using **Fluent Bit and Splunk** to centralize logs from Kafka, Spark, and application containers.

TECHNICAL SKILLS

Programming Languages: Java, Python, C++, C#, R, JavaScript, SQL, HTML, CSS, Bash, Shell Scripting, Scala
Cloud & DevOps: AWS (EC2, S3), Azure, GCP, Docker, Kubernetes, Git, GitLab, Jenkins, Terraform, DevOps, CI/CD Pipelines
Monitoring & Observability: Dynatrace, Splunk, Open Telemetry, Prometheus, Grafana, Kibana, SonarQube, Distributed Tracing, Structured Logging, Logging Schema Design
Data & Analytics Tools: MySQL, Oracle DB, SQL Server, MongoDB, Power BI, Tableau, SAS, Apache Spark, Datadog
Project Management & Support: Jira, ClearQuest, Zendesk, Agile (SAFe)
Streaming & Orchestration: Apache Kafka, Confluent Platform, Apache Airflow, AWS Glue, AWS Kinesis

PROJECTS

Student Enrollment Data Pipeline

- Designed SQL queries, views, and stored procedures to transform and integrate 500K+ records of registration, performance, and course data, enabling automated term-wise reporting across 30+ departments.
- Improved Power BI dashboard performance by 65% and reduced ad-hoc data request time by 70% through query optimization using CTEs, indexing, and execution plan tuning.