

Ragamai Senior Data Engineer

Email: ragamaic@gmail.com

Phone: +1 (602) 767-2576

LinkedIn: [Ragamai Ch](#)

Professional Summary:

- Around 8+ years of experience as a Senior Data Engineer, specializing in designing, developing, and implementing ETL pipelines for large-scale enterprise applications using Azure, Databricks, and Snowflake.
- Expertise in Azure Big Data technologies, including Azure Databricks, Data Lake Storage, Data Factory, Synapse Analytics, and Delta Lake, implementing Medallion Architecture for structured data processing.
- Strong hands-on experience in Databricks Workflows, Delta Live Tables (DLT), Unity Catalog, and SQL-based data processing, ensuring high-performance ETL transformations.
- Experience in developing batch and streaming data pipelines using PySpark, Spark SQL, and Python, optimizing data ingestion, transformation, and analytics.
- Good working experience with Azure Data Factory (ADF) for orchestrating and scheduling complex data workflows across cloud environments.
- Hands-on experience working with Delta Lake and Databricks Delta Tables, ensuring ACID-compliant, high-performance data processing.
- Worked with Kafka for event-driven architecture, implementing real-time data streaming pipelines for large-scale analytics solutions.
- Designed and deployed high-performance batch & streaming data workflows across Spark, Scala, Kafka, and Databricks.
- Strong understanding of SQL query optimization and performance tuning in Azure Synapse, Snowflake, and Databricks SQL for improving data processing efficiency.
- Exposure to Microsoft BI tools including SSIS and SSAS, along with working knowledge of DBT and FiveTran for modern data stack transformation and ingestion workflows.
- Experience in data governance, metadata management, and role-based security controls, ensuring compliance with enterprise data policies.
- Expertise in NoSQL databases such as MongoDB, DynamoDB, and HBase, designing and implementing distributed data storage solutions.
- Familiar with Agile Scrum methodology, sprint planning and SDLC best practices.
- Strong ability to work effectively in cross-functional teams, collaborating with business stakeholders and data analysts to translate business needs into scalable data solutions.
- Proficient in Python, SQL, Scala, and Java, with experience in ETL automation, API integrations.
- Experience in building data visualizations and dashboards using Power BI, enabling data-driven decision-making for business teams.
- Excellent interpersonal and communication skills, with a proactive approach to problem-solving, time management, and multitasking in fast-paced environments.

Technical Skills:

Big Data Technologies: HDFS, Map Reduce YARN, Hive, Pig, HBase, Impala, Sqoop, Oozie, Kafka, Drill, Spark and Avro.

Programming Languages: Python, Scala, SQL, Java, PL/SQL, Linux shell scripts

RDBMS: Oracle, MySQL, SQL server, Teradata, Snowflake.

Cloud Technologies: MS AZURE and AWS

Tools Used: SSIS, SSAS, DBT, FiveTran, Eclipse, Putty, Cygwin, MS Office, Elasticsearch, Logstash, Kibana, Splunk.

BI Tools: Tableau, Power BI.

Professional Experience:

Kroger – Blue Ash, Cincinnati

Senior Data Engineer

Jul 2025 – Present

Responsibilities:

- Supported Kroger's Manufacturing Modernization initiative by designing and developing pipelines to migrate legacy ERP data from multiple plants into a centralized Infor M3 system.
- Built and optimized ETL pipelines using Azure Databricks, PySpark, and Delta Lake following the Medallion Architecture (Bronze, Silver, Gold).
- Orchestrated and monitored ingestion workflows from on-premises and cloud sources using Azure Data Factory (ADF).
- Integrated ADF pipelines to ingest Infor M3 ERP data into Azure Data Lake Storage (ADLS) for downstream processing.
- Worked on using multiple ADF Triggers and worked on setting the Integration run times.
- Developed and maintained Delta Live Tables and Databricks Workflows for incremental data loads and transformations.
- Worked on optimizing spark jobs to improve the runtimes and bring down the SLA's and cloud compute cost.
- Used Terraform and CI/CD pipelines (GitHub Actions) to automate Databricks workspace, cluster, and resource provisioning across environments.
- Implemented Unity Catalog for governance, access control, and table lineage tracking throughout Databricks.
- Created schemas and access policies in Unity Catalog, managing table-level permissions and workspace entitlements.
- Configured Databricks Genie for testing and validating conversation queries, monitoring SQL Warehouse performance, and verifying user permissions.
- Worked on building visualizations using Databricks Dashboards and Power BI Dashboards.
- Worked on ingesting data from multiple sources and loaded the data into Data Lake and Vector search tables.
- Worked on building custom chatbots on the internal knowledge and integrated it with MCP and slack using Lang Chain, ML Flow which served as a SME bot for the Business and production support team.
- Designed and optimized data models in Databricks SQL to support analytics and Power BI dashboards.
- Maintained GitHub branches and deployment documentation to support version control and team collaboration.

Environment: Azure (Databricks, Data Lake Storage, Synapse Analytics, Data Factory), Spark, Databricks Workflows, Delta Live Tables (DLT), Medallion Architecture, Apache Spark (PySpark, Spark SQL), Delta Lake, JSON, CSV, Excel, Python (Pandas, NumPy), SQL, Unity Catalog, GenAI, Lang Chain, ML Flow, GitHub, CI/CD Pipelines, Power BI.

ADM – Erlanger, Kentucky

Senior Data Engineer 2025

Jun 2024 – May

Responsibilities:

- Designed and developed ETL pipelines in Azure Databricks, implementing Medallion Architecture for structured data processing.
- Hands on experience working with Azure Databricks and worked on completely implementing the new Lake House architecture with in Databricks.
- Automated data workflows using Databricks Workflows and Delta Live Tables (DLT) to improve efficiency and reliability.
- Worked with Unity Catalog, DLT, Delta tables and used DB-SQL extensively. Worked on creating Genie rooms for Generative AI use cases.
- Developed Spark applications using Scala and Java (Spring Boot) for batch data ingestion and real-time processing within Databricks.
- Assisted data analytics teams in implementing DBT models within Snowflake and Databricks environments, focusing on SQL-based transformation and testing.

- Integrated FiveTran as part of ingestion architecture for automated extraction from cloud-based SaaS sources into ADLS and Snowflake.
- Designed and deployed real-time streaming pipelines by integrating Kafka, Azure Event Hubs, and Databricks Structured Streaming to process business-critical data with sub-second latency.
- Migrated legacy batch pipelines to a cloud-native, modular architecture using Azure Data Lake Gen2, Delta Lake, and Databricks workflows.
- Implemented SDLC best practices including version control with GitHub, code reviews, CI/CD automation, and deployment pipelines in Azure DevOps.
- Worked on consuming the data from API's, Databases and File systems and stored the data on ADLS.
- Developed SQL-based ETL queries on source systems to validate output and ensure data integrity before transformation.
- Engineered forecasting models in Databricks to predict future sales, optimizing data-driven decision-making.
- Experience working with Delta table and performance tuning of the spark jobs on Databricks cluster.
- Conducted data validation and testing across multiple layers to ensure accuracy and consistency in analytical outputs.
- Performed data catalog management, optimizing tables, views, and metadata for improved data governance and compliance.
- Executed ad-hoc data cleansing tasks, eliminating redundant objects and improving system performance.
- Developed KPI-driven reports and dashboards, transforming raw data into actionable insights for business users.
- Worked in a cross-functional environment, collaborating with business analysts, managers, and leadership teams to align data solutions with business goals.
- Provided technical support and debugging expertise, resolving pipeline failures and improving system reliability.
- Expertise in data validation, testing, and performance tuning of large-scale data processing workloads.

Environment: Azure (Databricks, Data Lake Storage, Synapse Analytics, Data Factory), Spark, Databricks Workflows, Delta Live Tables (DLT), Medallion Architecture, Apache Spark (PySpark, Spark SQL), Delta Lake, JSON, CSV, Excel, Python (Pandas, NumPy), SQL, Unity Catalog, GitHub, CI/CD Pipelines, PowerBI.

BNY Mellon – New York City, New York

Jan 2022 - May 2024

Data Engineer

Responsibilities:

- Built data ingestion workflows to extract, transform, and load data from various source systems into Snowflake and cloud storage.
- Developed SQL-based data validation processes, ensuring accuracy and consistency across Snowflake, Azure Data Lake, and other storage systems.
- Implemented Change Data Capture (CDC) solutions using Snowflake Streams and Tasks, enabling efficient incremental data processing.
- Developed Spark and Spark Streaming jobs in Scala to process and transform large datasets for Snowflake integration.
- Utilized Snowpark for Python to develop scalable, object-oriented data transformation logic within Snowflake.
- Implemented data quality rules and data control mechanisms, including threshold alerts and exception handling to maintain data integrity.
- Migrated on-prem solutions from MapR to Azure cloud.
- Built and optimized real-time streaming pipelines using Confluent Kafka, integrating with Snowflake for analytics.
- Managed data governance, cataloging, and metadata optimization, improving data discoverability and compliance.
- Created data models and schema structures to optimize storage, indexing, and retrieval for Snowflake tables.
- Configured Kafka topics across environments and used Spark Structured Streaming for near real-time data ingestion and processing.
- Automated data validation and quality checks, improving downstream reporting accuracy.
- Developed real-time and batch data transformation scripts using Python, SQL, and Snowflake scripting.

- Integrated data pipelines with external systems, enabling seamless data exchange between Snowflake, Azure, and third-party platforms.
- Applied dimensional data modeling concepts (Star/Snowflake schemas, Fact/Dimension tables) to design efficient analytics structures.
- Implemented workflow orchestration using Apache Airflow and Autosys for scheduling and dependency management.
- Collaborated with cross-functional teams in an Agile environment to design and implement scalable, business-driven data solutions.

Environment: Azure (SQL, ADLS, BLOB Storage, Azure SQL Server, HDInsight Snowflake (Streams, Tasks, Stages, Snow SQL), Confluent Kafka, Spark, Scala, Python, SQL, Map Reduce, HDFS, Hive, Impala, Sqoop, Tableau, Power BI, Apache Airflow, Autosys, Teradata, Data Modeling (Star Schema, Snowflake Schema, Fact & Dimension Tables).

AT&T, Birmingham, AL
2021

Sept 2019 - Dec

Big Data Engineer

Responsibilities:

- Extracted the needed data from the SFTP servers into HDFS and Bulk Loaded the cleaned data into HBase.
- Supporting Continuous storage in AWS using Elastic Block Storage, S3, Glacier. Created Volumes and configured Snapshots for EC2 instances.
- Worked on migrating Hadoop solutions that are running on Cloudera Distribution of Hadoop to run on AWS.
- Estimated the Software & Hardware requirements for the Name Node and Data Node & planning the cluster.
- Written the Map Reduce programs, Hive UDFs in Java.
- Successfully Generated consumer group lags from Kafka using their API
- Worked on a python script to extract data from Netezza databases and transfer it to AWS S3
- Developed Lambda functions and assigned IAM roles to run python scripts along with various triggers (SQS, SNS)
- Designed and deployed end-to-end data pipelines using AWS Glue and AWS Lambda to automate ingestion, transformation, and storage of structured data in S3 and Redshift.
- Created and managed infrastructure components using Terraform for provisioning S3 buckets, IAM roles, and Lambda configurations.
- Developed CI/CD pipelines using GitHub Actions and integrated with Terraform for continuous delivery of IaC templates.
- Developed Automation Regressing Scripts for validation of ETL process between multiple databases like AWS Redshift, Oracle, Mongo DB, T-SQL, and SQL Server using Python.
- Architected several DAGs (Directed Acyclic Graph) for automating ETL pipelines.
- Created a Lambda Deployment function, and configured it to receive events from S3 buckets
- Experience in fact dimensional modelling (Star schema, Snowflake schema), transactional modelling and SCD (Slowly changing dimension)
- Developed Airflow DAGs in python by importing the Airflow libraries.
- Developed and Configured Kafka brokers to pipeline server logs data into Spark streaming.
- Played a key role in productionizing the application after testing by BI analysts.
- Maintain System integrity of all sub-components related to Hadoop.

Environment: Apache Hadoop, HDFS, Spark, Kafka, Hive, Cassandra, Map Reduce, Pig, Java, Flume, Cloudera CDH4, Oozie, Oracle, MySQL, AWS.

eClerx - India

June 2016 - May 2019

Data Engineer/ Hadoop Developer

Responsibilities:

- Created Hive tables, loading the data and writing hive queries that will run internally in a map reduce way. Developed a custom File System plugin for Hadoop so it can access files on Data Platform.
- Developed Simple to complex Map/reduce Jobs using Hive and Pig.
- Involved in ETL, Data Integration and Migration.
- Used Impala for querying the ad hoc datasets.
- Imported data using Sqoop to load data from Oracle to HDFS on regular basis.
- Written Hive queries for data analysis to meet the business requirements.
- Creating Hive tables and working on them using Hive QL. Experienced in defining job flows.
- Developed Java Map Reduce programs for the analysis of sample log file stored in cluster.
- Used Sqoop to import data into HDFS and Hive from other data systems.
- Setup and benchmarked Hadoop/HBase clusters for internal use.

Environment: Apache Hadoop, HDFS, Cloudera Manager, CentOS, Java, Map Reduce, Eclipse, Hive, PIG, Sqoop, Oozie and SQL.

Educational Details:

Bachelor's degree in computer science engineering, JNTUH

2012 - 2016