

Dhananjaya Kumar Venkata Mittapalli

Boston, MA | +1 8578328906 | mittapalli.d@northeastern.edu | [linkedin.com](#) | [github.com](#) | [Portfolio](#)

SUMMARY

Data Engineer with 4+ years of experience designing and implementing scalable data pipelines, ETL/ELT workflows, and cloud-based data architectures on AWS, Azure, and GCP. Expert in building automated data orchestration solutions using Apache Airflow, Python, and SQL to support real-time analytics, business intelligence, and machine learning initiatives. Proven track record of reducing processing time by 40%, improving data quality, and delivering high-performance data warehouse solutions. Skilled in big data technologies, workflow automation, and cross-functional collaboration with business stakeholders to drive data-driven decisions.

EDUCATION

Northeastern University, Boston, MA

Sept 2023 - Present

Master of Science - Data Science

Related Courses: Data management and processing, Supervised and Unsupervised Machine Learning, Algorithms, MLOps, LLM

PROFESSIONAL EXPERIENCE

Natixis Investment Managers

Boston, MA, USA

Data Engineer

Jan 2025 - Aug 2025

- Engineered and deployed end-to-end data pipelines using Python, SQL, and Rundeck to integrate trade activity data, AUM movements, and CRM data into Snowflake data warehouse, enabling real-time executive dashboards and business intelligence.
- Automated ETL workflows data orchestration processes for revenue pipeline refresh processes using workflow orchestration tools, reducing daily reporting cycles by 40% and improving data quality and reliability across all downstream systems.
- IMPACT:** Streamlined executive reporting workflow from manual data compilation to automated pipeline delivery, enabling real-time access to critical business metrics for investment decision-making and reducing data-to-insights time by 60% while maintaining 99%+ data accuracy across all revenue and AUM tracking systems.

SaksOFF5th

Bangalore, India

Supply Chain Business Insight Engineer

May 2022 - Aug 2023

- Designed and implemented supply chain data warehouse architecture capturing end-to-end KPIs from procurement to delivery, building dimensional data models and star schema while performing statistical analysis and data mining techniques to derive actionable business insights for operations teams.
- Built automated data pipelines and ETL workflows using SQL, Python, and Apache Airflow to streamline data ingestion from multiple sources and load processed data into centralized warehouse, reducing processing time by 30%.
- IMPACT:** Reduced report generation time from 2 days to 4 hours (75% improvement) through pipeline automation and optimization, enabling Chief Supply Chain Officer to launch multiple strategic business initiatives with real-time data insights.

TheMathCompany

Bangalore, India

Data Engineer

Sep 2020 - May 2022

- Spearheaded a team to **productionalize data migration solutions** with **AWS S3 staging** and automated the entire **ETL data pipeline** via **Airflow hosted on an EC2 instance**. Ensured the solutions to be both high available and scalable.
- Designed and developed Tableau and MicroStrategy dashboards to visualize marketing KPIs, enabling the marketing team to monitor performance effectively and make data-driven decisions.
- IMPACT:** Achieved **140% increase** in **execution speed** and maintained **less than 2% downtime** for all business critical triggers. Fostered a collaborative environment through clear and concise communication between the internal team and external clients. Expanded team size, secured and managed multiple new projects due to clients' high satisfaction with work.

TECHNICAL SKILLS

Languages	Python, R, SQL, Java, Scala, Bash, Shell Scripting
Orchestration	Apache Airflow, DataIKU, DBT, Rundeck, Fivetran, Databricks, Data Pipeline Development
Cloud Platforms	AWS (S3, EMR, EC2, Lambda, Glue, CloudWatch, SageMaker), Azure (Synapse Analytics, Data Lake, Data Factory), GCP (BigQuery)
Databases	Snowflake, Redshift, MySQL, PostgreSQL, MongoDB, NoSQL, Oracle, SQL Server
Big Data	Apache Spark, Hadoop, MapReduce, PySpark, Distributed Computing, Hive, HDFS
DevOps	Docker, Kubernetes, GitHub, CI/CD, Jenkins, Terraform, Ansible, Infrastructure as Code
Data Visualization	Tableau, PowerBI, Looker, MicroStrategy
Data Quality	Data Validation, DVC, Schema Validation, Data Monitoring, Data Governance

ACADEMIC PROJECTS

Multiclass Disease Classification

Sep 2024 - Dec 2024

Northeastern University

Boston, MA

- Built reproducible **ETL workflows** leveraging Docker, Airflow, and DVC for data versioning, orchestration, and deployment on **GCP**. Implemented preprocessing (resizing, grayscale, CLAHE), augmentation (rotations, zooms), and schema validation using TFDV to ensure data quality.
- Automated end-to-end machine learning pipeline with hyperparameter tuning, slice-based bias evaluation, and model performance monitoring using GitHub Actions for continuous integration and deployment, achieving over **90% classification accuracy** while reducing manual model iteration time by **50% through automation**.

Smart City

Feb 2024 - Apr 2024

Northeastern University

Boston, MA

- Engineered big data processing pipelines using **Python**, **Apache Spark**, and **Pandas** to analyze urban datasets addressing housing affordability, energy consumption patterns, and public transportation efficiency in Vancouver, processing millions of records and implementing distributed computing for scalable data transformations and analytics.