



Jerald John | | jeraldjohn765@gmail.com
937-303-5707 [LinkedIn](#)

Lead Data Engineer | Cloud | Supply Chain | Life Insurance | Medical & HR Data sets

Summary

Results-driven Cloud Data Engineer over 9 years of experience designing, architecting, and optimizing cloud-native data platforms and pipelines across GCP, Azure, and AWS. Specialized in ETL/ELT workflows, RAG (Retrieval-Augmented Generation) pipelines for AI/LLMs, and cloud infrastructure engineering to support enterprise-scale analytics, regulatory reporting, and HR/insurance data processing. Proven track record in migrating legacy systems to GCP, implementing real-time and batch pipelines, and ensuring compliance with SOX, HIPAA, and EEO regulations. Adept at building secure, scalable, and cost-optimized cloud data ecosystems that power data-driven insights across Supply Chain ERP, HR, and Life insurance domains.

Technical Skills

Cloud Platforms: GCP (BigQuery, Cloud Storage, Cloud Composer, Vertex AI, IAM, Monitoring), Azure Synapse, Azure Data Factory, Azure Data Lake, AWS Redshift, AWS Glue, AWS S3, Snowflake, Databricks

Programming & Data Processing: Python, SQL, Golang, PySpark, UNIX Shell Scripting

Data Engineering & AI: ETL/ELT workflows, RAG pipelines for LLMs, workday, vector databases (FAISS, Pinecone), embeddings, orchestration (Airflow, Composer)

Data Governance & Security: IAM, encryption, masking, lineage, audit logs, SOX/HIPAA/EEO compliance

Visualization & Reporting: Power BI, Tableau, Grafana, QuickSight, Alteryx

Workflow & DevOps: Jenkins, GitLab CI, Azure DevOps, CI/CD pipelines, JIRA, Confluence

Databases: PostgreSQL, SQL Server, Teradata, AWS Redshift

Agile Practices: Scrum, Kanban

Professional Experience

Arisoft Global

Lead Data Engineer | Remote | 08/2024 – Present

Led a team of 5 data engineers in designing and implementing cloud-native data pipelines on GCP, modernizing supply chain ERP datasets across procurement, inventory, and logistics.

Architected the enterprise data model in BigQuery, optimizing queries through table partitioning, clustering, and cost-efficient design — achieving a 45% improvement in performance and cost savings. Spearheaded the migration of SAP HANA and Teradata workloads to GCP, ensuring minimal downtime and maintaining compliance with SOX and HIPAA standards. Migrated ERP operational datasets from on-prem Teradata and SAP HANA to GCP BigQuery, applying partitioned & clustered tables for faster analytics and reduced query costs. Built RAG pipelines with Vertex AI + Big Query, enabling AI-driven search and audit automation.

Designed real-time streaming frameworks using Kafka + Dataflow for ERP transaction tracking (POS, inventory, shipments), enabling near real-time visibility for supply chain operations.

Oversaw CI/CD automation and infrastructure-as-code (IaC) via Jenkins and GitLab CI, embedding validation, lineage, and automated testing for high reliability in production.

Leveraged Python/PySpark, SQL, and Golang to analyze large datasets and data modeling, supporting strategic initiatives and cross-team decision-making.

Implemented CI/CD pipelines with Jenkins and GitLab CI, embedding validation, lineage, and unit testing for production reliability. Implemented Kubernetes (GKE) for containerized data processing workloads.

Led cost optimization initiatives by monitoring BigQuery slot utilization and Cloud Storage lifecycle policies, reducing monthly cloud spend by 30%..
Applied row-level security and masking policies in Power BI for SOX compliance. Optimized cloud infrastructure and cost by monitoring BigQuery slots and leveraging Cloud Storage lifecycle management.

Cardinal Health, OH

Sr Data Engineer | Aug 2023 - Jul 2024

Mentored the Developers and enhanced data models in GCP BigQuery to support healthcare and pharmaceutical analytics, ensuring high data quality and performance.
Designed and implemented Airflow DAGs to automate workflows and trigger BigQuery stored procedures, optimizing ETL processes.
Created, maintained, and documented technical specifications for data models to ensure consistency and cross-team usability across multiple PODs.
Partnered with product owners and data architects to gather, analyze, and translate complex business requirements into scalable data solutions.
Migrated data objects and pipelines to higher environments using Jira for version control, deployment tracking, and release management.
Built and improved Tableau dashboards and reports, enabling stakeholders to visualize KPIs and make data-driven business decisions.
Worked with Python, SQL, and BigQuery stored procedures to perform data transformations, cleansing, and optimization for analytical workloads.
Leveraged GCS, AtScale, and GCP services to build scalable BI and analytics solutions within a cloud-first data architecture.
Monitored and troubleshoot pipelines using Airflow, ensuring accuracy, reliability, and SLA adherence.

Infosys (McCamish Systems)

Data Engineer/Process Specialist | Feb/2019 – Aug/2021

Contributed to Infosys McCamish Insurance Platform (VPAS® Life), managing life insurance and annuity product data for global insurers.
Designed ETL pipelines to process policy, claims, and premium data; migrated legacy insurance datasets into Azure and GCP warehouses.
Partnered with actuarial teams to deliver analytics-ready datasets for premium forecasting and claims risk analysis. Integrated SQL-based data validation steps within AWS Glue ETL jobs to ensure high-quality, audit-ready data pipelines.
Delivered 30+ analyses weekly using SQL queries, visualizing user metrics and trends for life insurance claims data in Redshift, which enabled data-driven decisions for the care management team.
Developed insurance-specific data modeling supporting regulatory, actuarial, and profitability reporting.
Automated regulatory compliance reporting workflows, reducing manual effort by 40%.
Built Power BI dashboards for policy KPIs, claims ratios, and product performance.
Applied cloud-based data governance and access controls to ensure security and compliance with insurance regulations.

Software Engineer

Themosoft | Jun/2016 – Feb/2019

Designed ETL pipelines integrating HR data from Workday, SuccessFactors, and PeopleSoft into centralized Azure and GCP data lakes.
Automated ingestion and transformation of employee lifecycle data (onboarding, payroll, attrition, performance) into data marts for People Analytics.

Developed Power BI and Tableau dashboards for workforce planning, turnover, DEI metrics, and compensation analysis.

Leveraged Python, SQL, and Spark to analyze large datasets, supporting strategic initiatives and cross-team decision-making.

Automated data transformation processes with Pandas, NumPy, and PySpark; managed efficient data storage using SQL Alchemy with PostgreSQL and AWS Redshift.

Implemented Airflow-based orchestration for HR pipelines, ensuring on-time and validated reporting delivery.

Built predictive attrition and workforce forecasting models by partnering with HR data scientists.

Migrated on-prem HR datasets to Snowflake and Azure Data Lake, enabling real-time analytics for HR stakeholders.

Enforced data validation, lineage tracking, and compliance frameworks for EEO and SOX audits.

EDUCATION

Master in Computer Science – Christian Brothers University | Memphis, TN 2021 – 2023

Bachelor of ICE– Anna University | Coimbatore, India | 2013 – 2016

CERTIFICATIONS

Google Cloud Certified – Associate Cloud Engineer

SQL of Data Science Certification