

Summary

As a Data Engineer with over 4 years of experience designing and maintaining scalable data pipelines, **ETL workflows**, and **cloud data architectures** across healthcare domains. Skilled in building data solutions using **SQL**, **Python**, **Spark**, and cloud platforms including **Azure and AWS**. Proven ability to collaborate with cross-functional teams to automate data ingestion, improve data accuracy **by 40%**, and optimize performance of analytics solutions **by 30%**. Experienced in working with structured, semi-structured, and unstructured data in both **batch** and **streaming** environments. Strong foundation in **data warehousing**, **data lake integration**, and **dashboard** development for clinical and claims-based reporting.

Skills

- **Languages & Tools:** SQL, Python, Spark (PySpark), Bash, Git, Airflow, Azure DevOps, Power BI, Tableau
- **Cloud Platforms:** Azure (Data Factory, Synapse, Blob Storage, Data Lake), AWS (S3, Glue, Redshift, Lambda)
- **Data Technologies:** Hadoop, Hive, Kafka, Databricks, Snowflake, Delta Lake
- **ETL/ELT Tools:** Apache Airflow, Informatica, Talend, Azure Data Factory
- **Databases:** SQL Server, PostgreSQL, MySQL, Oracle, MongoDB
- **CI/CD & Orchestration:** Jenkins, GitHub Actions, Azure DevOps Pipelines
- **Others:** REST APIs, JSON, Parquet, Avro, Agile, JIRA

Education

- Masters in Computer Science from Texas Tech University, TX

Professional Experience

Data Engineer | Elevance Health | Aug 2024 - Current

- Developed and maintained scalable **Azure**-based data pipelines using **Data Factory** and **PySpark** to process **200K+** records per day with **99.9%** accuracy.
- Improved **ETL** performance **by 35%** through optimization techniques and modularizing transformation logic.
- Automated ingestion of **JSON** and CSV files from SFTP and **REST APIs**, reducing manual intervention **by 60%**.
- Designed data models in **Azure Synapse** and implemented views and stored procedures for analytics teams.
- Enabled analytics on patient claims and medical data by integrating data from on-prem to **Azure Data Lake**.
- Created **Power BI dashboards** to visualize operational **KPIs**, improving reporting speed **by 40%**.
- Developed data validation scripts using **PySpark** and **SQL**, which reduced data anomalies in production **by 45%**.
- Integrated real-time claims data processing using **Azure Event Hub** and Stream Analytics for faster fraud detection workflows.
- Participated in daily **Agile** ceremonies and provided input for sprint planning, improving story completion rate **by 20%**.
- Maintained version control and **CI/CD** pipelines using **Azure DevOps** and **GitHub** Actions for consistent deployments.
- Collaborated with **QA** and **Data Analysts** to test and validate data quality using **PyTest** and custom **SQL** test cases.

Tools & Tech: Azure Data Factory, Azure Synapse, Data Lake Gen2, PySpark, Power BI, Git, SQL Server, REST API, Airflow, JIRA

Data Engineer | Cerner India (Oracle Health) | Sep 2020 -Jul 2023

- Built robust data pipelines using **Spark** and **Airflow** to integrate clinical and **EHR** data from multiple sources into a centralized **data warehouse**.
- Designed and deployed **schema-on-read** structures in **Hadoop/Hive**, improving query flexibility for healthcare reporting.
- Used **AWS Glue** and **S3** to store and transform data for downstream analytics; reduced report generation time **by 30%**.
- Developed parameterized **ETL workflows** in Airflow, enabling configurable and reusable data ingestion jobs.
- Implemented incremental data loading using **CDC** logic, reducing data load time **by over 50%**.
- Created detailed technical documentation for all **ETL pipelines** and data dictionaries to support knowledge sharing.
- Scheduled automated **data quality checks** using Great Expectations and integrated results with Slack for alerting.
- Conducted root cause analysis for production failures, resulting in a **30%** decrease in pipeline downtime.
- Identify data quality issues and resolved over **90%** of critical data mismatches.
- Participated in **sprint planning**, peer code reviews, and documentation activities in **Agile** environment.
- Created **SQL** views and optimized **joins** and indexing to improve query speed **by 25%** in large datasets.

Tools & Tech: Spark, Hive, Airflow, AWS Glue, S3, Redshift, SQL, Git, Hadoop, PostgreSQL, Python, JIRA