

Kausha Trivedi

+1 774-486-9668 | kaushatrivedi12@gmail.com | <https://www.linkedin.com/in/kaushat/>

EDUCATION

Northeastern University , Boston, MA	Jan 2023 - Apr 2025
Master of Science in Computer Science	GPA: 3.85/4.0
Relevant Coursework: Algorithms, Distributed Systems, Advanced Software Engineering, Cloud Computing	
Indus University , Ahmedabad, Gujarat	Sept 2018 - May 2022
Bachelor of Science in Computer Science	GPA: 4.0/4.0

WORK EXPERIENCE

GBCS Group , Remote, USA	Jun 2024 - Dec 2024
Software Developer Intern	
- Architected scalable transaction processing system handling 50K+ daily transactions using React/Node.js and Web Workers, reducing Time-to-Interactive by 55% for 95th-percentile users - Implemented CI/CD pipeline with AWS Lambda and CloudWatch, enabling zero-downtime deployments and 30% faster MTTR - Pioneered AI-driven Redis caching with predictive key expiration algorithms, reducing P99 latency by 40% under 10K+ concurrent load	
Elsner Technology , Ahmedabad, Gujarat	May 2021 - Nov 2022
Software Developer	
- Reduced AWS spend by 35% (\$216K annual) through Terraform auto-scaling and ML-based spot instance forecasting, presenting strategy to CTO - Designed fault-tolerant microservices architecture (Redis caching + Circuit Breaker pattern) boosting throughput by 40% while eliminating 70% of cascading failures - Automated secrets rotation via HashiCorp Vault and Kubernetes webhooks, eradicating credentials from 50+ microservices - Led cross-functional AI adoption guild evaluating TorchServe vs TensorFlow Serving, implementing solution reducing model deployment latency by 60% - Mentored 3 junior engineers through architecture reviews, resulting in 25% faster feature delivery and 40% reduction in critical bugs	

PROJECTS

Advanced Job Portal , React, Node.js, Elasticsearch	
- Built real-time search with Elasticsearch BM25+ relevance scoring and React virtualized lists, reducing query latency by 40% for 100K+ listings - Designed React virtualization and Redis caching layer supporting 50K+ daily transactions with 55% faster Time-to-Interactive - Secured platform with JWT token rotation and Redis rate limiting, blocking 15K+ monthly brute-force attacks	
ML-Powered Resume Analyzer , Python, FastAPI, BERT, spaCy	
- Developed BERT-based NLP pipeline achieving 92% accuracy on 50K+ resumes with custom spaCy NER components, boosting skill extraction F1-score by 18% - Engineered FastAPI microservice with concurrent processing, reducing parsing latency from minutes to <3 seconds at 500 RPM load - Built React visualization dashboard adopted by 200+ hiring managers, reducing screening time by 30%	
Distributed Real-Time Engine , C++, Rust, ZeroMQ	
- Created AVX-512 vectorized processing pipelines with Rust memory safety layers, outperforming standard implementations by 5x - Architected low-latency (<5ms) ZeroMQ messaging system handling 1M+ events/sec with exactly-once delivery guarantees - Implemented persistent write-ahead logging and snapshot isolation ensuring 99.99% uptime during network partitions	

TECHNICAL SKILLS

Languages	TypeScript, Python, Java, C++, SQL
Frontend	React, Redux, Web Vitals Optimization, Performance Tuning
Backend	Node.js, Spring Boot, REST, GraphQL, gRPC
Cloud & DevOps	AWS (Lambda, S3, EC2), Docker, CI/CD, Terraform, Kubernetes
ML/Data	spaCy, scikit-learn, TensorFlow, BERT, Hugging Face, PyTorch
Databases	PostgreSQL, DynamoDB, Redis, Elasticsearch
Practices	Microservices, TDD, MLOps, System Design, Agile