

Lahari Guttikonda

+1 7165331061 | lahariguttikonda26@gmail.com



PROFESSIONAL SUMMARY:

Results-oriented Data Engineer with 4+ years of experience designing, developing, and optimizing data pipelines and ETL workflows across cloud platforms including AWS, Azure, Databricks, and Snowflake. Skilled in Big Data technologies such as PySpark, Apache Airflow, and Kafka, with expertise in Medallion architecture, dimensional modeling, and data governance using Unity Catalog and DBT. Proven ability to deliver secure, scalable, and high-performance data solutions that enable real-time analytics, business intelligence, and ML deployment. Experienced in building automated CI/CD pipelines and cross-platform integrations to drive operational efficiency and data-driven decision-making across industries.

EDUCATION

Engineering Science (Data Science), Master of Science (MS) University at Buffalo, SUNY

December 2023

SKILLS

- Languages: Python, SQL, PySpark, Spark SQL, Scala, T-SQL, Linux.
- Cloud platforms: Azure Data Lake, Azure Data Factory, Azure Databricks, Azure Synapse, Azure SQL Data warehouse, AWS S3, AWS Redshift, AWS Glue, AWS Lambda, AWS Athena, AWS EMR.
- ETL: Kafka Streams, Databricks, Azure Data Factory, AWS Glue, Snowflake, DBT, Apache Airflow, SSIS.
- Databases: Postgres, MySQL, SQL Server, Oracle dB, Cassandra, MongoDB
- Data Visualization & Analytics: Tableau, PowerBI, Databricks SQL Dashboards, Pandas, Matplotlib, Seaborn
- Software Development: Databricks Asset Bundles, Azure DevOps, GIT, SDLC, Agile, Scrum, Test-Driven Development.
- Certifications: **AWS certified Data Engineer Associate , Azure Associate Data Engineer**

EXPERIENCE

Truist Bank | Data Engineer

February 2024 – Present

- Designed and implemented scalable data pipelines to ingest customer transaction data, credit history, account activities, loan applications, and financial product data from Postgres, Oracle DB, APIs, and third-party sources into S3 as a staging layer.
- Migrated core banking and financial transaction data from Snowflake to Databricks with minimal disruption, ensuring data consistency through validation checks, schema mapping, and compliance with financial data regulations.
- Utilized S3 as the source of truth and DLT's to ingest banking transactions, loan processing records, and customer profile data into Databricks' Bronze layer, orchestrating batch, real-time, and incremental workflows via Apache Airflow.
- Implemented real-time data streaming pipelines using DLT's and Change Data Capture (CDC) techniques to enable low-latency updates for financial transactions, fraud monitoring, and real-time account activities.
- Re-engineered 2500+ Snowflake SQL scripts for compatibility with Spark SQL, optimizing queries related to customer transactions, credit scoring, financial reporting, and fraud detection by simplifying nested queries and improving processing speeds.
- Developed a custom data validation framework in Databricks that runs daily to ensure schema integrity, validate transaction counts, and identify anomalies or NULL values in customer and financial datasets, ensuring data quality for regulatory reporting.
- Utilized Unity Catalog in Databricks for financial data governance, security, and asset cataloging, enforcing Role-Based Access Control (RBAC) and ensuring compliance with banking sector data privacy regulations (e.g., GLBA, SOX).
- Enabled data scientists to develop predictive models for credit risk scoring, customer churn prediction, fraud detection, and loan default forecasting. Deployed models using MLOps in Databricks for real-time financial decision-making.
- Optimized Databricks clusters for scalable financial data analytics using Photon acceleration, autoscaling, and resource utilization monitoring with Ganglia, enhancing performance for large-scale transaction processing and financial reporting.
- Automated workflows and transformations using DBT for financial data marts, with CI/CD pipelines for DBT models and Databricks notebooks via GitHub, accelerating development and improving accuracy in financial data pipelines.

Li-Matt Soft Solutions | Data Engineer

January 2020 - August 2022

- Developed and implemented structured streaming ETL pipelines on Databricks to process large volumes of real-time GPS and location data from devices, ensuring data completeness and accuracy using PySpark in Databricks notebooks.
- Utilized Delta Lake to manage the GPS and location data efficiently, ensuring ACID transaction support and schema enforcement. Handled complex queries and schema evolution with Spark SQL to maintain high-quality, reliable location data.
- Configured Azure Data Factory (ADF) pipelines to ingest real-time GPS and location data from multiple sources such as GPS devices, mobile apps, and cloud data storage, ensuring reliable extraction and timely processing of location data.
- Implemented a Medallion architecture to structure the location data for scalability, fault isolation, and maintainability.
- Employed PySpark to transform, clean, and aggregate location data for downstream geospatial analytics and applications.
- Conducted data quality checks to validate the accuracy and integrity of GPS and location data, while optimizing data processing workflows to reduce processing time by 30%, ensuring low-latency delivery of real-time data for analytics and tracking.
- Integrated data security and compliance features using Azure Key Vault and Databricks secret scopes to manage sensitive credentials, ensuring secure handling of location and GPS data in accordance with privacy and industry regulations.
- Monitored pipeline performance using Azure Monitor, troubleshooting and identifying bottlenecks in real-time data processing
- Configured CI/CD pipelines with Azure DevOps to automate the deployment of Databricks notebooks and ADF pipelines, ensuring seamless integration and deployment, while collaborating with teams in Agile environments to refine development processes and enhance GPS data processing.