

Summary:

- Data Engineer with over 5+ years of expertise in designing, optimizing, and deploying scalable data pipelines and architectures across enterprise-level environments. Skilled in **Big Data technologies** such as **Hadoop, Spark, and Hive** for efficient **data processing** and **analytics**.
- Proficient in **Python, Java, and SQL**, with hands-on experience in **data warehousing** solutions including **Snowflake, Redshift, and Azure Synapse**, enabling structured **data storage, query optimization**, and efficient analytics.
- Expert in **ETL development**, leveraging tools like **Informatica, Talend, SSIS, and Apache Airflow** to streamline **data integration, transformation, and automation**, ensuring high-quality, real-time insights.
- Extensive experience in cloud platforms—**AWS, Azure, and GCP**—building serverless **ETL pipelines** with **AWS Glue, Redshift, and Lambda**, optimizing processing speed, and reducing latency.
- Strong **DevOps** and **CI/CD** expertise, **utilizing Jenkins, GitHub, Docker, Terraform, and Kubernetes** to automate deployments, enhance scalability, and ensure efficient workflow management.
- Hands-on experience with real-time **data streaming** technologies such as **Apache Kafka, Amazon Kinesis, and Spark Streaming**, improving **data ingestion, reliability, and low latency analytics**.
- Proven ability in implementing **scalable architectures, automating workflows**, and optimizing distributed systems, enabling data-driven decision-making, **advanced analytics**, and **AI/ML-driven** business intelligence.

Skills:

- **Programming Languages:** Python, Java, Scala, SQL, PL/SQL, T-SQL, Unix Shell Scripting
- **Big Data Technologies:** Hadoop, Spark, Hive, Pig, Sqoop, MapReduce, HBase, EMR
- **Methodologies:** SDLC, Agile, Waterfall
- **DevOps Tools:** Git, Jenkins, Docker, Kubernetes
- **ETL Tools:** SSIS, Informatica PowerCenter, Erwin, Talend, DataStage
- **Databases:** SQL Server, PostgreSQL, MySQL, Oracle, Snowflake, DynamoDB, MongoDB
- **Pipelines:** Apache Airflow, AWS Step Functions, Luigi, Prefect, Oozie Streaming Technologies: Amazon Kinesis, Apache Spark, Apache Kafka
- **Libraries:** Pandas, NumPy, Matplotlib, SciPy, Scrapy, TensorFlow, PyTorch, Scikit-learn, NLTK, Plotly, Keras, PyMC3
- **Data Visualization:** Microsoft Excel, Power BI, Tableau, IBM Cognos, QlikView, QuickSight, Seaborn, SSRS
- **Cloud Platforms:** AWS (EC2, S3, Redshift, Glue), Azure (Azure Data Factory, Azure Databricks), GCP
- **Data Warehousing:** Amazon (DynamoDB, RDS, Athena), Azure (Synapse, Blob, Data Lake), BigQuery, Teradata, Snowflake

Education

- Masters in Data Engineering from University of North Texas - Denton, Texas.

Certification

- AWS Certified Data Engineer Associate

Professional Experience

Meta | Data Engineer | Jun 2024 - Present

- Engineered and optimized high-throughput **data pipelines** using **Python** and **Meta's internal data infrastructure**, improving **data processing** efficiency **by 30%** and enabling seamless **integration** across petabyte-scale distributed systems.
- Automated **80%** of routine **data pipeline** tasks **using Python**, resolving **over 200 data integrity** issues.
- Optimized and accelerated **data transformation** processes, **slashing latency by 40%** and **enhancing data flow efficiency** across massive datasets by architecting high-performance **Spark** based solutions within **Meta's data framework**, leveraging **AWS Glue** for serverless **ETL**.
- Developed and maintained over **20 automated data orchestration** workflows in **Apache Airflow**, utilizing complex **DAGs** to streamline **data pipeline** execution and decrease processing time **by 40%** for intricate **data transformation** tasks within **Meta's analytical ecosystem**.
- Architected and implemented event-driven **data solutions** using **AWS Lambda, integrated with Meta's real-time data streaming** platforms and **AWS services (API Gateway, SNS, SQS, S3, CloudWatch)**, enabling immediate data-driven actions and notifications.
- Engineered and optimized high-performance **SQL queries** within **Snowflake** and **Meta's data warehousing systems**, executing comprehensive data analysis to eradicate discrepancies
- across large-scale fact and dimension tables, leveraging **star** and **snowflake schema** principles to maximize efficiency and accuracy.

- Automated deployment processes using **Jenkins**, **GitHub**, and **Bash/Shell scripting** within **CI/CD pipelines**, accelerating continuous integration and delivery, reducing deployment time **by 50%**, and ensuring robust and reliable releases within **Meta's production environment**.

Fortis Healthcare | Data Engineer | Aug 2019 - Dec 2022

- Established and deployed **data pipelines** with **AWS Redshift** and **Glue**, automating **data integration** and **transformation operations**, increasing query performance **by 40%**, and optimizing analytical workflows.
- Organized **data processing** in **Python** and **Pandas**, using **MongoDB** for efficient storage and retrieval, and lowering data manipulation time **by 30%** for **large-scale datasets**.
- Implemented **Big data processing** was optimized on **Hadoop** utilizing **YARN** and **MapReduce**, which improved resource allocation and reduced job execution time **by 35%** for **large-scale data analysis** activities.
- Executed scalable and high-performance **NoSQL** solutions on **AWS** with **DynamoDB**, resulting in **30%** faster **data retrieval** and seamless scalability for **cloud-based** applications.
- Orchestrated complex **data pipelines** using **Apache Airflow**, automating workflows and reducing task execution time **by 35%**, ensuring seamless data integration and monitoring.
- Utilized **Docker** and **integrated** it with **CI/CD pipelines**, automating containerized application deployments and decreasing release times **by 50%**, resulting in consistent environments and shorter delivery cycle.
- Engineered and deployed **robust ETL processes** utilizing **SSIS** to automate **data extraction**, **transformation**, and **loading**, increasing **data integration** efficiency **by 30%** across several corporate systems.
- Manufactured distributed messaging solutions using **Kafka** for real-time **data streaming** and **ZooKeeper** for **cluster management**, increasing system stability and lowering message delivery latency **by 25%**.
- Customized and improved **large data processing** workflows using **Hive** and **HiveQL** for querying, as well as **Pig** and **Pig Latin** for data transformation, resulting in a **30%** increase in **processing speed** in a distributed setting

