

Mani Shankar Reddy | Data Engineer

Overland Park, Kansas | +1 (913)-456-9167 | manishankarreddyy113@gmail.com | [LinkedIn](#)

SUMMARY

Data Engineer with 2+ years of experience designing and optimizing scalable data pipelines on AWS and Azure. Proficient in ETL development using Spark, PySpark, Hive, and Airflow, with expertise in processing structured and semi-structured data (CSV, JSON, Parquet). Skilled in migrating and managing data across cloud platforms (S3, Redshift, Snowflake, PostgreSQL) while improving performance through partitioning, caching, and Spark-Scala optimizations. Hands-on experience with Azure Data Factory, Databricks, and Power BI for workflow automation and analytics reporting. Streamlined deployments using CI/CD pipelines (Azure DevOps, Git) and reduced manual efforts via Shell scripting. Strong collaborative problem-solver committed to delivering efficient, data-driven solutions.

TECHNICAL SKILLS

Programming: Python (PySpark, Pandas), Java, JavaScript, Html, CSS, SQL, Scala (Spark-Scala), Shell Scripting

Big Data Tools: Apache Spark, Databricks, Hive, Snowflake, Airflow, Delta Lake

Cloud Platforms: AWS (S3, Redshift, Glue), Azure (Data Factory, Databricks)

Databases & Warehousing: PostgreSQL, Redshift, MongoDB, Data Warehousing

Data Pipelines: ETL Design, Parquet/JSON/CSV

DevOps: Git, CI/CD (Azure DevOps), Docker (Basic)

BI & Visualization: Power BI, Tableau, Excel

PROFESSIONAL EXPERIENCE

Data Engineer | Cloudpolitan Technologies | Hyderabad, India

Dec 2021 - July 2023

- Built and managed scalable ETL pipelines using Spark, Hive, and AWS S3, processing terabytes of data efficiently to support analytics and reporting needs.
- Developed Spark applications in Scala and PySpark to integrate data across S3, PostgreSQL, and Redshift, ensuring reliable and consistent data movement.
- Migrated ETL workflows from PostgreSQL to Redshift, preserving data integrity and improving query performance across business-critical datasets.
- Processed structured and semi-structured data using Delta Lake, Parquet, and JSON formats within cloud environments to power downstream applications.
- Orchestrated multi-step batch jobs using Airflow and Oozie, automating dependencies and improving pipeline reliability across distributed systems.
- Implemented CI/CD workflows and data quality validations in Snowflake, increasing deployment efficiency and ensuring trusted analytics delivery.

Data Engineering Intern | Cloudpolitan Technologies | Hyderabad, India

Jun 2021 - Nov 2021

- Supported development of scalable data pipelines using Apache Spark and Hive, contributing to improved data processing and Efficiency for structured datasets.
- Assisted in automating data preprocessing tasks with Shell scripting, Hive, and Pig, reducing manual workflows and enhancing team productivity.
- Worked with AWS S3 to store and retrieve large datasets, enabling more efficient data handling and easier access for downstream analytics.
- Helped validate ETL scripts and participated in test runs, ensuring data integrity and catching issues early in the development cycle.
- Built Excel reports using pivot tables, VLOOKUP, and conditional formatting to summarize KPIs for business stakeholders.
- Contributed to building Power BI dashboards that visualized key metrics, helping business teams make data-informed decisions.
- Collaborated with engineers to troubleshoot and optimize Spark jobs, gaining exposure to performance tuning techniques in distributed environments.

Projects

Sales Analytics Dashboard | Power BI, AWS Redshift, Python, MySQL

- Developed interactive dashboards for cloud-based sales tracking using Power BI and Redshift.
- Automated ETL pipelines with Python to load data into AWS S3 and Redshift.
- Designed dimensional models in MySQL and scheduled daily backups to AWS.
- Implemented IAM roles to enforce secure access across teams.

Healthcare Data Analysis | PySpark, AWS EMR, PostgreSQL, Tableau, AWS KMS

- Processed large-scale patient datasets using PySpark on AWS EMR clusters for analytical insights.
- Built dynamic Tableau dashboards connected to AWS RDS (PostgreSQL) for healthcare reporting.
- Secured PHI data using AWS KMS encryption and dotenv based environment credential management.
- Optimized Spark job configurations on AWS to reduce processing time by 40%.

Certifications

- AWS Certified Data Engineer – Associate (Amazon Web Services) | [link](#)
- AWS Certified Solutions Architect – Associate (Amazon Web Services) | [Link](#)
- Machine Learning Pipelines with Azure ML Studio (Coursera) | [Link](#)
- Python Data Structures (Coursera) | [Link](#)
- Database Programming with SQL (Oracle Academy)
- Technology Virtual Experience Program (Deloitte)
- Python & Data Analytics Program and Data Visualization Workshop | ShapeAI in collaboration with DSC & GDC Ranchi.

EDUCATION

Masters in University of central Missouri | Warrensburg, MO

May 2025