

Meghana Kota

Data Scientist

kotameghanareddy@gmail.com | 4697561209 | Princeton, Texas | [GitHub](#)

SUMMARY

Data Scientist with hands-on experience building scalable data pipelines, automating feature engineering, and supporting machine learning workflows in cloud environments. Skilled in transforming raw data into training-ready datasets, deploying batch inference pipelines, and collaborating with data science teams to operationalize models. Experienced across Azure, Databricks, PySpark, SQL, and Snowflake with strong foundations in data modeling, CI/CD, and MLOps practices. Focused on delivering reliable, reproducible, and well-governed data systems that enable predictive analytics and business decision-making.

TECHNICAL SKILLS

Languages & Programming: Python (Pandas, NumPy, PySpark), SQL, R, Scala, C++, Java, XML, JSON, YAML

Cloud & Data Platforms: Microsoft Azure, Snowflake, BigQuery, AWS

Machine Learning & MLOps: Feature Engineering, Model Training Support, Batch Scoring Pipelines, Scikit-Learn, XGBoost, PyTorch, Azure ML, MLflow, Model Monitoring

ETL/ELT & Workflow Automation: Azure Data Factory, Databricks Workflows, PySpark, Azure Logic Apps, Function Apps, Apache Airflow

Data Engineering Frameworks: Apache Spark, Hadoop, Hive, Delta Lake, Kafka

DevOps & Automation: CI/CD (Azure DevOps, GitHub Actions), Docker, Terraform, Monitoring

Visualization & BI: Power BI, Tableau, Streamlit, Excel (Pivot Tables, VBA basics)

Soft Skills: Agile/Scrum, Stakeholder Collaboration, Requirements Analysis, Code Review, Technical Documentation, Problem Solving

PROFESSIONAL EXPERIENCE

Data Scientist, Elevance, USA

June 2025 – Present

- Built data pipelines in **Azure Data Factory, Databricks, and PySpark** to prepare healthcare datasets for machine learning workflows, supporting **daily analytics and modeling activities**.
- Developed **feature engineering scripts, data validation checks, and automated preprocessing steps**, helping the data science team **reduce model preparation time by 30%**.
- Assisted in deploying **XGBoost and PyTorch models** in Azure ML by creating scoring pipelines, batch inference jobs, and monitoring scripts, improving **model run time by 20%**.
- Added data lineage documentation and quality rules in **Azure Purview**, improving traceability and supporting **internal compliance reviews**.
- Built CI/CD workflows for ML data and scoring pipelines using **Azure DevOps and GitHub Actions**, reducing manual deployment work and improving **release reliability**.

Graduate Research and Teaching Assistant, University of North Texas, USA

May 2024 – May 2025

- Engineered scalable data pipelines using **Python, Spark, and SQL** to process **million-record institutional datasets**, improving data refresh cycles by **35%**.
- Built **Snowflake & BigQuery data marts** powering dashboards used by 200+ university leaders and faculty for strategic analytics.
- Automated ingestion and preprocessing in **GCP**, reducing manual intervention by **40%** and improving data reliability.
- Taught and mentored 350+ graduate students on Azure, Databricks, ETL/ELT, and CI/CD practices, improving project success rate to **98%**.
- Developed technical documentation and metadata catalogs to support **data governance, reproducibility, and audit compliance**.

Application Development Analyst, Accenture, India

August 2021 – July 2023

- Engineered **15+ automated data pipelines** using Python, SQL, and Logic Apps to collect, prepare, and transform business data for machine learning consumption across finance & supply chain systems.
- Built ML data prep workflows & feature pipelines that improved data quality and accuracy by **35%** and increased model training stability across forecasting use cases.
- Supported deployment and monitoring of **forecasting & anomaly-detection models** for enterprise clients, cutting operational manual review time by **40%**.
- Developed **automated scoring pipelines** using Azure Functions + Databricks, enabling real-time data enrichment and delivering outputs to dashboards and API consumers.
- Created reproducible model training environments with **Azure DevOps CI/CD**, reducing model deployment time by **60%** and

improving reliability in QA & production environments.

- Collaborated with data scientists, architects, and product owners in Agile teams to translate business ML use-cases into scalable data + MLOps solutions.

PROJECTS

Chicago Crash Data Analysis

January 2025 – May 2025

- Analyzed over **500K** crash records from Chicago (2020–2024) to identify key risk factors influencing crash severity using Python and machine learning.
- Developed and evaluated **Multinomial Logistic Regression and XGBoost** models, improving classification accuracy to **72%** using **SMOTE** for class balancing.
- Forecasted 2025 crash volume using Prophet, predicting a **21%** increase and informing safety policy recommendations.

Interactive dashboards highlighting US Native lands

January 2024 – May 2024

- Cleaned and transformed a 10,000 record dataset related to US Native Lands using **Tableau Prep** and **Python** to prepare data for in-depth analysis.
- Built and published two interactive Tableau dashboards highlighting demographic, social, economic, and housing trends, improving data accessibility for policymakers and researchers.
- Executed exploratory analysis to uncover regional disparities and patterns, contributing to informed storytelling through dynamic visualizations.

EDUCATION

University of North Texas

August 2023 – May 2025

Master of Science in Advanced Data Analytics (GPA: 3.9)

JNTUH College of Engineering

August 2017 – July 2021

Bachelor of Technology in Electronics and Communication Engineering (CGPA: 7.85)

CERTIFICATIONS

- Google Data Analytics Professional Certificate
- Microsoft Azure Fundamentals (AZ-900)
- Microsoft Power Platform Fundamentals (PL-900)

LEADERSHIP AND INVOLVEMENT

- Submitted a paper for the **SNAS** Conference 2025 in Connecticut titled "**Assessing Artificial Intelligence's Transformative Impact on Higher Education**," exploring the influence of AI on academic environments and learning methodologies.
- Co-authored and presented a research paper titled "**Exploring User Bias and Response Variability in Generative AI Systems**" at the **Southwest Decision Sciences Institute (SWDSI)** Conference 2025, Tulsa, OK, focusing on response patterns and ethical concerns in AI-generated outputs across multiple platforms.
- Represented college in the **Zone-B Basketball Competition**, achieving runner-up honors and demonstrating teamwork and competitive spirit.