Venkata Sai Varun Mooraboina

Springfield, MO
+1-213-600-8781 saivarun1125@gmail.com https://github.com/SaiVarun1125 https://github.com/SaiVarun1125

PROFESSIONAL SUMMARY

Data Engineer with considerable experience optimizing scalable data architectures and ETL pipelines across cloud platforms. Skilled in enhancing data processing speed and reliability, with key projects including:
- Building distributed data processing architectures to boost data processing speed by 47% at Molina Healthcare.
- Designing and implementing real-time data streaming pipelines to improve reporting accuracy. - Developing reusable ETL frameworks for faster deployment of data pipelines.
Applying expertise to support robust big data systems through strong engineering best practices.

TECHNICAL SKILLS

- **Languages:** Python, SQL, Scala, C#, JavaScript
- **Big Data:** Spark (PySpark, Spark SQL), HDFS, Kafka, Hive, HBase, Hadoop
- **Cloud:** AWS (S3, Glue, Redshift, EMR, Lambda), Azure (Databricks, Data Lake)
- **ETL Tools:** PySpark, AWS Glue, Databricks, Apache Nifi, Airflow
- **Data Streaming/Orchestration:** Spark Streaming, Kafka, Airflow
- **Databases:** Snowflake, MySQL, PostgreSQL, MongoDB
- **Visualization:** PowerBI, Tableau
- **CI/CD:** Git, Jenkins, Docker
- **Data Engineering:** ETL, Data Ingestion, Schema Design, Data Modeling, Data Analysis, SQL Queries

PROFESSIONAL EXPERIENCE Molina Healthcare Jan 2023 - Present

Molina Healthcare **Jan 2023 - Present**
Data Engineer *Long Beach, CA*

- Engineered a distributed data processing architecture using Databricks and PySpark, optimizing large-scale ETL processes, aligning with established best engineering practices to achieve a 47% boost in data processing speed.
- Led the development of real-time data streaming pipelines using AWS Kinesis and Glue, enhancing data ingestion and reporting accuracy, crucial for growth and multi-dimensional analysis.
- Designed and deployed reliable data lakes using AWS S3 and Azure Data Lake, resulting in increased scalability for the big data systems and a reduction in storage costs by 30%.
- Built reusable ETL frameworks on Databricks and Spark to facilitate quick deployment of new data pipelines and effective transformations, significantly reducing onboarding time for new datasets by 35%.
- Developed CI/CD pipelines using Jenkins and Git to automate testing and deployment of data workflows, supporting robust and scalable data platform integration, cutting deployment times by 25%.
- Collaborated with data scientists to create predictive analytics models by sourcing and preparing data using PySpark and AWS Redshift, enhancing data model reliability and fostering faster insights.
- Integrated and processed high-volume data from streaming sources into Snowflake, using schema design principles to create real-time dashboards in PowerBI, which improved visibility into key performance indicators.

Syscon Solutions **May 2021 - Jun 2022**
Data Engineer *Hyderabad, India*

- Designed and built highly optimized data pipelines utilizing AWS Glue, PySpark, and EMR for efficient data ingestion and transformations, boosting speeds by 30% for reliable big data methodologies.
- Developed an end-to-end data integration solution, consolidating over 10 disparate data sources into AWS Redshift, reducing data latency and enhancing systems for multidimensional query analysis by 25%.
- Migrated legacy data systems to AWS S3 and Azure Data Lake, implementing scalable schema designs, which led to a 30% reduction in operational costs while supporting extensible data framework solutions.
- Implemented real-time streaming data pipelines using AWS Kinesis for robust data processing, significantly improving real-time analytics and informed decision-making processes.
- Automated ETL workflows using Apache Airflow, reducing manual data processing and improving data analysis reliability by 45%, aligning with best practices in reliable data job execution.
- Built Spark Streaming applications that handled large-scale data ingestion and processing, promptly reducing average delay of data availability, essential for core product analysis, from hours to minutes.
- Enhanced financial data processing workflows using PySpark and SQL, implementing durable data modeling improvements to increase financial reporting accuracy by 12% and reduce processing time by 27%.

PROJECTS Data Migration to Snowflake

Data Migration to Snowflake

- Led a migration project from legacy on-prem databases to Snowflake using PySpark and Databricks, improving query processing speeds by 47.

AWS-based Real-Time Streaming Pipeline

- Built a real-time data streaming solution utilizing AWS Kinesis, EMR, and PySpark, enabling real-time data ingestion and processing for faster decision-making.

CERTIFICATIONS

- **AWS Associate Data Engineer**
- **Microsoft Certified: Azure Data Engineer Associate**
- **Oracle Cloud Infrastructure Certified Foundations Associate**

EDUCATION

University of Central Missouri
Bachelors, Computer Science