

NAGA RENUKA KANDI

knrenu4601@gmail.com | Boston, USA | <https://www.linkedin.com/in/naga-renuka-kandi/> | (+1) 774-636-8864

Associate-level AI/ML Engineer with 2 years of experience in Python, SQL, and cloud analytics. Skilled in building AI-powered prototypes, predictive models, and data pipelines through hackathons, academic projects, and enterprise analytics. Passionate about responsible AI and applying emerging technologies to real business problems.

EDUCATION

Master of Professional Studies, Analytics, Concentration in Applied Machine Intelligence September 2023 – April 2025
Northeastern University [Boston, MA], GPA: 3.91/4

Bachelor of Electrical and Electronics Engineering, CGPA: 7.98/10 June 2018 - July 2022
VNR Vignana Jyothi Institute of Engineering and Technology, [Hyderabad, India]

Relevant Courses: Predictive Analytics, Data Mining Applications, Enterprise Analytics, Python & Analytics Technology, Data Management & Big Data, Communication/Visual Data Analysis, Applications of AI, Probability and Statistics.

SKILLS

Programming & AI: Python (Pandas, NumPy, Scikit-learn, TensorFlow, PyTorch), SQL, R

Machine Learning & AI Applications: Predictive Modeling, Time Series Forecasting, Statistical Analysis, A/B Testing, Generative AI (Gemini API, Perplexity Sonar API), Prompt Engineering

Visualization & Reporting: Tableau, Seaborn, Plotly, Data Storytelling, Dashboard Design

Cloud Platforms: AWS (EC2, S3, Glue, SageMaker, RDS, CloudFormation), Snowflake (basic)

Data Engineering & Pipelines: Databricks, Flintrock, ETL, SQL Query Optimization, Data Wrangling

Tools & Collaboration: GitHub, JIRA, Agile/Scrum, Jupyter Notebooks, Streamlit

WORK EXPERIENCE

SOFTWARE ENGINEER I Feb 2022 - August 2023
JP Morgan Chase & Co.

- Led large-scale **data migration to AWS cloud**, collaborating with engineering teams to ensure data accuracy, integrity, and scalability, enabling reliable reporting infrastructure and reducing downtime for analytics systems.
- Developed **automated SQL pipelines** for extracting, cleaning, and analyzing production data, improving reporting accuracy, and ensuring stakeholders had timely financial insights for daily operational decision-making.
- Wrote and optimized **SQL queries with multi-table joins**, resolving data inconsistencies, improving backend processing speed, and enabling faster analytics reporting for compliance, operations, and business intelligence teams.
- Diagnosed and resolved over 30+ recurring production data-related issues, improving system stability, application responsiveness, and data reliability for analytics-driven platforms supporting enterprise-wide reporting teams.
- Partnered with **DevOps engineers** to optimize CI/CD pipelines, integrating automated testing that reduced release cycles by 20%, accelerating analytics delivery, and ensuring higher quality deployments for critical reporting systems.

Key Projects:

Capstone Project – Data Analyst | Big Sky Franchise Group Jan 2025 – Mar 2025

- Built predictive Python models to forecast lead conversions, deal success, and follow-up timing, enabling franchises to optimize sales strategy, strengthen pipeline efficiency, and improve decision-making accuracy.
- Designed Tableau dashboards with AI-driven scoring, regional segmentation, and marketing insights, providing stakeholders with interactive visualizations that improved overall conversion performance across 30+ locations.

OptiPrompt – AI Prompt Optimizer | Shark Hack 2025 April 2025

- Developed an AI-powered prompt optimization tool using Gemini API and Streamlit, enabling real-time side-by-side LLM response comparisons, few-shot learning evaluation, and a clean interactive UI for improved clarity and efficiency.
- Engineered a real-time prompt evaluation interface with side-by-side LLM response comparison, improving user clarity & model alignment. Explored few-shot prompting, response analysis, and UX-driven AI tooling for sustainable adoption.

Research Sphere – Real-Time AI Research Assistant | Perplexity AI Hackathon May 2025

- Built a lightweight GenAI-powered research assistant using Perplexity Sonar API and Streamlit, delivering real-time insights, citations, and dynamic query follow-ups for researchers and knowledge workers.
- Demonstrated rapid 0→1 prototyping with Python by iterating on prompt design, integrating API feedback, and testing response evaluation methods, showcasing agile innovation under hackathon constraints.

Iowa Liquor Sales Prediction Analysis October 2024

- Designed machine learning models achieving 94% R^2 by analyzing 1.87M records, reducing inventory waste by 20%, improving revenue predictions, and uncovering demand trends for more accurate retail planning.
- Applied feature engineering and visualization techniques with Python (Pandas, Scikit-learn) to uncover demand trends, optimize pricing strategies, and deliver actionable business insights for improved revenue prediction.

CERTIFICATIONS

- AWS Certified Cloud Practitioner.
- R Statistics, Essential Training by LinkedIn Learning.
- AWS Certified Solutions Architect – Associate