

Nagarjun Chidaravalli Shanthakumar

Data Engineer

CA | 901-231-7086 | Nagarjun@myitjobmails.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Data Engineer with 4+ years of experience designing and implementing robust data pipelines and analytics solutions. Proficient in **Python**, **SQL**, and **Scala**, with hands-on expertise in **Apache Spark**, **Hadoop**, **Apache Airflow**, and **Apache Kafka** for efficient data processing. Skilled in developing **Data Lakes** and **Data Warehouses** using **Azure** and **Databricks** to support business intelligence. Experienced in creating impactful visualizations with **Tableau** and **Power BI** and automating workflows with **Jenkins**, **Git**, and **Docker** in **CI/CD** environments. Proven ability to optimize data workflows and deliver solutions in **Agile** settings, driving operational excellence and data-driven decision-making.

SKILLS

Methodology: SDLC, Agile, Waterfall

Programming Languages: C, C++, Python, R, SQL, Scala, Java, HTML, CSS

IDEs: PyCharm, Jupyter Notebook, IntelliJ IDEA, Visual Studio, NetBeans

Big Data Ecosystem: Hadoop, MapReduce, Hive, Apache Spark, Pig, HDFS, Yarn, Data Lake, Data Warehouse

Visualization and Reporting Tools: Power BI, Tableau, MATLAB, SAS, Bloomberg Terminal, Argos, Matplotlib

ETL Tools: SSIS, Apache Kafka, Airflow, DBT

Cloud Technologies and Infrastructure: GCP, Azure

Development and Operations Tools: Git, GitLab, Jenkins (CI/CD), Docker, Kubernetes, Grafana, Splunk, Jira, GitHub

Databases: MySQL, PostgreSQL, MongoDB, Neo4j, RDBMS, Snowflake

Operating Systems: Windows, Linux

EDUCATION

Master's in Information Systems

May 2025

California State University, Long Beach

Bachelor's in Computer Science and Engineering

July 2021

JSS Science and Technology University, Mysuru

EXPERIENCE

Morgan Stanley, CA | Sep 2024 – Current | Data Engineer

- Enabled faster and more accurate financial reporting by architecting and deploying **ETL** pipelines using **Apache Airflow** and **SSIS** to integrate disparate financial data sources, improving data availability by **35%** and reducing reporting delays.
- Reduced operational costs by designing a **Snowflake**-based **Data Warehouse** with highly optimized **SQL queries**, cutting query execution time by **25%** and lowering cloud compute expenses tied to financial analytics workloads.
- Accelerated strategic decision-making by delivering executive-level **Tableau dashboards**, providing real-time insights on revenue, cost allocations, and budget utilization, improving finance team responsiveness and scenario modeling accuracy.
- Enhanced financial data consolidation and compliance through scalable **Azure Data Factory** pipelines, ensuring seamless integration of cloud and on-premise **ERP**, billing, and ledger systems—minimizing manual reconciliations and audit risks.
- Lowered data infrastructure expenses by 40% by optimizing **Spark** workloads on **Databricks** using **Delta Lake** and adaptive query execution, streamlining large-scale processing of transaction and financial forecasting data.
- Developed incremental **ETL** pipelines using **CDC** and timestamp logic, reducing data processing load by **40%**, and integrated **Git-based CI/CD** for automated, version-controlled deployments.

Infinite Infolab, India | Jan 2020 – July 2023 | Data Engineer

- Architected and optimized end-to-end big data pipelines within the **Hadoop ecosystem** (**HDFS**, **MapReduce**, **Hive**), leveraging advanced partitioning and compression strategies to reduce batch processing latency by **30%**.
- Engineered a fault-tolerant, high-throughput real-time data ingestion framework using **Apache Kafka** integrated with **Azure Data Lake Storage Gen2**, enabling seamless and reliable streaming from heterogeneous legacy systems.
- Designed and deployed dynamic, drill-down **Power BI** dashboards with **DAX** measures and role-based security, enabling executive leadership to track operational KPIs in real-time and reduce reporting latency.
- Automated robust data validation and quality checks using **Databricks Delta Live Tables (DLT)** with declarative expectations, ensuring schema compliance, referential integrity, and reducing manual QA overhead by 45% through seamless **CI/CD integration**.
- Led cross-functional **Agile** data engineering sprints with **Jira** and **GitLab**, managing backlog grooming, **sprint planning**, and continuous integration of feature branches, resulting in consistent delivery of scalable data solutions aligned.
- Designed and managed high-performance **MongoDB** clusters optimized for semi-structured data workloads, implementing sharding and indexing strategies to support real-time, read-optimized analytics across customer-facing applications.

CERTIFICATIONS

AWS Certified Cloud Practitioner

Tableau Desktop Specialist