

Neha Gatla | Data Engineer

Email: nehagatla1838@gmail.com | Phone: +1 (913) 206-2870 | LinkedIn: linkedin.com/in/nehagatla

Professional Summary

Data Engineer with 5+ years of experience designing, optimizing, and automating data pipelines across healthcare, finance, consumer goods, and life sciences. Skilled in **SQL, Python, Java, and PySpark** with expertise in **data warehousing, cloud platforms (AWS, Azure, IBM Cloud), and AI/ML pipeline integration**. Strong background in **data governance, lean process improvement, and performance tuning**. Adept at leveraging tools such as **Cloud Pak for Data, Db2 Warehouse, Watsonx, Informatica, and Airflow** to deliver scalable, secure, and efficient enterprise solutions.

PROFESSIONAL EXPERIENCE

PWC | Data Engineer | Kansas, Missouri, USA | February 2024 – Present

- Applied **lean methodologies** to streamline ETL workflows in Azure Data Factory, eliminating redundant processing and improving efficiency by 40%.
- Designed and automated **AI and ML data pipelines** with Spark and Databricks, enabling predictive analytics for financial datasets.
- Integrated **REST-based and SOAP-based APIs** along with message queues, automating ingestion of semi-structured and legacy system data.
- Built scalable data models in **Snowflake and Synapse**, optimizing query performance for analytical reporting.
- Leveraged **object-oriented Python and Java** for reusable transformation frameworks, reducing manual overhead.
- Partnered with vendors and consultants to migrate legacy **Oracle and Teradata workloads** to modernized databases such as MongoDB and Cosmos DB.

CareVet | Data Engineer | Missouri, USA | July 2023 – January 2024

- Automated **data wrangling** from multiple sources including APIs, MongoDB, Oracle databases, and flat files into AWS S3 using Glue and Spark.
- Developed real-time pipelines with **Kafka Streams and Spark Streaming**, enabling faster veterinary data analytics.
- Implemented **ML model tracking and reproducibility** with GitLab CI/CD, ensuring consistent deployment across environments.
- Built scalable **data lake governance frameworks**, ensuring schema consistency and compliance across ingestion pipelines.
- Applied **process improvement techniques** by automating infrastructure provisioning with CloudFormation and Ansible, cutting deployment time by 50%.
- Collaborated with clinicians and consultants to deliver **predictive models** for patient outcomes using Python and scikit-learn.

Roche | Application Developer/Data Engineer | Hyderabad | June 2021 – November 2022

- Migrated legacy ETL workflows into **Azure Data Factory and Hadoop ecosystems**, enabling more scalable clinical data integration.
- Designed Spark Databricks jobs with Scala, reducing analytics latency by 35% and improving trial data efficiency.

- Built OLAP cubes and designed **data warehouses in Oracle and Teradata**, supporting regulatory analytics and multidimensional reporting.
- Created **AI-ready datasets** by blending structured and unstructured healthcare data using Alteryx and Pig scripting.
- Applied **Watsonx and Cloud Pak for Data concepts** to enable scalable model deployment and integration with business processes.
- Standardized **metadata management practices**, improving lineage tracking and documentation for compliance across teams.

Hindustan Unilever | Data Engineer | Hyderabad, India | July 2019 – May 2021

- Optimized **Snowflake and Hadoop warehouses**, applying indexing and partitioning strategies to reduce costs and improve speed.
- Automated **end-to-end CI/CD deployments** with Jenkins, Git, and custom Python scripts, ensuring high availability in AWS environments.
- Designed scalable pipelines using **Kafka, Hive, and Hadoop Distributed File System**, enabling batch and near-real-time data processing for consumer analytics.
- Applied **object-oriented principles in Python and Java** to build modular ETL scripts with reusable business logic.
- Leveraged exposure to **Db2 Warehouse and IBM Cloud technologies** in modernization efforts for large-scale analytics.
- Partnered with stakeholders to identify **waste in manual reporting processes**, implementing Power BI dashboards for decision support.

Technical Skills

- **Big Data & Frameworks:** Spark, Flink, Kafka, Hadoop, Hive, Hudi, Sqoop
- **Databases:** Azure Synapse, Oracle, Cassandra, MongoDB, PostgreSQL, MySQL, Snowflake
- **Programming:** Python, SQL, Java, Scala, PL/SQL
- **Cloud Platforms:** AWS (S3, Lambda, Glue, EMR, RDS), Azure (Data Factory, Synapse, ADLS, Databricks)
- **ETL & Orchestration:** Dataflow, Airflow, Azure Data Factory, Cloud Composer
- **DevOps & CI/CD:** Terraform, GitHub Actions, Azure DevOps, Jenkins, Git
- **API & Integration:** REST APIs, SOAP APIs, API Gateway, Cloud Functions
- **Visualization & Reporting:** Power BI, Tableau, Looker
- **Other Tools:** Selenium, Pytest, Postman, JIRA, Confluence

Education

University of Missouri – Kansas City (UMKC), Kansas City, MO, USA

master's in electrical and computer engineering

Certifications

- **Microsoft Certified: Azure Data Fundamentals**
- **AWS Certified Data Engineer – Associate**