

SUMMARY

- **Data Professional with 3+ years of experience** in designing, analysing, and optimizing Data Sets, ETL pipelines, real-time streaming solutions, and cloud-native data platforms across healthcare and finance, ensuring high availability and scalability.
- Proficient in AWS Glue, Databricks, Delta Lake, PySpark, and SQL, with a proven record of improving data processing efficiency by 60% and enabling faster, more accurate analytics for critical business decisions.
- Strong background in data modeling, pipeline automation, and performance tuning, with hands-on expertise in building end-to-end data products that empower BI, machine learning, and advanced analytics teams.

TECHNICAL SKILLS

Languages:	Python, SQL (T-SQL, PL/SQL, NoSQL)
Libraries & Frameworks:	Pandas, NumPy, SciPy, Scikit-Learn, TensorFlow, Matplotlib, Seaborn
Databases:	Oracle, SQL Server, Redshift, PostgreSQL, MySQL
Big Data & ETL Tools:	Hadoop, AWS Glue, Apache Spark, Databricks, SSIS, Airflow
Cloud Platforms:	AWS (Glue, Lambda, S3, Redshift, EMR, EC2), Azure (ADF, Synapse)
Visualization & Reporting:	Power BI, Tableau, SSRS, Advanced Excel
Tools & Platforms:	Jupyter, Visual Studio, Git, MDM, Azure DevOps, GitHub Actions, CloudFormation
Data Modeling:	OLAP, SSAS Cubes, KPIs, Star & Snowflake Schema, Dimensional Modeling
Development Methodologies:	SDLC, Agile, Scrum, Kanban, Waterfall

WORK EXPERIENCE

Junior Data Engineer – Lumen Technologies, USA

Aug 2025 - Till Now

- Implemented end-to-end data ingestion and processing pipelines in PySpark running on Azure Databricks to normalize, cleanse, and transform claims datasets for ingestion into the HealthRules Payor (HRP) system, handling up to 10 million records daily.
- Developed semantic Synapse views and Power BI datasets and dashboards (DirectQuery and Composite models) to surface KPIs such as claim processing time, paid vs. denied amounts, and utilization metrics for business stakeholders.
- Improved reporting performance by 30% through targeted optimizations, partitioning, and statistics tuning in Synapse, predicate pushdown in Spark transformations, materialized views for high-cost aggregations, and query plan analysis.

Data Engineer – TCS, India

Apr 2021 – Jul 2023

- Orchestrated ELT pipelines using Azure Data Factory (ADF) to load incremental and historical claims data into Azure Synapse Analytics, leveraging Copy, Mapping Data Flows, and Databricks notebooks for transformations and partitioned writes.
- Modeled payer reporting schemas using dimensional modeling (star schema) in Synapse, designed fact tables (claims, payments), and conformed dimensions (member, provider, plan) to optimize analytic queries and reduce JOIN complexity.
- Built CI/CD pipelines in Azure DevOps to automate build, test, and deployment of Spark jobs, ADF ARM templates, and SQL artifacts, and integrated unit testing, YAML pipelines, and environment promotion to accelerate releases and enforce governance.
- Designed ETL/ELT pipelines using AWS (S3, Redshift, Glue, Lambda, EMR, Elastic Beanstalk), reducing ingestion time for datasets from APIs, databases, and SFTP, while configuring Elastic Beanstalk for scalable, high-availability deployments.
- Optimized data models using Star Schema and Slowly Changing Dimensions (SCD Type 1 & 2) in SQL Server and Snowflake to support scalable reporting solutions.
- Built centralized S3 data lake with partitioned storage & lifecycle automation, saving storage costs & boosting retrieval speeds.
- Designed Medallion Architecture layers to optimize transformations in the Lakehouse.
- Developed parameterized notebooks and Jobs API for flexible orchestration and enabled real-time analytics via streaming pipelines with S3 triggers & Delta sinks.

Data Analyst – TCS, India

Feb 2020 – Mar 2021

- Automated ETL pipelines using Python, SQL, and Azure Data Factory, streamlining Power BI data extracts and reducing report refresh time by 40%, leading to faster business insights.
- Designed and developed interactive dashboards in Power BI and Tableau to track claims processing and service performance, improving decision-making efficiency for business stakeholders.
- Applied Power Query and DAX to perform data cleansing, pivot/unpivot transformations, and advanced calculations, ensuring 99% data accuracy across BI reporting.
- Integrated and validated data from Oracle and flat files into centralized data models, enabling a single source of truth for analytics.
- Configured row-level security and access control policies in Power BI/Tableau, safeguarding sensitive financial and claims data while enabling secure role-based access.
- Performed advanced performance tuning on SQL queries and BI dashboards, reducing execution time by 35% and improving dashboard load times, enhancing user adoption.

EDUCATION

Master of Science in Information Assurance: St. Cloud State University, St. Cloud, MN

| Dec 2025

Bachelor of Science in Computer Science and Engineering: Anurag Engineering College, India

| Jul 2020