

PAKSHAL SHETH

DATA ENGINEER

MA |901-205-9384|pakshal2501@gmail.com |[LinkedIn](#)

SUMMARY

Data Engineer with 4+ years of experience architecting scalable, end-to-end ELT/ETL pipelines and managing large-scale data ecosystems in cloud environments. Proficient in Apache Spark, Databricks, Snowflake, Redshift, AWS Glue, and Airflow, with proven success in building high-performance data workflows for healthcare and business analytics. Skilled in developing robust data models, CI/CD pipelines (Terraform, Docker, Kubernetes), and real-time streaming solutions using Kafka and Spark Streaming, ensuring seamless data integration, reliability, and compliance (HIPAA). Experienced in delivering continuous improvements and rapid iterations within Agile and Scrum frameworks, collaborating cross-functionally to optimize data pipeline performance, enforce data governance, and provide actionable insights through Tableau and Power BI dashboards.

SKILLS

Programming & Scripting: Python, SQL, Bash, PySpark

Data Engineering & ETL Tools: Apache Spark, Apache Airflow, DBT, SSIS, Databricks, Delta Lake, Kafka, Kinesis, AWS Glue, Terraform

Cloud Platforms: AWS (S3, EC2, RDS, Redshift, Glue, Lambda, EMR, IAM), Azure (Data Factory, Databricks, Data Lake)

Data Warehousing & Storage: Snowflake, Amazon Redshift, Google BigQuery, PostgreSQL, MySQL, SQL Server, MongoDB, HDFS

Data Modeling & Pipeline Orchestration: Dimensional Modeling, Star/Snowflake Schema, CI/CD, GitHub Actions, Jenkins, Docker, Kubernetes

Data Analysis & Machine Learning: Pandas, NumPy, Scikit-learn, TensorFlow, SciPy, Seaborn, Matplotlib

Visualization & Reporting Tools: Tableau, Power BI, AWS QuickSight

Version Control & DevOps Tools: Git, GitHub, GitLab, Docker, Kubernetes, Jenkins

Project Management & Agile: Agile, Scrum, Jira, Confluence

Operating Systems: Linux, Unix, Windows

EDUCATION

Masters in Information Technology

University of Massachusetts Boston

Aug 2022 – May 2024

Bachelors in Computer Science

Sal Engineering and Technical Institute (Gujarat Technological University)

June 2017 – May 2021

EXPERIENCE

Cigna, MA | Nov 2023 - Current | c

- Designed and deployed robust ETL pipelines using Apache Airflow, Python, and SQL to process healthcare claims, EHR, and pharmacy data within the Hadoop and Spark ecosystems, improving processing efficiency by 40% while ensuring HIPAA compliance.
- Developed, optimized, and orchestrated batch data processing workflows using Apache Spark (PySpark, Spark SQL) within Databricks, applying Delta Lake and Databricks Workflows to process large datasets efficiently, boosting pipeline performance by 50% and ensuring scalable, fault-tolerant data pipelines.
- Engineered real-time pharmacy inventory pipelines using Apache Kafka and schema validation, enabling timely data-driven interventions that reduced medication stockouts by 25%.
- Built dynamic Tableau dashboards that automated reporting for healthcare operations and compliance teams, cutting manual effort by 30% and improving data-driven decision-making across departments.
- Improved query and ETL performance by implementing data partitioning, clustering, and indexing strategies across PostgreSQL, MySQL, and Redshift, reducing retrieval times by up to 50%.
- Constructed and refined dimensional and relational data models using Snowflake and Redshift, implementing star and snowflake schemas to support efficient querying and reporting across healthcare data domains.
- Elevated data quality by 20% through implementing comprehensive data cleaning, validation, and transformation processes using Pandas and NumPy, ensuring accuracy and consistency across healthcare datasets.
- Automated and optimized CI/CD pipelines using Terraform, GitHub Actions, Docker, and Kubernetes, streamlining the deployment of healthcare data engineering workflows and reducing release errors by 30% while ensuring secure, HIPAA-compliant infrastructure.

Infinite Infolab, India | Dec 2019 – July 2022 | Data Engineer

- Architected and implemented end-to-end ELT pipelines using SQL, Python, Apache Spark, and Hadoop, orchestrating seamless extraction of streaming and transactional data from 8+ diverse sources, loading it into cloud data warehouses (Snowflake, Redshift), and transforming it for analytics and reporting—sustaining 99.4% pipeline uptime to ensure continuous, reliable data availability for critical business operations.
- Utilized PySpark and Spark SQL to manipulate large-scale datasets, performing transformations, aggregations, and joins to prepare data for analytics and machine learning pipelines.
- Built and managed data lakes and warehouses using AWS S3, Redshift, and Snowflake, enabling centralized access to structured and semi-structured data across the organization.
- Enhanced data quality by 20% through robust cleaning and validation processes using Pandas and NumPy.
- Designed and deployed interactive dashboards in Tableau, automating dynamic visualizations and reducing manual reporting time by 30% for business intelligence teams.
- Improved query and ETL performance by implementing data partitioning, clustering, and indexing strategies across PostgreSQL, MySQL, and Redshift, reducing retrieval times by up to 50%.