

SUMMARY:

- **Data Engineer with 3+ years of experience** in Data Analysis, Data Warehousing, Big Data ,Hadoop, Spark, PySpark, Hive, HDFS, and NoSQL platforms MongoDB, Cassandra.
- **Cloud expertise** across AWS, S3, EMR, Glue, Lambda, Kinesis, Redshift, SageMaker, etc., Google Cloud Platform (BigQuery, GCS, Cloud Functions), and Azure services for scalable data pipeline development.
- **Proficient in Python, SQL, and ETL tools** Informatica PowerCenter, Talend, Apache NiFi, Airflow, Oozie for data ingestion, transformation, and orchestration.
- Skilled in **data modeling, data governance, and migrations** from RDBMS to NoSQL, as well as building complex Spark, MapReduce, and streaming jobs for large-scale analytics.
- **BI & CI/CD experience** with Tableau, Power BI, Cognos, Git, Jenkins, Docker, and Agile Scrum methodology to deliver high-quality, automated data solutions.

TECHNICAL SKILLS:

- **Tools:** Informatica PowerCenter, Informatica Power Exchange, IICS, Informatica Data Quality, Talend, Apache NiFi, Apache Airflow, Oozie, Apache Kafka, Kibana
- **CI/CD & IaC:** Git, Jenkins, Docker, Terraform, Kubernetes
- **Languages:** Python, Scala, SQL, PL/SQL, T-SQL, Unix Shell Scripting, PowerShell, HTML, XML
- **Databases:** Oracle, SQL Server, Teradata, MySQL, PostgreSQL, MongoDB, Cassandra, BigQuery, Amazon Redshift
- **Cloud Platforms:****AWS:** EMR, S3, EC2, Lambda, Glue, Redshift, RDS, SageMaker, DynamoDB, Kinesis, Athena, Lake Formation
- **Azure:** Data Factory (ADF), Databricks, Synapse Analytics
- **GCP:** BigQuery, GCS, Cloud Functions
- **Big Data Tools:** HDFS, Spark (PySpark, Spark SQL, Spark Streaming), Hive, MapReduce, Apache Pig, Sqoop, HBase, Elasticsearch
- **Data Warehousing & ETL:** AWS Glue, Apache Beam, Databricks, Talend, Informatica PowerCenter, Azure Data Factory
- **BI & Reporting Tools:** Tableau, Power BI, Cognos, SSRS, Kibana

PROFESSIONAL EXPERIENCE:

Goldman Sachs, Texas | Apr 2025 – Present
AWS Data Engineer

- Designed and implemented **end-to-end data solutions** on AWS, EMR, S3, Redshift, Lambda for storage, integration, processing, and visualization, improving analytics efficiency.
- Built and optimized **PySpark-based ETL pipelines** on EMR, integrated AWS Glue, Athena, and Kinesis for large-scale batch and streaming data processing.
- Led **data migration from RDBMS to NoSQL** (MongoDB, DynamoDB) using AWS DMS, AWS Glue, and Lake Formation for enterprise data lake deployment and governance.
- Developed and orchestrated **real-time streaming & workflow automation** using AWS Step Functions, Lambda, Apache Airflow, and SageMaker, enabling scalable ML model execution.
- Created **BI dashboards and analytics solutions** with Tableau, Power BI, Kibana, and Elasticsearch, delivering actionable insights for business stakeholders.
- Implemented **CI/CD pipelines** with Jenkins and GitLab, built security frameworks with IAM policies, and provided training on Big Data technologies to cross-functional teams.

Environment: AWS ,EMR, S3, RDS, Lambda, Glue, Redshift, DynamoDB, Kinesis, Athena, SageMaker, Lake Formation, Apache Spark ,PySpark, Spark SQL, Spark Streaming, Hadoop, Kafka, HBase, Hive, Informatica PowerCenter, Elasticsearch, Kibana, Tableau, Power BI, Python, PL/SQL, Jenkins, GitLab, Apache Airflow

Walmart, New York, NY | Jan 2024 – Mar 2025
Data Engineer

- Partnered with Product Engineering and business stakeholders to define objectives, gather requirements, and develop scalable data solutions aligned with **BI** and project goals using **Agile & Scrum methodology**.
- Migrated on-prem **SQL** Server environments to **Azure Cloud Azure Synapse Analytics, Azure SQL DB** ensuring seamless data transitions and adherence to data governance standards.
- Designed and developed complex data pipelines and transformations using **Azure Data Factory (ADF), PySpark**, and **Azure Databricks** for efficient data processing.
- Built and maintained streaming and batch data ingestion pipelines using **Kafka, Hadoop, HDFS**, and **Spark** for high-volume data processing.
- Created detailed source-to-target mapping documents and implemented **ETL** processes for consistent and accurate data movement across multiple layers.
- Developed **Power BI dashboards** and **visualizations** to deliver actionable insights and support business decision-making.
- Automated workflows, optimized performance, and provided 24/7 production support to ensure system stability and quick incident resolution.
- Led estimation efforts, reviewed complexities, and collaborated with cross-functional teams to continuously refine and improve data engineering solutions.

Environment: Azure, ADF, Databricks, Synapse Analytics, SQL DB, Hadoop, Spark, Kafka, HDFS, Python, PySpark, ETL, Power BI, Agile/Scrum.

IBM, India | Aug 2021 – Jun 2023
Data Engineer

- Designed and developed end-to-end **AWS** and **GCP**-based data solutions, leveraging services like **AWS S3, Redshift, Glue, EMR**, and **GCP DataProc, BigQuery**, and **Bigtable**.
- Built and optimized **ETL pipelines** using **AWS Glue, Apache Spark, Hive**, and **Python** to process structured, semi-structured, and unstructured data.
- Migrated and integrated on-prem Hadoop systems to Google Cloud Platform, ensuring optimized data storage, retrieval, and governance compliance.
- Implemented large-scale data processing workflows in **Amazon EMR** and **Hadoop clusters**, applying partitioning, bucketing, and query optimizations in Hive.
- Developed and automated data pipelines using **Python, Apache Beam**, and **Cloud Dataflow** for validation, transformation, and archival in Google Cloud Storage.
- Automated workflows with **Oozie** and **Cron Jobs**, streamlining Hive job scheduling, **UDF** execution, and other recurring tasks.
- Managed and processed data across a **200 PB data lake**, ensuring performance tuning, scalability, and data integrity across multiple data layers.
- Collaborated in an agile environment with **CI/CD pipelines, Kubernetes**, and **cross-functional** teams to deliver scalable and high-performance data engineering solutions.

Environment: AWS ,EMR, S3, Redshift, Lambda, Glue, GCP , DataProc, BigQuery, GCS, Dataflow, Bigtable, Hadoop 3.3, Spark 3.1, Python, Hive, Oozie, CI/CD, Apache Beam, REST API, Agile, HBase, Data Lake, Kubernetes.

EDUCATION:

Master's in Computer and Information Systems | **May 2025**
Saint Louis University