



Professional Summary

Dynamic **Azure Data Engineer** with 4+ years of experience in designing and optimizing **cloud-based data solutions** for healthcare and financial industries. Skilled in **Azure Data Factory, Synapse Analytics, Databricks, Snowflake**, and Power BI for **ETL workflows** and real-time data ingestion. Proficient in SQL, Python, Scala, and PySpark, with expertise in **data warehousing** and schema design. Strong in CI/CD pipelines using **Azure DevOps**, GitHub, and containerized workflows with Kubernetes and Docker. Experienced in data security, governance, and compliance using **Azure Purview** and Key Vault. Hands-on with **Snowflake query optimization** and machine learning integration for predictive analytics. Adept at Agile methodologies, ensuring, iterative and high-quality delivery.

Education

Master of Professional Studies in Data Science | University of Maryland Baltimore County | MD

GPA: 3.6

Bachelor of Technology in Computer Science and Engineering | Lovely Professional University | India

GPA: 3.5

Technical Skills

- **ETL & Data Integration:** Azure Data Factory, SSIS, Apache Airflow, Talend, Snowflake ETL Workflows, PolyBase
 - **Big Data & Cloud:** Azure Data Lake, Azure Synapse Analytics, Azure Databricks, Snowflake, Cosmos DB, Apache Spark
 - **Machine Learning & AI:** Azure Machine Learning, Databricks MLflow, Spark MLlib, Feature Engineering
 - **Databases:** SQL Server, Azure SQL DB, Snowflake, Oracle, MySQL, MongoDB, Cassandra
 - **Reporting:** Power BI, SSRS, Tableau
 - **Data Processing & Real-Time Streaming:** Azure Stream Analytics, Kafka, Polybase
 - **Data Modeling & Warehousing:** Star Schema, Snowflake Schema, Data Partitioning, Clustering, Time Travel
 - **Data Modelling & Analysis:** Erwin, Excel, SSAS, Alteryx
 - **Programming & Scripting:** SQL, T-SQL, U-SQL, PySpark, Python, Scala, Java, Shell Scripting
 - **DevOps & Orchestration:** Azure DevOps, GIT, Jenkins, Hive, Docker, Kubernetes
 - **Collaboration & Workflow Automation:** Jira, Confluence, Azure Logic Apps, Azure Functions
-

Certifications:

Azure Data Engineer Associate (DP-203) from Microsoft

AWS Solutions Architect Associate from Amazon (AWS)

Professional experience:

Kaiser Permanente

Azure Data Engineer

March 2023 – Present

- Designed, developed, and optimized **enterprise-grade ETL pipelines** using **Azure Data Factory (ADF), Snowflake**, and **Databricks**, ensuring highly efficient **data ingestion and transformation**, leading to a 50% improvement in data accessibility.
- Implemented **real-time data processing** and analytics solutions using **Azure Stream Analytics, Event Hubs**, and **Apache Spark on Databricks**, optimizing low-latency **data flows for business insights**.
- Architected, optimized, and administered **Snowflake environments**, leveraging clustering, **query tuning**, and **workload management** for efficient data warehousing.
- Developed and deployed machine learning pipelines using **Azure Machine Learning** and **Databricks MLflow**, supporting **predictive analytics and AI-driven insights**.
- Designed and implemented **data models (Star and Snowflake schemas)** to support analytics workflows, improving query performance by 30%.

- Designed and implemented **Azure Data Lake** Storage solutions, optimizing structured and **unstructured data storage** for scalable analytics and **archival purposes**.
- Built and maintained scalable **NoSQL and relational databases** using **MongoDB, Cosmos DB, and Azure SQL**, ensuring fault-tolerant and **high-throughput transactional processing**.
- Automated and orchestrated **data workflows** using **Apache Airflow** and **Azure Data Factory**, improving process automation and **job execution monitoring**.
- Containerized and deployed ETL workloads using **Docker and Kubernetes** on **Azure Kubernetes Service (AKS)**, enhancing scalability and **cross-platform compatibility**.
- Engineered Kafka-based **real-time streaming pipelines**, enabling scalable, event-driven architectures.
- Designed and optimized Snowflake **data models**, incorporating **zero-copy cloning, time travel, and Snowpipe** for real-time ingestion.
- Developed and deployed interactive BI **dashboards** in **Power BI** and **Tableau**, driving data-driven decision-making across multiple business units.
- Automated **Azure Data Factory orchestration**, reducing manual intervention by 40% and increasing **pipeline reliability**.
- Optimized **Azure Databricks clusters**, enhancing resource allocation and execution efficiency.
- Integrated **Azure Synapse** with **Snowflake** and **Power BI**, enabling faster analytics and **unified reporting**.
- Implemented CI/CD automation with **Azure DevOps**, accelerating pipeline deployments and reducing **software release cycles by 45%**.
- Established **enterprise-wide data governance and compliance strategies** using **Azure Purview** and **Key Vault**, ensuring security controls and metadata **management**.
- Spearheaded **performance tuning initiatives** across **Azure and Snowflake**, achieving query performance improvements and reducing **cloud storage costs**.

Environment: Azure Data Factory (ADF), Azure Stream Analytics, Apache Spark, Databricks, MongoDB, Apache Airflow, Docker, Kubernetes (AKS), Kafka, Snowflake, Tableau, Power BI, Azure Synapse, Azure DevOps, Azure Purview, Machine Learning

Access Healthcare, India

Azure Data Engineer

May 2020 - August 2022

- Analyzed **production applications** and built data solutions using **Azure PaaS services**, improving business intelligence visualization.
- Conducted **impact assessments**, ensuring minimal disruptions while optimizing **business processes**.
- Designed and implemented **ETL processes** with **Azure Data Factory**, transforming diverse data sources.
- Ingested and processed **large datasets** into **Azure Data Lake, Azure SQL Database, and Azure Synapse** using ADF and Databricks.
- Developed **REST APIs** for analytics data retrieval, enabling seamless **cross-platform integration**.
- Created **ADF pipelines** with Linked Services and Datasets, ensuring reliable data flows.
- Worked with **Talend ETL and Snowflake** cloud database systems, ensuring efficient data transformation.
- Optimized **Spark Databricks clusters**, reducing execution time by **25%**.
- Automated **ETL deployments** using JSON scripts, improving project efficiency by **40%**.
- Developed and managed **CI/CD pipelines** in **Azure DevOps and VSTS**, streamlining build and release processes.
- Implemented **Azure Logic Apps** for automated data ingestion workflows.
- Leveraged **Azure Functions** to build scalable **serverless compute solutions**.
- Designed **real-time data processing solutions** using **Azure Stream Analytics, Event Hub, and Service Bus Queue**.
- Enforced **data governance frameworks** with **Azure Purview**, enhancing compliance efforts.

Environment: Azure PaaS, **Azure Data Factory (ADF)**, Azure Data Lake, **Azure SQL Database**, Azure Synapse, **Databricks**, Apache Spark, PySpark, **REST APIs**, Blob Storage, JSON, **Azure DevOps**, VSTS, Azure Logic Apps, Azure Functions, Azure Stream Analytics, Event Hub, **Service Bus Queue**, Azure Purview.