

NAME : Krishna priya Yalamanchi

Phone number : 501-436-3590

Email Id : yalamanchikrishnapriya4@gmail.com

Senior Data Engineer Resume

Professional Summary:

- **Senior Data Engineer** with a **9+ Years of experience** in creating and improving data pipelines, **data lakes**, and **cloud platforms** for healthcare, finance, and retail.
- Experienced in data lifecycle management, including ingestion, **ETL/ELT**, **warehousing**, streaming, analytics, and **BI** for large companies.
- Skilled at using **AWS**, **Azure**, and **GCP** to build secure, available, and cost-effective data platforms.
- Built real-time streaming solutions with **Apache Kafka**, **Spark Streaming**, and **Flink**, which sped up processing for urgent applications.
- Experienced with **Snowflake**, **Redshift**, **Synapse**, and **Databricks**, creating cloud analytics environments for business teams.
- Developed **data lake** houses and governance to improve data quality, compliance, and **access** for analytics.
- Knowledgeable in **Python**, **SQL**, **PySpark**, and **Scala**, building **ETL** frameworks that cut development time by over 40%.
- Worked with data scientists and analysts to supply datasets, ensuring accurate models and faster **predictive** analytics.
- Skilled in using **Airflow** and **Azure Data Factory** to manage workflows, ensuring reliable data delivery.
- Successfully cut cloud infrastructure costs, saving companies over \$1M per year through workload balancing and resource management.
- Led data migration projects, moving systems to **AWS** and **Azure** while maintaining uptime.
- Worked with product owners to align data projects with revenue, compliance, and efficiency goals.
- Experienced in mentoring engineers, reviewing code, and promoting practices for **CI/CD** in data projects.
- Known for problem-solving skills, such as automating pipeline monitoring to reduce manual work by 70%.
- **Dedicated** to secure and compliant data solutions, with knowledge of **HIPAA**, **GDPR**, and financial regulations.

Technical Skills:

Frameworks / Libraries / Tools	Apache Spark, PySpark, Databricks, Apache Kafka, Flink, Airflow, Azure Data Factory, SSIS, SSRS, SSAS, Delta Lake, Collibra
Programming & Scripting	Python, SQL, Scala, Java, Shell Scripting, PL/SQL, T-SQL

Databases / Warehouses	Snowflake, Redshift, Azure Synapse, BigQuery, Oracle, SQL Server, Teradata, PostgreSQL, MySQL
Cloud Platforms	AWS (S3, Glue, EMR, Lambda, Redshift, Lake Formation), Azure (Data Lake, Synapse, Databricks, ADF), GCP (BigQuery, Dataflow, Pub/Sub)
DevOps & CI/CD	Git, GitHub Actions, Jenkins, Terraform, Docker, Kubernetes, Control-M
APIs & Protocols	REST, SOAP, GraphQL, HL7, FHIR, EDI, JSON, XML, Avro, Parquet
Visualization / BI Tools	Tableau, Power BI, Crystal Reports, QlikView

Work Experience:

Client: Molina Healthcare, WA

Mar 2025 –

Present

Role: Senior Data Engineer

Responsibilities:

- Created a **data Lakehouse** using **AWS S3**, **Glue**, and **Snowflake**, which boosted our analytics and cut query times by 55%.
- Developed fraud detection systems with **Kafka**, **Spark Streaming**, and **Flink**, managing over 2 million transactions a day with very quick response times.
- Automated **ETL** processes with **Airflow** and **Python**, which decreased manual data work by 40% and allowed for almost real-time reporting.
- Improved **Snowflake** setup by adjusting clustering, partitioning, and caching, leading to a 30% decrease in computing expenses and faster queries.
- Worked with business teams to set key metrics for financial risk models, making sure we had good data for machine learning.
- Headed a data governance project, adding **AWS Glue Data Catalog** and **Collibra** to handle metadata and track data origins.
- Analyzed **AWS** costs and found ways to save \$600,000 a year on storage and computing.
- Moved old **Teradata** and **Oracle** systems to **Snowflake**, updating over 50TB of data while keeping 99.9% uptime.
- Put in place role-based security in **Snowflake** and **AWS Lake Formation**, which helped meet **SOX** and **GDPR** standards.
- Worked with data scientists to put ML models into action, boosting prediction accuracy by 18%.
- Set up **CI/CD** workflows with **GitHub Actions** and **Terraform**, automating how we set up our infrastructure and cutting down setup time by 80%.
- Built a data quality system with **PyDeequ** that automatically checks data, reducing bad data loads by 25%.

- Promoted teamwork with **BI** teams, speeding up self-service analytics with **Power BI** and **Tableau**.

Environment: AWS (S3, Glue, EMR, Lambda, Redshift, Lake Formation), Snowflake, Apache Kafka, Spark, Flink, Python, SQL, Airflow, GitHub Actions, Terraform, Tableau, Collibra.

Client: American Equity, Des Moines, IA
2025

Nov 2022 – Feb

Role: Data Engineer

Responsibilities:

- Designed a **HIPAA-compliant data lake** on **Azure Data Lake** and **Databricks**, putting clinical, claims, and **EHR** data in one spot for company-wide analysis.
- Created **PySpark ETL pipelines** that process over 10TB of patient data each day, cutting down data intake times by 60%.
- Set up **Azure Data Factory pipelines** that integrate more than 40 different data sources, which improved data sharing and made it easier to get data ready for analysis.
- Worked with compliance teams to make sure PHI masking and encryption rules were followed, keeping us in line with **HIPAA** rules.
- Came up with data mart solutions in **Azure Synapse**, which allowed for real-time healthcare use dashboards.
- Added **FHIR** and **HL7** standards into the data intake **pipelines** to ensure data works well across healthcare systems.
- Made automated pipeline monitoring dashboards in **Power BI**, dropping the time it takes to spot downtime from hours to minutes.
- Improved **Databricks** clusters, which cut costs and reduced runtime by 45%.
- Used **Delta Lake** to put in place data versioning and lineage, making sure clinical research studies can be reproduced.
- Worked with doctors and researchers to give them specific datasets for AI models that stratify patient risk.
- Moved old **SQL Server** and **SSIS** tasks to the **Azure** cloud, which lowered infrastructure maintenance costs by 50%.
- Pushed for delivery methods, heading up sprint planning, backlog improvements, and daily scrums for a data team of 7.
- Built a data quality and anomaly detection setup that flags mistakes in **medical** claims with 90% accuracy.
- Managed data sharing deals with outside partners, exchanging data safely with regulatory groups and research centers.

- Helped out on the organizational data governance committee, writing standards for data cataloging and lifecycle management.

Environment: Azure Data Factory, Databricks, Azure Synapse, SQL Server, Delta Lake, HL7, FHIR, Power BI, Python, PySpark, Git, Agile, HIPAA-compliant frameworks.

Client: BNY Mellon, NYC, NY
2022

Apr 2021 – Oct

Role: Data Engineer

Responsibilities:

- Built a data warehouse using **AWS Redshift**, bringing together data from point-of-sale, supply chain, and **CRM systems** for a full view of our customers.
- Designed dashboards for real-time sales using **Kinesis** and **Lambda**, cutting reporting delays from 12 hours to under 5 minutes.
- Made **ETL pipelines** with **Python** and **Glue**, which automated data collection from over 25 retail systems and cut manual data preparation by 70%.
- Used **Spark** and **SQL** to build customer segmentation models, boosting the return on targeted campaigns by 22%.
- Worked with inventory teams to create **predictive restocking** pipelines, which lowered stock-outs by 15%.
- Moved **Oracle** retail databases to **Redshift**, speeding up reporting and lowering licensing costs.
- Developed self-service **BI** data marts for marketing and finance, making decision-making quicker across departments.
- Improved our data quality monitoring, which caught errors before they affected analytics systems.
- Worked with various teams to design **ETL** workflow orchestration with **Airflow**, making things more reliable.
- Improved **Redshift** queries and schema design, which cut reporting time from 20 minutes to under 2 minutes.
- Trained analysts on **SQL** and **BI** best practices, which helped people use self-service tools more.
- Automated an **S3**-based archiving system, cutting storage costs by using lifecycle policies.
- Worked with vendors to set up secure API integrations for e-commerce data collection.
- Improved fraud detection by adding third-party APIs, boosting detection accuracy by 30%.
- Played a role in data governance efforts, standardizing data definitions across the retail setup.

Environment: AWS Redshift, S3, Glue, Lambda, Kinesis, Airflow, Python, SQL, Oracle, Tableau, Power BI.

Client: Empower Retirement, Greenwood, CO
2021

Jan 2019 – Mar

Role: Junior Data Engineer

Responsibilities:

- Created and kept up **ETL pipelines** using **SQL Server** and **SSIS**. I brought together shipment tracking, warehouse, and **ERP** data into one place.
- Made daily batch **pipelines** that automatically matched up freight costs, cutting down hand work by 60%.
- Worked with **BI** analysts to come up with reporting cubes in **SQL Server** Analysis Services (**SSAS**), which helped create shipment tracking dashboards that updated almost in real-time.
- Made **SQL** queries and indexes run faster, which brought down warehouse reporting times from 30 minutes to 5 minutes.
- Added data validation checks when data was coming in, which lowered data quality problems down the line by 40%.
- Worked with operations managers to design models that **predict** fleet use, boosting delivery success by 12%.
- Made stored procedures and functions to make business logic the same across different systems.
- Worked with vendors to set up **EDI**-based data integrations, making sure partner logistics data was sent without problems.
- Helped the money teams with one-off data requests, which helped make sure we were following **SOX** reporting rules.
- Made automated alerting and monitoring scripts to let stakeholders know when shipments were late or data **pipelines** failed.
- Moved old Excel-based reporting to **SQL Server** Reporting Services (**SSRS**), which lowered mistakes from hand reporting.
- Made a data archival framework to move old data to cheaper storage, which lowered costs by 25%.
- Worked with **QA** teams to do full pipeline testing, which made sure we were meeting SLA standards.
- Took part in **Agile** events, helping with sprint planning, backlog grooming, and retrospectives.

Environment: SQL Server, SSIS, SSRS, SSAS, T-SQL, EDI, Excel VBA, ERP/Logistics Systems, Windows Server, Agile.

Client: KLA Tencor, India
2018

Nov 2014 – Aug

Role: Data Analyst / Engineer

Responsibilities:

- Created and kept up **ETL pipelines** using **SQL Server** and **SSIS**. I brought together shipment tracking, warehouse, and **ERP** data into one place.

- Made daily batch **pipelines** that automatically matched up freight costs, cutting down hand work by 60%.
- Worked with **BI** analysts to come up with reporting cubes in **SQL Server** Analysis Services (**SSAS**), which helped create shipment tracking dashboards that updated almost in real-time.
- Made **SQL** queries and indexes run faster, which brought down warehouse reporting times from 30 minutes to 5 minutes.
- Added data validation checks when data was coming in, which lowered data quality problems down the line by 40%.
- Worked with operations managers to design models that **predict** fleet use, boosting delivery success by 12%.
- Made stored procedures and functions to make business logic the same across different systems.
- Worked with vendors to set up **EDI**-based data integrations, making sure partner logistics data was sent without problems.
- Helped the money teams with one-off data requests, which helped make sure we were following **SOX** reporting rules.
- Made automated alerting and monitoring scripts to let stakeholders know when shipments were late or data **pipelines** failed.
- Moved old Excel-based reporting to **SQL Server** Reporting Services (**SSRS**), which lowered mistakes from hand reporting.
- Made a data archival framework to move old data to cheaper storage, which lowered costs by 25%.
- Worked with **QA** teams to do full pipeline testing, which made sure we were meeting SLA standards.
- Took part in **Agile** events, helping with sprint planning, backlog grooming, and retrospectives.

Environment: Oracle, SQL Server, PL/SQL, Python, Shell Scripting, Tableau, Crystal Reports, Control-M, Access, ERP Systems.