

RISHABH DINESH SINGH

New York City, NY | rishabhdineshsingh@gmail.com | (857) 693-8587 | [linkedin.com/in/rishabhdineshsingh/](https://www.linkedin.com/in/rishabhdineshsingh/) | [Portfolio](#)

SUMMARY

Passionate Data Engineer with 3 years of experience in designing and orchestrating scalable **ETL pipelines** and cloud applications. Proficient in **Python**, **SQL**, **AWS**, and **Apache Airflow**, with a strong background in data modeling and warehousing. Committed to driving data-driven decision-making through optimized data architecture and automation.

PROFESSIONAL EXPERIENCE

Data Engineer | Aretove Inc | New York City, NY

Sep 2024 – Present

- Engineered a data pipeline using **Airbyte** and **Python** to extract **Salesforce** and Student Information System data, achieving 98% data accuracy in the **Snowflake** Data Warehouse
- Architected dimensional models from raw student data in **Snowflake** by utilizing **DBT**, enhancing data reusability by 25%
- Developed reusable **DBT macros** for data transformation in Snowflake, reducing development time by 40%
- Streamlined project management using **JIRA** to track tasks and sprints, enhancing team productivity by 15%
- Implemented dynamic task dependencies in **Airflow** DAGs to optimize resource utilization, improving pipeline throughput by 30%
- Designed **Power BI** dashboards to visualize student performance metrics, improving decision-making for administrators by 20%

AI Data Engineer Co-op | Schell Brothers | Boston, MA

May 2023 – Dec 2023

- Engineered a Gen-AI chatbot leveraging **Python** and LangChain, automating customer support processes and reducing workload by 75%, while achieving 99% accuracy in data pipeline extraction across diverse web page sources
- Streamlined customer inquiry classification using logistic regression and **NLP**, achieving 90% accuracy and 30% faster processing
- Developed ML Ops pipelines using **Airflow** to retrain models with historical data, boosting accuracy by 0.4% with bi-weekly updates
- Designed AI-powered SQL optimization system using LLAMA Index, **AWS Redshift** and NLP, reducing query response time by 33%

Data Engineer | Accenture | Mumbai, India

Sep 2019 – Jul 2022

- Implemented **SQL**-based data modeling techniques to optimize batch data transformation, enhancing cross-functional data accessibility and integrity by 30%
- Optimized **AWS Glue** ETL jobs with parallel processing, reducing execution time by 40% and enhanced data pipeline performance
- Automated data quality checks in DBT to validate transformations, ensuring compliance with documented data governance policies.
- Collaborated across multiple teams within an **Agile** framework, utilizing CI/CD processes and **Jenkins** to improve workflows, resulting in 20% higher stakeholder satisfaction
- Consolidated diverse data sources into a centralized **warehouse**, streamlining pipelines to cut retrieval times by 40%
- Improved the **Collibra** Catalog to enhance **data governance**, reducing table metadata search time by 25%

EDUCATION

Master of Science in Information Systems | Northeastern University – Boston, MA

Sep 2022 – Aug 2024

Bachelor's in Engineering – Information Technology | University of Mumbai – Mumbai, India

Aug 2015 – May 2019

ACADEMIC PROJECTS

NYT Sentiment Pulse: Real-Time News Analysis & Insights | *Microsoft Fabric, Microsoft PowerBI*

Jan 2024 - May 2024

- Engineered data pipeline using **Data Factory**, transforming NYT data into **One Lakehouse**, reducing data retrieval time by 40%
- Developed **PySpark** sentiment analysis model, enhancing data insights with 95% accuracy in real-time sentiment tracking
- Created semantic model in **Power BI**, enabling interactive reports that improved trend analysis efficiency by 25%
- Implemented **Data Activator** alerts, increasing responsiveness to critical news trends by 30%

AlssueFlow | *FastAPI, Docker, Airflow, Snowflake*

Sep 2023 - Dec 2023

- Devised a data pipeline updating **Milvus** embedding database post-user searches, ensuring 100% coverage of unscreened issues
- Revolutionized data storage in **Snowflake**, accelerating retrieval efficiency by 30% and elevating overall system performance
- Spearheaded **Docker** containerization of AlssueFlow app, achieving 99.9% uptime with daily Airflow DAG updates.

SKILLS

Database: Oracle, PostgreSQL, MySQL, MongoDB, SQL Server, AWS DynamoDB, AWS Redshift

Programming Languages: Python, SQL, R, Spark, Scala, Java, Shell Script

Cloud Technologies: Amazon Web Services, Snowflake, Docker, AWS Lambda, Microsoft Azure, Microsoft Fabric

Data Engineering: Data Build Tool, Hadoop, Informatica, Ab Initio, AWS GLUE, Jenkins, JIRA, Airbyte

Automation: Apache Airflow, Dagster, Jenkins, Terraform, Docker, Kubernetes, AutoSys Script, Kafka