

sSai Kiran Dangeti

saikiran.dangeti14@gmail.com | +1 (269) 267-2927

<https://www.linkedin.com/in/sai-kiran-dangeti/>



SUMMARY

- Over 5 years of experience as a Data Engineer, specializing in designing, developing, and implementing scalable data architectures and data pipelines for high-performance systems across AWS, Azure, and GCP environments
- Extensive hands-on experience with cloud data warehousing technologies, including Snowflake, Redshift, BigQuery, and Azure Synapse, supporting large-scale data storage and analytics solutions.
- Proficient in building and optimizing ETL/ELT pipelines using Apache Kafka, Apache Spark, and Apache Airflow to ensure efficient data integration, transformation, and real-time streaming for business-critical applications.
- Expertise in handling and processing structured, semi-structured, and unstructured data using tools such as Spark, Hive, HBase, Flume, and Sqoop in big data environments.
- Strong experience working with major cloud platforms including AWS (EMR, Glue, S3, Lambda), Azure (Azure Data Lake, Azure Data Factory, Azure SQL), and GCP (BigQuery, Cloud Functions, Pub/Sub).
- Advanced skills in Python (PySpark, Pandas, NumPy), SQL (T-SQL, PL/SQL), and shell scripting for data automation and tool development, as well as managing and optimizing large-scale data operations.
- Expertise in data governance, quality checks, and lineage using industry-standard frameworks and tools, ensuring data integrity and compliance across data workflows.
- Experienced in building and deploying data pipelines in real-time and batch processing scenarios, leveraging event-driven architectures using Apache Kafka and AWS Kinesis for event stream processing.
- Hands-on experience with cloud-native DevOps practices, CI/CD pipeline automation using Git, Jenkins, Azure DevOps, and Docker for continuous integration and deployment of data solutions.
- In-depth knowledge of distributed computing, data partitioning, and optimization techniques in Hadoop ecosystems (HDFS, MapReduce, Hive) and Spark environments, ensuring efficient data processing and job scheduling.
- Skilled in the design of data models, data marts, and dimensional models (Star Schema, Snowflake Schema) for BI solutions and reporting, using both relational and NoSQL databases (MongoDB, Cassandra, Teradata).
- Experience in managing and orchestrating data processing workflows using scheduling tools such as Airflow, Oozie, Cron, and Control-M to ensure timely execution of ETL jobs and batch processes.
- Expertise in integrating various data sources and APIs, including third-party systems and web services, for seamless data ingestion and synchronization across enterprise systems.
- Proficient in designing and deploying machine learning pipelines and feature engineering workflows, ensuring data is appropriately pre-processed for ML and predictive analytics models.
- Solid understanding of data privacy, security, and compliance regulations, with the ability to implement best practices for secure data handling in cloud and on-prem environments.
- Familiarity with BI tools (Tableau, Power BI, Looker) for developing dashboards and data visualization solutions that empower business users to make data-driven decisions.
- Strong communicator and collaborator, skilled in working with cross-functional teams including data scientists, business analysts, and stakeholders to align data solutions with business goals and strategies.
- Continuously learning and staying up to date with emerging technologies in data engineering, cloud computing, and AI, ensuring that solutions remain innovative and efficient.

TECHNICAL SKILLS

Big Data Ecosystem	HDFS, Yarn, MapReduce, Spark, Kafka, Kafka Connect, Hive, Airflow, Stream Sets, Sqoop, HBase, Flume, Pig, Ambari, Oozie, Zookeeper, Nifi, Sentry
Cloud Environment	Amazon Web Services (AWS), Microsoft Azure
Databases	MySQL, Oracle, Teradata, MS SQL SERVER, PostgreSQL, DB2
AWS	EC2, EMR, S3, Redshift, EMR, Lambda, Kinesis Glue, Data Pipeline
Microsoft Azure	Databricks, Data Lake, Blob Storage, Azure Data Factory, SQL Database, SQL Data Warehouse, Cosmos DB, Azure Active Directory, Azure Synapse Analytics
DevOps and Deployment	Git, Docker, Kubernetes, CI/CD pipeline, Terraform, MLFlow

Programming Languages	Python (Pandas, SciPy, NumPy, Scikit-Learn, Stats Models, Matplotlib, Plotly, Seaborn, Keras, TensorFlow, PyTorch), PySpark, T-SQL/SQL, PL/SQL, HiveQL, Scala, UNIX Shell Scripting
Tools	Tableau, PowerBI, Excel, Jira, Informatica, MS Power Query, Google Data Studio
Soft Skills	Communication, Problem Solving, Teamwork, Time Management

WORK EXPERIENCE

Zoetis | Kalamazoo, MI

Sr. Data Engineer

December 2023 – Present

- Designed and developed scalable and cost-effective architecture in AWS Big Data services for data collection, ingestion, storage, processing, and visualization.
- Developed PySpark Data Ingestion framework to load claims data into HIVE tables with data cleansing, aggregations, and deduplication.
- Created end-to-end data pipelines in distributed environments using Big Data tools, Spark, and Tableau for visualization.
- Designed and maintained Alteryx workflows to automate ETL processes and streamline data integration.
- Developed CloudFormation Templates (CFTs) for infrastructure migration across environments.
- Leveraged Spark features like in-memory processing, broadcast variables, and efficient joins to minimize latency.
- Wrote Python topology scripts to generate CloudFormation templates for EMR cluster creation.
- Utilized AWS services including Athena, Redshift, Glue, EMR, Lambda, S3, and Step Functions.
- Built Spark ETL jobs handling 450 GB of data daily and integrated them with AWS Glue.
- Structured S3 buckets and data lakes to optimize Glue crawler efficiency and metadata management.
- Used Glue Data Catalog and Lake Formation for schema validation and governance.
- Developed Spark applications in Scala and Python for processing data from RDBMS and streaming sources.
- Employed Spark SQL and PySpark for batch and streaming ETL jobs.
- Applied performance optimization techniques such as partitions, in-memory caching, and effective joins.
- Wrote unit test cases for Spark code as part of CI/CD process using Bitbucket and Bamboo.
- Built Databricks Spark jobs for cleansing, validation, standardization, and transformation.
- Migrated data from on-premises systems to AWS S3 and developed Python scripts for API-based data extraction.
- Ingested and transformed data using AWS Lambda, Glue, Step Functions, and custom Python scripts.
- Wrote YAML files for each data source and automated Glue table stack creation.
- Developed Lambda functions with IAM roles and triggers including SQS, EventBridge, and SNS.
- Automated Lambda deployments and integrated with S3 events.
- Wrote UNIX shell scripts and scheduled cron jobs for process automation.
- Built Informatica mappings with expressions, filters, joins, and lookups for data consistency.
- Implemented Sqoop for incremental DB2 to Hive loads; connected to Tableau for reporting.
- Created Spark applications for multi-format data transformation using PySpark and Spark-SQL.
- Used Spark Streaming with Kafka to load real-time data into HDFS, HBase, and Cassandra.
- Collected and transformed near-real-time S3 data for building learner models in HDFS.
- Used Apache NiFi to copy local files into HDP.
- Performed data modeling using Star and Snowflake schemas with Erwin.
- Automated data loading into HDFS using Oozie workflows.

Environment: AWS (EMR, S3, Glue, Lambda, Athena, Redshift, Step Functions, Lake Formation), PySpark, Python, Scala, Hive, Kafka, HDFS, Informatica, Tableau, Alteryx, Apache NiFi, Sqoop, Spark Streaming, Teradata, JIRA, Bitbucket, Bamboo, Linux, Oracle PL/SQL

- Collaborated with global and cross-functional teams to translate business requirements into scalable ELT solutions aligned with best architectural practices.
- Designed, built, and maintained robust ELT pipelines using Azure Data Factory, Azure Synapse Analytics, and SQL Server, ensuring seamless data integration from various source systems.
- Executed data extraction, transformation, and loading (ETL/ELT) strategies for both structured and semi-structured data, supporting enterprise Data Warehouse and Lakehouse initiatives.
- Orchestrated and scheduled ADF pipelines, datasets, and data flows to automate data ingestion into Azure Data Lake, Blob Storage, and Azure SQL Databases.
- Supported data validation, profiling, and cleansing processes to ensure high-quality and accurate data assets.
- Built scalable and reusable data transformation pipelines using DBT, SQL, and PySpark, enabling faster development and promoting repeatable workflows.
- Participated in Agile ceremonies and adopted iterative development practices for continuous delivery and improvement.
- Managed code repositories, version control, and CI/CD pipelines for Azure Data Factory and data workflows using Azure DevOps.
- Collaborated with system owners and IT teams to coordinate integration points, access controls, and secure data movement.
- Contributed to maintaining and supporting Allegion-style global data assets and documentation, including data lineage and cataloging.
- Applied data modeling best practices, developed star and snowflake schemas to support reporting and analytics.
- Used Azure Key Vault for secure credential management and implemented data security standards for enterprise compliance.
- Ensured performance optimization of data pipelines and cost-efficiency of Azure resource utilization.
- Demonstrated flexibility in adapting to changing project scopes and supported both development and operational data tasks.

Environment: Azure Data Factory, Azure Synapse Analytics, Azure SQL Database, SQL Server, Azure Data Lake, Blob Storage, DBT, Snowflake, PySpark, Python, Git, Azure DevOps, Agile/Iterative Development

Aspire Systems | Chennai, India**Associate Data Analyst****June 2020 – December 2021**

- Developed and maintained ETL workflows using Informatica PowerCenter 10.x, ensuring seamless data extraction, transformation, and loading across multiple source systems.
- Designed, developed, and optimized complex ETL mappings, mapplets, sessions, and workflows to automate and streamline data movement across diverse systems.
- Wrote advanced SQL queries, stored procedures, views, and triggers using Oracle and DB2, optimizing performance for large data sets and ensuring efficient data retrieval.
- Collaborated with cross-functional teams to design and implement ETL processes to extract data from various data sources such as Salesforce, Mainframe, Oracle, DB2, and flat files.
- Implemented performance tuning techniques for ETL workflows, enhancing data processing speeds and reducing resource utilization across the data pipeline.
- Contributed to the design and development of Snowflake schema and Star Schema modeling for efficient data storage and retrieval, ensuring high-quality data for reporting and analytics.
- Built and maintained interactive dashboards and visualizations using Tableau, providing actionable insights and enabling business users to perform deep data analysis.
- Participated in Agile ceremonies such as sprint planning, stand-ups, and retrospectives, ensuring smooth project execution and alignment with business objectives.

- Worked with SQL Server Integration Services (SSIS) for ETL processes, leveraging SSRS and SSAS for reporting and analytical solutions to meet business intelligence needs.
- Collaborated with team members to implement best practices in ETL development, improving code maintainability, and reducing technical debt.
- Worked closely with the team to transition to cloud-based data solutions, exploring tools like Azure Data Factory and AWS Glue for future scalability and modernization.

Environment: Informatica PowerCenter, Tableau, SQL, Oracle, DB2, Control-M, T-SQL, SSIS, SSRS, SSAS, Salesforce, Mainframe.

Associate Software Engineer

June 2019 – June 2020

- Collaborated with cross-functional teams including developers, analysts, and business stakeholders to gather and analyze requirements, ensuring alignment with business needs.
- Designed and developed ETL processes using Informatica, implementing complex transformations such as Source Qualifier, Expression, Aggregator, Lookup, Filter, Router, Joiner, and Update Strategy to ensure seamless data flow.
- Created and deployed Informatica jobs across environments, ensuring smooth deployment and troubleshooting, minimizing downtime during data migration.
- Developed analytical reports using Tableau, blending data from multiple sources and providing key insights for business decision-making.
- Built and maintained SSIS packages for ETL processes, automating data extraction, transformation, and loading.
- Participated in data modeling and created dimensional models using Star and Snowflake schemas, optimizing data for business intelligence and reporting.

Environment: Informatica PowerCenter, SQL, Oracle, DB2, Tableau, SSIS, SSRS, SSAS, Control-M

EDUCATION

Master of Science - Data Science

January 2022 - December 2023

Western Michigan University, Kalamazoo, MI

Courses: A.I., Machine learning, Applied linear models, Data mining, Database systems, and Advanced R.