

Summary

Data Engineer with 3+ years of experience designing cloud-native, AI-integrated data solutions across AWS, Azure, and GCP. Specialized in building scalable data pipelines, real-time inference workflows, and deploying LLMs for enterprise applications. Skilled in MLOps, vector databases, and automated ML pipelines using Spark, Databricks, and OpenAI APIs.

Skills

- Programming:** C, Java, Python, R, Scala, SQL, PySpark
 - Big Data & ETL Tools:** Apache Spark, Hadoop, Hive, Sqoop, Flume, Apache Kafka, NiFi, Talend, Informatica, DataStage, Apache Beam, Flink, Druid
 - AutoML Platforms:** Custom AutoML development using PySpark, Airflow, Python; automated model training and tuning.
 - Cloud Technologies:** AWS (S3, Redshift, Glue, Lambda, EMR), Azure (Databricks, Data Factory, Synapse, ADLS, Cosmos DB, Cognitive Services, OpenAI), Snowflake, GCP Big Query, SAP HANA
 - Databases:** PostgreSQL, MySQL, MongoDB, Cloud SQL (GCP), Apache Spark, Vector Databases
 - Machine Learning:** Logistic Regression, Decision Trees, Random Forest, PyTorch, AWS SageMaker
- AI/ML Frameworks:** scikit-learn, TensorFlow (basic), MLflow (for tracking), PyTorch (basic exposure) Fast AI, OpenAI API, Large Language Models (LLMs), Reasoning Models
 - BI & Visualization:** Tableau, Power BI, Excel
 - DevOps & IaC:** Docker, Kubernetes, Terraform, Jenkins, Git
 - Statistical Analysis:** Linear Regression, ANOVA, Chi-Square, SciPy, Seaborn, Scikit-learn, TensorFlow
 - Project Management:** Agile, Waterfall.
 - Tools:** Tableau, Power BI, Advanced Excel, Visual Studio, Jupyter Notebook, Matplotlib, Seaborn, CI/CD Pipelines, Orchestration Tools, Data Engineering Frameworks

Experience

Data Engineer
Accenture

06/2024 to Present
USA

- Designed and deployed cloud-native data pipelines using Azure Synapse, AWS Redshift, and Informatica, enabling scalable, high-performance analytics for enterprise AI use cases.
- Engineered scalable data pipelines and vector databases to support Fast AI and OpenAI-powered LLM applications, integrating reasoning models for enterprise-grade analytics solutions; optimized inference workflows on Azure Kubernetes Service (AKS) and Azure ML, reducing model latency by 35%.
- Led end-to-end deployment of LLM-based data intelligence platforms on Azure, leveraging Azure Data Factory, Azure Synapse, and OpenAI services to build secure, compliant, and high-availability solutions for multi-modal reasoning across structured and unstructured data.
- Built and operationalized AI-ready datasets by integrating feature engineering and model outputs into automated PySpark workflows, improving ML model training efficiency by 35%.
- Implemented MLOps workflows with Azure ML and GitHub Actions, supporting model versioning, automated retraining, and continuous monitoring of production models.
- Integrated real-time inference pipelines for fraud detection and customer personalization using Spark Structured Streaming and SageMaker endpoints, reducing decision latency by 40%.

Data Engineer
Cognizant Technology Solutions

06/2022 to 06/2023
TN, India

- Developed and tested ETL mappings and workflows using Informatica PowerCenter, improving data processing efficiency by 25% and ensuring scalable data integration across systems.
- Architected Azure-native data solutions using Azure Databricks, Azure Cognitive Services, and Azure OpenAI, enabling secure model deployment and orchestration for NLP and reasoning models across large-scale data lakes, ensuring SLA compliance and robust governance.
- Engineered and optimized data pipelines using Azure Data Factory, AWS Glue, and Spark Streaming, reducing latency in high-frequency data environments and enhancing data reliability.
- Wrote and optimized advanced PySpark scripts for large-scale data transformation tasks, significantly improving execution time and processing accuracy across structured and unstructured datasets.
- Designed and deployed cloud-based data solutions across AWS (EC2, S3, Lambda, Redshift) and Azure (Functions, Synapse), increasing system performance and supporting high-volume analytics workloads.
- Scaled and automated an AutoML platform using Spark, Airflow, and Python, reducing model training time for big data by 30% and enabling predictive insights for operational teams.
- Built and deployed a production-grade SaaS AI platform (Mate verse) integrating feature engineering, regression, and time-series modules with Flask, Go, and Python, enabling cross-industry adoption.

Associate Data Engineer
BharatPe

01/2021 to 05/2022
Remote, India

- Led the migration of large-scale datasets from legacy on-premise systems to cloud platforms (AWS and Azure), improving data availability and scalability.
- Designed and deployed robust ETL pipelines using PySpark, Spark SQL, and Databricks to cleanse, validate, and transform high-volume datasets.
- Built and orchestrated DAGs in Apache Airflow for workflow automation, optimizing incremental data loads and reducing latency in real-time processing.
- Tuned Spark jobs by leveraging in-memory processing, broadcast joins, and custom aggregations to enhance performance and reduce job execution time by up to 40%.
- Engineered scalable data processing frameworks to handle JSON, XML, Avro, ORC, and Parquet file formats, ensuring optimal storage and query performance.
- Created and maintained Hive tables for batch processing and analytical reporting, ensuring data quality and integrity across all environments.
- Implemented data ingestion solutions with Spark, Hive, and Sqoop, facilitating seamless data flow from relational and NoSQL sources.
- Managed Databricks clusters for both batch and streaming workloads, optimizing resource usage and improving operational efficiency.
- Collaborated with cross-functional teams to define data quality metrics, monitor pipeline health, and ensure compliance with data governance standards.

Certifications

- Microsoft Certified: Fabric Data Engineer Associate – Microsoft (2024)
- Azure Data Factory – LinkedIn Learning (2023)
- Data Science Foundation – Great Learning (2022)
- Java Programming – Great Learning (2021)
- Analyzing and Visualizing Data with Microsoft Excel – Great Learning (2021)
- Machine Learning with Python and R – Robokalam (2021)

Education

Master's: Advanced Data Analytics
University of North Texas

2025
United States