📄 SAI TEJA ANKURU

📍 Boston, MA | 📞 (716) 247-3810 | 📧 saiteja.ankuru@gmail.com

🧠 PROFESSIONAL SUMMARY

Versatile AI/ML Engineer with 5+ years of progressive experience architecting and deploying end-to-end AI systems across logistics, energy, and healthcare. Transitioned from data engineering to full-cycle machine learning and NLP, delivering intelligent solutions for real-time forecasting, route optimization, and decision support. Strong track record in scaling ML pipelines, integrating APIs, and leading multi-cloud deployments.

- Spearheaded production-grade ML and AI integrations that drove measurable business gains.
- Built scalable data platforms on Azure, AWS, and Databricks with Spark and PySpark.
- Architected deep learning pipelines using TensorFlow, HuggingFace, and LLM frameworks.
- Implemented NLP systems for clinical data and logistics decision-making.
- Delivered MLOps-driven workflows with containerization, CI/CD, and automated retraining.
- Designed visual analytics solutions using Tableau and Power BI for executive insights.
- Reduced operational inefficiencies through applied AI across logistics and energy sectors.
- Actively contributed to cross-functional teams and stakeholder communication.
- Consistently delivered projects with business alignment, agility, and innovation.
- Quick learner with deep curiosity for GenAI, multi-modal learning, and autonomous agents.

🛠️ TECHNICAL SKILLS

Languages: Python, SQL, R, Java, Bash
Machine Learning & AI: Scikit-learn, TensorFlow, PyTorch, Keras, XGBoost, LightGBM, CatBoost
NLP/LLMs: HuggingFace, spaCy, NLTK, LangChain, BERT, GPT, sentence-transformers
Data Engineering: PySpark, Apache Spark, Hadoop, Kafka, Delta Lake, Apache Airflow
Cloud Platforms: Azure (ADF, Synapse, Data Lake), AWS (S3, SageMaker, Lambda), GCP (BigQuery)
Visualization & BI: Tableau, Power BI, Plotly, Looker, Seaborn, Matplotlib
MLOps & DevOps: Docker, Kubernetes, MLflow, Flask, GitHub Actions, Jenkins, FastAPI
Big Data Tools: Databricks, Hive, HDFS
Databases: PostgreSQL, MySQL, SQL Server, MongoDB
Tools & Platforms: VS Code, Jupyter, Git, Confluence, Trello, REST APIs, Postman

---

💼 PROFESSIONAL EXPERIENCE

AI/ML Data Scientist

Leapgen | Remote, USA
Nov 2024 – Present

- Spearheaded the development of an energy usage prediction platform using LightGBM and Flask.
- Designed and deployed an end-to-end NLP pipeline using HuggingFace Transformers and spaCy for extracting energy-related insights from customer reports and sensor logs.

- Deployed containerized ML models using Docker, significantly streamlining the release process and reducing the need for manual intervention.
- Engineered a hybrid inference pipeline (batch + real-time) using Apache Airflow.
- Delivered stakeholder-facing web application that enabled energy forecasting on-demand.
- Achieved near real-time predictions with minimal latency, influencing client investment strategy.
- Conducted advanced EDA and feature selection from 1M+ time-series records.
- Integrated CI/CD workflows using GitHub Actions for end-to-end pipeline automation.
- Mentored peers on MLOps best practices and Dockerized deployments.
- Reduced model error through targeted feature engineering and cross-validation.
- Presented insights and architecture to non-technical stakeholders and secured buy-in for scale-up.

Data Scientist Intern – AI Track

MSR Cosmos | Washington, USA
April 2023 – Nov 2024

- Designed 5+ enterprise-grade ETL pipelines using PySpark, Azure Data Factory, and Delta Lake.
- Streamlined ingestion of 100+ GB daily data into Azure Data Lake for machine learning usage.
- Enabled rapid prototyping and model iteration in Databricks notebooks.
- Built fault-tolerant DAGs using Apache Airflow for mission-critical batch jobs.
- Developed automated test harnesses for pipeline validation and data quality monitoring.
- Partnered with AI team to produce curated datasets for classification and regression use cases.
- Orchestrated modular workflows to support reuse across departments.
- Implemented row-level security and governance policies via Synapse Analytics.
- Contributed architectural insights in design sprints, influencing data mesh strategy.
- Received internal recognition for improving latency in critical ETL paths.

Associate Data Scientist

Wavelabs Technologies | Hyderabad, India
Aug 2021 – Dec 2022

- Led ML model development for route optimization, improving delivery speed by hours across routes.
- Integrated external APIs (traffic, weather) for dynamic routing intelligence.
- Developed forecasting models using Gradient Boosting and time-series techniques.
- Built dashboards in Tableau to monitor transport cost reduction and performance KPIs.
- Conducted A/B testing on delivery logic to validate model impact on field ops.
- Partnered with DevOps team to build a CI/CD pipeline for model refresh and release.
- Automated retraining jobs based on new logistics data streams.

- Designed internal tools to analyze warehouse-to-destination path efficiency.
- Mentored two junior engineers and created onboarding playbooks.
- Drove adoption of predictive analytics across supply chain teams and stakeholders.

Data Scientist Intern

Wavelabs Technologies | Hyderabad, India
Jan 2020 – Aug 2021

- Contributed to ingestion pipelines processing logistics, sensor, and event data.
- Supported Spark-based workflows for processing large-scale delivery logs.
- Built early-stage data validation and deduplication scripts in Python.
- Developed Power BI dashboards to visualize route delays and exception reports.
- Documented data flow diagrams and maintained technical wikis.
- Shadowed senior data engineers on large-scale HDFS deployments.
- Gained hands-on exposure to Kafka streaming and Hive queries.
- Automated schema drift detection in raw warehouse feeds.
- Participated in sprint planning and QA sessions to support production stability.
- Designed small internal tools to simulate data under different weather scenarios.

🎓 EDUCATION

Master of Science in Data Science
University at Buffalo, SUNY – Buffalo, NY
May 2024

- Entered U.S. on F1 in Jan 2023
- Focus Areas: Machine Learning, Deep Learning, NLP, Cloud-Based AI Systems
- Capstone: Built a clinical NLP pipeline using BERT to extract medical findings from EHR text

Bachelor of Engineering in Computer Engineering
Vasavi College of Engineering – Hyderabad, India
Aug 2021

- Gained 1.5 years of hands-on industry experience during 3rd and 4th year
- Projects in ETL optimization and early-stage ML forecasting

---

🧪 PROJECT HIGHLIGHTS

- **Clinical NLP Pipeline**: Developed transformer-based model using BERT to classify clinical symptoms from EHR text, achieving high F1-score.
- **Energy Prediction Web Platform**: Architected end-to-end ML pipeline (LightGBM + Flask + Docker) to forecast energy usage based on historical consumption.

- **Route Optimization Model**: Designed and validated predictive models for logistics efficiency, integrated into Tableau for real-time tracking.
- **Databricks MLflow Integration**: Enabled full experiment tracking, model versioning, and team collaboration in a reproducible ML lifecycle.

📜 CERTIFICATIONS

- **AWS Certified Machine Learning – Specialty**
- **Microsoft Azure Data Scientist Associate**
- **Databricks Data Engineer Foundations**