

Sai Teja Sirikonda

Senior Azure Data Engineer

E-Mail: sirikondasai007@gmail.com

Phone: 312-857-8733

LinkedIn: www.linkedin.com/in/sai-teja-sirikonda-461093324

PROFESSIONAL SUMMARY:

- 11+ years of IT experience as an Azure Data Engineer, Big Data Developer, and Data Warehouse Developer, with a focus on designing efficient and scalable data solutions to support business intelligence and advanced analytics.
- Specialized in wide range of Azure services, including Azure Data Lake Storage, Azure Synapse Analytics, Azure Data Factory, Azure Databricks, Azure SQL Database, Azure Cosmos DB, Azure Stream Analytics, Azure Functions, and Azure Key Vault, to build robust, scalable, and secure data solutions.
- Skilled in planning and executing data migration projects from on-premises databases such as SQL Server, Oracle to cloud based solutions such as Azure SQL Database, Azure Synapse Analytics
- Extensive experience in developing robust data pipelines in Azure Data Factory (ADF) to automate and orchestrate ETL workflows and utilized Azure Functions and Logic Apps within ADF pipelines for custom processing and orchestration.
- Demonstrated proficiency in leveraging Azure Synapse Analytics to design and implement scalable data warehousing solutions, enabling integration, advanced analytics, and optimized performance for complex data processing workflows.
- Expert in Azure Data Lake Storage Gen2 as a unified data repository, integrating structured and unstructured data from various sources to support big data analytics and data warehousing solutions, with extensive experience in data lake architecture on Azure Data Lake Storage.
- Developed and placed into place highly scalable data pipelines with Azure Data Factory (ADF), combining various sources such as REST APIs for centralised analytics and SQL Server and Oracle.
- Developed and deployed Spark-based ETL dataflows within Azure Databricks for batch and streaming analytics using PySpark and Delta Lake.
- Strong experience in Azure Databricks workspaces utilizing python and Pyspark notebooks and python jobs to perform data cleaning, aggregation, data processing, complex data transformations, including data normalization and denormalization.
- Proficient in employing Spark SQL and DataFrame operations to reshape data according to business requirements and engineered dynamic SQL queries to enable ad-hoc analysis capabilities for business users.
- Engineered Lakehouse architectures with Delta Lake for unified batch and streaming data, supporting ACID transactions, schema evolution, and time travel features.
- Hands-on experience with Lakehouse Architecture using Delta Lake, enabling ACID transactions, schema evolution, time travel, and zero-copy cloning for reliable data engineering.
- Skilled in AI/ML integration with Azure Databricks and Azure Machine Learning for predictive analytics, demand forecasting, anomaly detection, and real-time decision-making.
- Optimized Azure Synapse Analytics dedicated SQL pools through advanced partitioning, index management, and query tuning to boost performance for enterprise workloads.
- Strong knowledge of data governance and compliance frameworks using Azure Purview, Unity Catalog, Microsoft Fabric, and Microsoft Defender for Cloud.
- Hands-on experience in designing and implementing hybrid data architectures across Azure, Snowflake, and on-premises systems, ensuring seamless data integration and interoperability.
- Experience in integrating Azure Event Hubs with Azure Stream Analytics to enable real-time analytics and monitoring, capturing high-velocity data streams from IoT devices and POS systems for scalable and reliable data ingestion.
- Specialized in designing and implementing ETL pipelines in Azure using Python, leveraging libraries such as Pandas and NumPy for data cleansing, transformation, and loading into Azure data services.
- Implemented real-time ingestion and analytics solutions using Azure Event Hub, Stream Analytics, and Azure Functions to process high-velocity IoT and transactional data.
- Specialized in PolyBase to seamlessly integrate data from Storage, Azure Data Lake, and on-premises Servers into Azure SQL Data Warehouse, leveraging Azure SQL Database for scalable storage and structured data management.
- Skilled in Azure Key Vault to securely manage and store encryption keys, secrets, and certificates and managed identities and governance using Azure Active Directory (Microsoft Entra ID), including user authentication and Single Sign-On.
- Extensive experience in Azure SQL Database instances to store metadata and integrated Azure SQL Database with other Azure services such as Azure Data Factory for ETL processes.
- Hands on experience in orchestrating data integration pipelines using Azure Synapse Pipelines to ingest, prepare, and visualize data from sources including Azure Blob Storage, Azure SQL Database, and external data feeds within Azure Synapse Analytics using tools like Power BI or Azure Synapse Studio.
- Expert in integrating Cosmos DB with Azure services such as Azure Functions, Azure Databricks, and Azure Synapse Analytics to build scalable data processing and analytics pipelines, leveraging its flexible schema and globally distributed architecture for scalable NoSQL applications.
- Implemented event-driven workflows in Azure Logic Apps to trigger actions based on events from Azure Event Grid, Azure Service Bus and Azure Functions to process data streams, perform ETL operations, and handle data transformations.

- Experienced in monitoring dashboards in Azure Monitor to track data pipeline performance metrics, enabling proactive optimization and troubleshooting.
- Skilled in integrating data from diverse sources including RDBMS and file formats like CSV, JSON, XML, Excel, Avro, Parquet and REST APIs into unified data pipelines.
- Proficient in using Apache Hadoop for distributed storage and processing of large data sets, including expertise in HDFS, MapReduce, YARN, Hive and the Hadoop ecosystem.
- Extensive experience with Apache Spark for large-scale data processing and real-time analytics, utilizing Scala, python, Spark SQL and Spark Streaming
- Experience in setting up and managing Apache Kafka for real-time data streaming and integrated Kafka with Azure Event Hubs, Azure Data Factory, and Azure Stream Analytics for seamless data flow and processing.
- Specialized in partitioning and bucketing techniques in technologies such as Hadoop, Apache Cassandra, and Azure Cosmos DB to distribute large datasets across multiple storage units.
- Skilled in database languages such as SQL and PL/SQL, including writing triggers, stored procedures, functions, views, and cursors and experienced in optimizing PL/SQL code for performance, including query tuning, cursor management, and minimizing context switches between SQL and PL/SQL.
- Proficient in utilizing Azure Fabric for building distributed applications and microservices architectures, ensuring seamless scaling, high availability, and fault tolerance, along with extensive experience in SQL Server Management Studio (SSMS) for database management, query optimization, and troubleshooting, enabling efficient management of on-premises and cloud databases across Azure environments.
- Familiarity with MySQL databases for various applications, implementing schema design, query optimization, and performance tuning, and expertise in database schemas and models using SQL DBM, ensuring accurate representation and proficient in data modeling, including snowflake and star schemas to optimize data organization and retrieval efficiency.
- Managed code repositories on Bitbucket and Git for version control, facilitating collaborative development and code review processes, and worked with CI/CD pipelines using Azure DevOps integrated with Jenkins to automate build, test, and deployment processes for continuous integration and delivery.
- Implemented a robust data processing pipeline using Apache Kafka for real-time stream processing, orchestrated with Apache Airflow on Azure Kubernetes Service (AKS) for scalable and efficient task scheduling and monitoring.
- Proficient in implementing Agile practices using Jira, including sprint planning, scrum meeting, and retrospectives, to enhance team efficiency and project delivery and proficiency in shell scripting for automation tasks in data pipelines.
- Strong communication, analytical and problem-solving skills with the ability to quickly adapt to new environments and learn new technologies and expertise in various phases of project life cycles (Design, Analysis, Implementation and testing).

TECHNICAL SKILLS:

Azure Technologies:	Azure Data factory (ADF), Azure Databricks, Azure Synapse Analytics, Azure Data Lake Storage Gen2, Azure Storage, Azure Event Hub, Azure IOT Hub, Azure SQL Database, Azure Data Catalog, Azure Logic Apps, Azure Functions, Azure Blob Storage, Azure DevOps, Azure Virtual Machine, Azure Active Directory, Azure Monitor, Azure Purview, Azure Cosmos DB, Azure Key Vault, Azure Repos, Azure HD Insight, Azure Kubernetes Services, Azure Stream Analytics
Big Data Technologies:	HDFS, Map Reduce, YARN, Hive, SQOOP, Apache Spark, Kafka, Airflow, Spark Performance Tuning, Polybase, HQL, OOOIE, Zookeeper, Pig, Flink, Spark Streaming, Cloudera, Hortonworks, Databricks Delta, Apache NiFi, Snowflake, Azure Synapse, Amazon Redshift, dbt, Informatica PowerCenter, Talend.
Databases:	MS SQL Server 2016/2014/2012, Azure SQL DB, Azure Synapse, MS Excel, Oracle 11g/12c, MS SQL, Cassandra, MS Access, Cosmos DB, MongoDB, PostgreSQL, Star Schema, Snowflake Schema, Dimensional Modelling, Snowflake, Partitioning, Clustering, Materialized Views.
CI/CD Tools:	GIT, GitHub, Jenkins, Bitbucket, GitLab, Azure DevOps, Jenkins, Terraform, Kubernetes, Docker, Helm, GitHub Actions, YAML Pipelines, CloudFormation, Bicep, RBAC, IAM Policies, Security & Compliance.
Programming Languages:	Java, Python, PySpark, Shell Script, SQL, Scala, HiveQL, PL/SQL, T-SQL.
IDE & Build Tools:	PyCharm, Eclipse, Visual Studio, SSIS, SSAS, SSRS, SSMS.
Reporting Tools:	Tableau, Power BI, Ant, Maven, Terraform, Microsoft Fabric.
Operating Systems:	Windows (XP/7/8/10), UNIX, LINUX, UBUNTU, CENTOS.
Methodology:	Apache Airflow, Prefect, Agile, Scrum.

EDUCATION:

- Masters – Governor's state university, University Park, Illinois. (December 2013)
- Bachelors – Vignan institute of technology, Hyderabad, Telangana, India. (May 2011)

CERTIFICATIONS:

- Microsoft Certified Azure Data Engineer Associate (DP-203)
- Microsoft Certified Fabric Data Engineer Associate (DP-700)
- Data Bricks Certified Data Engineer Associate

PROFESSIONAL EXPERIENCE:

Optum, Chicago, IL.

January 2022- Current

Senior Azure Data Engineer

Responsibilities:

- Utilized Azure Data Factory (ADF) for data ingestion, implementing activities such as copy data, look-up, for-each, and if condition activities, with scheduled and storage event triggers.
- Automated data ingestion from multiple on-premises SQL Servers with self-hosted integration runtime, into Azure Data Lake using Azure Data Factory (ADF), facilitating seamless data integration and consolidation.
- Utilized scalable data storage solutions using Azure Data Lake Storage (ADLS) to handle large volumes of structured and unstructured data and utilized Azure Databricks and Azure HDInsight for processing and transformations with Pyspark.
- Incorporated Azure Data Lake Storage Gen2 (ADLS) and Azure SQL Database for efficient data storage and retrieval, decreasing data access latency by 30%.
- Leveraged Azure Databricks with Pyspark and python for data transformation, performing validation, profiling and cleansing using Spark jobs, and optimized data manipulation with Data Frames over s for efficient query execution.
- Designed and implemented AI-driven analytics solutions leveraging Azure ML and Databricks to detect fraud, optimize supply chain operations, and analyze customer behavior, resulting in faster insights and improved decision-making.
- Established a robust data and analytics platform using Azure Synapse Analytics, enhancing data quality, decision-making reliability, and operational efficiency and integrated Azure Active Directory and Azure Key Vault, improving resource and identity management efficiency.
- Automated the deployment of Azure resources using Infrastructure as Code tools, including Terraform and ARM templates, streamlining production workflows and compliance enforcement.
- Improved data processing by setting up parallel and incremental loading in Azure Data Factory, which reduced delays and handled large amounts of fast-moving data effectively.
- Implemented Lakehouse architecture using Delta Lake with schema evolution, time travel, and historical versioning to ensure reliable and scalable data pipelines.
- Created custom connections with Azure Event Hub, automating setup tasks and using Python scripts to handle real-time data streaming.
- Implemented Azure API Management to securely integrate third-party logistics and enterprise business systems, enforcing authentication, throttling, and monitoring policies for reliable data exchange.
- Enhanced data governance and compliance by implementing data lineage, masking, and auditing using Purview, Unity Catalog, and Microsoft Fabric.
- Set up Azure DevOps CI/CD pipelines with templates, environment-specific settings, secrets management, and approval steps, ensuring faster, safer, and more reliable deployments.
- Integrated Azure Logic Apps with Azure Blob Storage, Azure SQL Database, and Azure Data Lake Storage for batch data processing, enabling automated data workflows and notifications.
- Built reusable and flexible pipelines in Azure Data Factory using parameters and dynamic settings, making them easy to maintain and use across different business needs.
- Automated data governance with Microsoft Purview and Databricks Unity Catalog to track metadata, enforce policies, manage sensitive data, and give better visibility across the organization.
- Utilized Azure Event Grid to create event-driven architectures, allowing Azure Functions to react to data events from services like Azure Blob Storage, Azure SQL Database, and Cosmos DB, ensuring real-time batch processing capabilities.
- Fine-tuned Apache Spark jobs on Azure Databricks by adjusting cluster configurations and optimizing data caching strategies within PySpark and Python Databricks notebooks and clusters, resulting in reduced processing time for data.
- Implemented partitioning strategies on Azure Synapse Analytics to optimize data warehouse storage and query performance for complex analytical workloads and designed complex data models.
- Integrated Azure Machine Learning models with Databricks and Synapse for predictive maintenance, demand forecasting, and real-time anomaly detection.
- Implementation and deployment of AI/ML models for business applications (demand forecasting, anomaly detection) with Azure ML, AKS, Databricks, MLflow.
- Implemented monitoring, alerting, and troubleshooting using Azure Monitor and Application Insights to proactively detect and address issues, ensuring optimal system uptime, performance, and reliability across data pipelines.
- Optimized SQL queries and stored procedures on Azure SQL Database, improving query response times by 30% through index tuning and query plan analysis, while also creating and optimizing complex stored procedures, functions, and views to facilitate comprehensive data transformations and robust business logic implementation within batch pipelines.
- Implemented Delta Lake and Delta Live Tables on Azure Databricks to manage and optimize batch data processing with ACID transactions, leveraging Python, PySpark, Databricks, Spark SQL, and Data Frames API for complex data transformations, aggregations, and enrichment tasks, thereby enhancing data quality and insights.
- Hands-on Azure IoT Hub/Event Hub integration for real-time sensor/RFID/GPS/temperature data in supply chain, logistics, or warehouse solutions.

- Utilized dedicated SQL pools for high-performance batch data processing, leveraged Synapse Spark pools for big data batch processing tasks, and queried data lakes using serverless SQL pool for ad-hoc batch analysis.
- Implemented Change Feed Processor for distributed batch updates processing and integrated with Azure Synapse Link for seamless batch analytics and developed and managed data policies and standards using Azure Purview to maintain data quality and consistency across various data sources.
- Real-world Lakehouse/Delta Lake projects leveraging advanced features (ACID, time travel, cloning, schema evolution) and their enterprise benefits.
- Extracted and processed data from Integrated Banking Solution, Cash Dispensing Terminals, Digital Banking Platforms (transaction Processing), and Customer Relationship Management (CRM) suites, using Azure Data factory ensuring comprehensive coverage of all critical banking operations.
- End-to-end data governance and compliance programs (GDPR, HIPAA, FDA, HACCP) featuring data masking, lineage, Purview, Unity Catalog, audit, and secure key management.
- Management of hybrid cloud/on-prem enterprise migrations with performance tuning and minimal downtime using orchestration tools (ADF, PolyBase, SSIS, ARM).
- Utilized Azure Blob Storage's Block Blob feature within Azure Data Lake Store Gen2 (ADLS) to efficiently handle large scale batch uploads, achieving ingestion rates of 20TB/hour for unstructured data.
- Engineered executive analytics dashboards using Azure Synapse and Power BI, leveraging advanced DAX, power query and data modelling techniques to improve strategic planning accuracy by 30%, while utilizing Power BI for comprehensive data visualization including executive summaries, customer insights, transaction analysis, and financial performance reports.
- Utilized Azure Resource Manager (ARM) for deploying batch processing infrastructure and tracked batch pipeline development using agile methodologies with Azure Boards.
- Broader CI/CD, DevOps, and Infrastructure-as-Code practices using Terraform, Azure DevOps, Jenkins for full analytics/BI pipeline automation and deployment.
- Facilitated knowledge-sharing sessions on best practices in data engineering and azure technologies, fostering a culture of learning and innovation within the team, while engaging stakeholders including business analysts and data scientists to gather requirements and refine user stories, ensuring continuous alignment of data engineering efforts with business needs.
- Collaborated in Agile sprints as part of a cross-functional team, participating in daily stand-ups and bi-weekly sprint planning sessions to prioritize and deliver data solutions aligned with business goals.

Environment: Azure Data factory, Azure Databricks, Azure Data Lake Storage Gen2, Azure Synapse Analytics, Azure Logic Apps, Azure Functions, Azure Monitor, Azure SQL Database, Azure Cosmos DB, Azure Resource Manager, Azure Repos, Azure HDInsight, Azure Active Directory, Azure Blob Storage, Cosmos DB, Power BI.

3M, Saint Paul, MN.

March 2017 – December 2021

Azure Data Engineer

Responsibilities:

- Implemented and managed linked services in Azure Data Factory (ADF) to securely connect and integrate with various data sources and destinations such as Azure SQL Database, Azure Blob Storage, and on-premises databases.
- Utilized Azure Data Factory (ADF) scheduling capabilities, including tumbling window, event-based, and time-based triggers, to automate the execution of data pipelines, ensuring timely data processing and integration.
- Utilized Azure Event Hub to create robust real-time streaming pipelines, capturing and processing live data streams from sources such as point-of-sale systems, IoT devices, and e-commerce platforms.
- Built orchestration pipelines in Azure Data Factory with reusable templates, making ETL workflows modular, scalable, and easy to maintain.
- Integrated Azure Event Hubs with Azure Stream Analytics for real-time data processing and facilitated the migration of existing Apache Kafka applications to Azure Event Hubs without code changes.
- Developed incremental and delta data loading in ADF using watermark tables and change-tracking logic, reducing processing time and cloud costs.
- Developed automated workflows using Azure Functions for seamless data migration across azure services, and utilized Azure Logic Apps for real-time processing, triggering actions and alerts based on streaming data events.
- Leveraged Azure Databricks tools such as Workflows, Unity Catalog, Delta Live Tables, and Databricks SQL to enhance ETL processing, performance, and data governance.
- Used Azure Event Hubs Premium with dynamic scaling, redundancy, and event retention to handle streaming data for IoT and fraud detection.
- Architected and implemented a collaborative environment using Azure Databricks, leveraging version-controlled notebooks for team-wide code sharing, iterative development, and scheduling PySpark jobs for automated execution to enhance operational efficiency and scalability of data workflows.
- Automated workflows through REST API, CLI, Python, Terraform, and Directed Acyclic Graphs (DAGs) for seamless integration and deployment of data pipelines and clusters.

- Delivered real-time analytics in Azure Synapse by combining Apache Spark and serverless SQL pools, enabling quick insights from both streaming and batch data for supply chain and operations.
- Utilized Synapse Link in Azure Synapse Analytics for real-time analytics on operational data from Cosmos DB, leveraging serverless SQL pools for ad-hoc querying of streaming data and dedicated SQL pools for high-performance processing.
- Utilized Azure Data Lake Gen2 (ADLS) as an output sink for Azure Stream Analytics jobs, implementing partitioning and output strategies for optimized write performance, configuring fine-grained POSIX-like ACLs for secure access to streaming data.
- Configured Azure Cosmos DB Time-to-Live (TTL) for automatic data expiration in streaming scenarios, implemented hot partition mitigation techniques for even data distribution, and tuned request units (RUs) for optimal performance.
- Deployed and managed Kafka clusters on Azure Kubernetes Service (AKS) for high-throughput data streaming and implemented scalable event processing using Azure Kubernetes Service and Azure Event Hubs.
- Integrated Azure Monitor with Azure Event Hubs, Azure Databricks, and Azure Data Factory to monitor end-to-end streaming ETL workflows, optimize performance, and utilize Azure Monitor REST API for programmatic access to monitoring data.
- Led the migration of a legacy data warehouse to Azure SQL Database, resulting in a 20% reduction in query execution time, and utilized DMS for efficient data transfer, minimizing downtime during migration.
- Implemented Python-based libraries utilizing Scikit-learn and PyTorch frameworks on Azure Databricks, optimizing for scalable data processing and advanced analytics capabilities through scheduling, ETL automation, and scripting.
- Identified and resolved query performance issues in Synapse SQL pools by analyzing query execution plans, optimizing SQL code, and adjusting resource allocation.
- Implemented data storage strategies such as tiered storage and performance tuning for large-scale data workloads in Azure Data Lake Storage Gen2 (ADLS) using Pyspark to optimize cost and performance.
- Managed schema evolution and data versioning in Azure Data Lake Storage (ADLS) using Delta Lake, Pyspark and Databricks SQL Analytics to accommodate changes in streaming data structures without disrupting downstream processes.
- Tuned Azure Synapse Analytics by leveraging materialized views, optimizing table distribution and partitioning, and using dedicated SQL pools for high-performance query processing.
- Developed ETL solutions for incremental data loading using Azure Data Factory and Change Data Capture (CDC) techniques, leveraging PySpark for data transformation and integrated Slowly Changing Dimension (SCD) managed data with BI tools like Tableau to deliver accurate and historical insights.
- Employed Azure DevOps with Terraform to define and provision infrastructure as code and automated the deployment of Azure resources, such as Azure Data Factory, Azure Databricks, and Azure SQL Database, ensuring consistency, repeatability, and scalability during data migration, using Python scripts for custom deployment tasks.
- Designed and implemented Continuous Integration and Continuous Deployment (CI/CD) pipelines using Azure Pipelines integrated with Docker to automate the build, test, and deployment processes of data pipelines and migration workflows.
- Coordinated with cross-functional teams, including data scientists, business analysts, and DevOps engineers, to gather requirements, refine user stories, and ensure alignment with project goals.
- Regularly communicated project status, risks, and milestones to both technical and non-technical stakeholders while developing dashboards for real-time visualization of pipeline performance and migration progress.

Environment: Azure Data Factory (ADF), Azure Event hub, Azure Stream Analytics, Azure Functions, Azure Logic Apps, Azure Databricks, Azure Synapse Analytics, Azure Data Lake Gen2, Azure Cosmos DB, Azure Kubernetes Service, Azure HDInsight, Azure Monitor, Azure SQL Database, Azure DevOps.

Zurich Insurance, Schamburg, IL.

December 2015 - February 2017

Big Data Engineer

Responsibilities:

- Responsible for analyzing the business requirement and estimating the tasks and preparing the mapping design documents for Confidential Point of Sale (POS) and direct sales across all GOE's.
- Utilized Cloudera distributions to build scalable data platforms, leveraging Hadoop ecosystem tools like HDFS, MapReduce, Hive, and HBase to manage and process large-scale datasets efficiently.
- Developed Spark applications using Scala, implementing Apache Spark data processing projects to handle insurance data, and proficient in leveraging Scala's functional programming paradigms for creating scalable and performance applications.
- Used Spark Streaming APIs to perform near real-time transformations and actions from Kafka, persisting results to Cassandra, and familiar with PySpark's Spark Streaming module for real-time processing of data streams from sources like Kafka, Flume, or TCP sockets.
- Managed and administered Hadoop Distributed File System (HDFS), ensuring scalability, reliability, and high availability of data storage for big data processing.
- Utilized DDL commands to define and manage Hive schemas, including external tables, partitions, and bucketing, and wrote complex DML queries to extract, transform, and load data.
- Experienced in writing and optimizing HiveQL queries for data transformation, aggregation, filtering, and joining operations on structured and semi-structured data stored in Hadoop, optimizing 5% of data querying efficiency.
- Implemented distributed locks and synchronization primitives using Zookeeper, and proficiently used Apache Sqoop for efficiently transferring bulk data between Hadoop and external data sources.

- Utilized Apache Pig Latin, a high-level scripting language, for data transformation, ETL (Extract, Transform, Load), and analytics tasks on large datasets stored in Hadoop.
- Experienced in designing, deploying, and managing HBase for real-time and skilled in designing and scheduling workflows using Apache Oozie to automate and coordinate Hadoop jobs and data processing pipelines.
- Improved query performance across SQL, HiveQL, and Spark SQL by analyzing execution plans, implementing appropriate indexing and partitioning strategies, optimizing joins, and applying data manipulation practices in PL/SQL code.
- Expertise in SQL Server for managing and optimizing relational databases, designing schemas, and writing complex SQL queries, along with skilled optimization of Teradata SQL queries using performance tuning techniques, indexing strategies, and workload management.
- Integrated Spark SQL with various Spark components and external data sources like Hive and JDBC/ODBC for data processing and analytics, while also leveraging Spark RDD API for low-level data manipulations and transformations.
- Experienced in UNIX shell scripting (Bash, Korn shell) for automating system tasks, file manipulation, process management, and data processing workflows.
- Skilled in SQL for querying, managing, and manipulating data within relational database management systems (RDBMS) such as Oracle, SQL Server, and MySQL.
- Experienced in using DAX for defining relationships, calculated columns, measures, and hierarchies, and integrated PySpark with Python libraries for advanced analytics workflows.
- Implemented ETL processes for importing and exporting data between MS Access, Excel, and SQL Server, facilitating data integration and reporting.
- Skilled in Oracle Database for designing, developing, and administering relational databases, ensuring data integrity, security, and performance including ETL operations.

Environment: Hadoop, Hive, Pig, Spark, Hbase, Zookeeper, Scala, Kafka, Cassandra, HDFS, MapReduce, YARN, Kafka, HiveQL, Sqoop, Oozie, SQL, Spark SQL, PL/SQL, SQL Server, Teradata, Unix Shell Scripting, DAX, MS Access, Oracle.

Citi, Austin, TX

January 2014- November 2015

Data Warehouse Developer

Responsibilities:

- Experience in Windows Server administration, including installation, configuration, and management of server environments for hosting applications and databases.
- Expertise in designing, developing, and optimizing relational databases in MS SQL Server 2014 and Oracle Database, with proficiency in writing complex SQL queries, stored procedures, functions, and triggers for data analysis and reporting.
- Implemented data standardization rules and cleansing transformations in Informatica Data Quality to improve data accuracy and consistency, along with hands-on experience in developing ETL processes using Teradata utilities like BTEQ, Fast Load, Multiload, and Fast Export for efficient data extraction, transformation, and loading.
- Proficient in using SQLDBM to design and visualize database schemas, including tables, relationships, and constraints, for data warehouse solutions.
- Implemented star and snowflake schema designs in OLAP systems to optimize query performance for interactive and ad hoc analytical queries while also implementing concurrency control mechanisms in OLTP databases.
- Extensive experience in using Informatica PowerCenter 9.1 for designing, developing, and deploying ETL processes, data mappings, and workflows, as well as configuring SSIS data flow tasks such as Data Conversion, Conditional Split, Aggregate, Lookup, and Derived Column for efficient data transformations.
- Skilled in writing MDX (Multidimensional Expressions) and DAX (Data Analysis Expressions) queries for data analysis, calculations, and reporting within SQL Server Analysis service (SSAS).
- Familiar with integrating Microsoft SQL Server Reporting Services (SSRS) reports into SharePoint portals and other applications for centralized report access and distribution.
- Experienced in Oracle Database 11g administration, covering installation, configuration, backup and recovery, and performance tuning, and implemented SharePoint for document management, team collaboration and workflow automation.
- Skilled in using SQL Profiler traces to identify database issues and improve performance and created complex SQL queries for data analysis, reporting, and BI tasks, with a focus on improving query speed and data access.
- Automated ETL processes using shell scripts for data extraction, transformation, and loading, and experienced in SharePoint for document management, team collaboration, workflow automation, and enterprise content management.
- Created interactive dashboards and visualizations using Tableau by connecting to various data sources and designing data models for business intelligence, and skilled in utilizing advanced Excel functions, formulas, pivot tables, and macros for data manipulation, analysis, and reporting purposes.

Environment: Windows Server administration, MS SQL Server, Oracle, Informatica data quality, Teradata, Informatica Power center, SSAS, SSRS, SSIS, Shell Scripts, Tableau.