Sakshi Sherikar | Data Scientist

| +1 (334) 865-5502 | sakshisherikar9@gmail.com | LinkedIn | GitHub

SUMMARY

Results-driven Data Scientist with 3 years of experience in designing and deploying scalable machine learning solutions across cloud platforms (AWS, GCP, Azure) and big data ecosystems (Spark, Hadoop) including Generative AI solutions. Expert in applying transformer-based models for NLP tasks such as text generation, summarization, and integrated into scalable data science workflows. Proven expertise in predictive modeling, NLP, and time-series forecasting, improving model accuracy, and automating workflows to save time. Adept at translating complex data into actionable insights through Tableau/Power BI dashboards and MLOps tools. Strong collaborator with experience in Agile environments, cross-functional stakeholder alignment, and end-to-end data pipeline optimization. Passionate about ethical AI and data-driven decision-making.

TECHNICAL SKILLS

Programming & Query Languages: Python (Pandas, NumPy, PySpark), SQL, R, Scala
Gen AI & ML: Scikit-learn, TensorFlow, PyTorch, XGBoost, NLP (spaCy, BERT, NLTK), Time-Series Forecasting, Hugging Face, LangChain
Big Data & Cloud: AWS (S3, Redshift, SageMaker, Lambda), GCP (BigQuery, DataFlow), Azure (Data Lake), Hadoop, Spark
Data Engineering: ETL/ELT Pipelines, Airflow, Databricks, Delta Lake, Lakehouse Architecture
Data Visualization: Tableau, Power BI, Matplotlib, Plotly, Dash
Databases: MySQL, PostgreSQL, AWS Redshift, MongoDB
Model Deployment & MLOps: Docker, Kubernetes, MLflow, CI/CD, REST APIs
Statistics & Experimentation: A/B Testing, Bayesian Methods, Hypothesis Testing
Collaboration & Tools: Git, Agile (Scrum), Jira, Jupyter Notebooks, Cloudera Data Science Workbench (CDSW)

PROFESSIONAL EXPERIENCE

Data Scientist Intern | Cloudera — Jan 2025 – Present | Remote, USA

- Designed and deployed a predictive maintenance model for IoT device telemetry using PySpark and TensorFlow on the Cloudera Data Platform (CDP), improving equipment failure detection accuracy by 31% and helping reduce annual downtime-related costs by 15%.
- Spearheaded the transition to a Lakehouse architecture using Apache Iceberg, Spark SQL, and Hive Metastore, enabling unified metadata access and improving cross-team data accessibility by 2.5x while reducing cloud storage expenses by 22%.
- Fine-tuned BERT-based transformer models for multilingual sentiment analysis on customer feedback, improving classification accuracy by 44% and driving a 22% increase in CSAT scores through data-backed service improvements.
- Developed and maintained reusable ML application templates (AMPs) for fraud detection and demand forecasting, standardizing pipelines across teams and cutting the average model development timeline by 50%.
- Built real-time Tableau dashboards integrated with Cloudera DataFlow pipelines and automated cross-platform ETL workflows connecting AWS S3 with on-prem Hadoop, enabling business executives to track KPIs with 45% faster refresh cycles, reducing data ingestion latency by 30%, and supporting near-real-time analytics across multiple business units.
- Applied Generative AI techniques by fine-tuning transformer models (BERT, Hugging Face, PyTorch) for multilingual sentiment analysis on customer feedback, improving classification accuracy by 44% and boosting CSAT scores by 22%.

Data Scientist | IBM — Jan 2021 – Jun 2023 | India

- Designed and implemented three predictive models (logistic regression, random forest, XGBoost) for classification use cases, achieving up to 96% accuracy by optimizing feature selection and hyperparameter tuning, thereby increasing model reliability by 22%.
- Automated end-to-end data preparation workflows using Python libraries (Pandas, NumPy), deployed them on AWS S3/Redshift environments, and reduced manual data wrangling time by over 60% across projects.
- Created interactive Power BI dashboards for internal reporting and stakeholder reviews, reducing report generation times by 40% and increasing usage of data-driven decision-making by leadership by 50%.
- Scaled ML experimentation by connecting AWS Redshift data sources to model training pipelines, enabling faster retraining and improving data accessibility by 45% for the analytics team.
- Conducted A/B testing on digital marketing strategies using StatsModels, lifting in conversion rates, and delivered bi-weekly performance reviews to cross-functional teams, ensuring alignment of model results with business KPIs and reducing analytics-to-operations disconnect by 30%.
- Supported migration of legacy SQL-based workflows to IBM Cloud infrastructure, improving query efficiency by 35% through index optimization, query refactoring, and partitioning strategies.

EDUCATION

Master of Science in Computer Science — **Aug 2023 – May 2025**
Auburn University at Montgomery – AL

PROJECTS

Generative AI Chatbot for Customer Support *(Python, Hugging Face, LangChain, Pinecone, OpenAI APIs)*

- Designed and deployed a RAG-based chatbot using GPT and Pinecone, fine-tuning embeddings on 50K+ support tickets to boost automated resolution by 87% and cut hallucination rates.
- Integrated with Slack and MS Teams, reducing Tier-1 support workload by 40% and improving average response time by 65% through optimized data preprocessing and clustering.

AI-Powered Resume Builder for Job Seekers *(Python, Transformers (BERT, GPT), Streamlit, FastAPI, AWS, MongoDB, Docker)*

- Built and deployed an ATS-friendly GenAI resume builder leveraging BERT/GPT pipelines for skill extraction, summary generation, and role matching, improving recruiter response rates significantly.
- Developed a Streamlit-based UI with FastAPI backend, deployed on AWS using Docker, and integrated MongoDB for scalable user data management and real-time resume rendering.

CERTIFICATIONS

• IBM Data Science Internship • Machine Learning with Python