

Sathish Kumar Reddy Daggula

Tel: +1(314) 325-9624 | Email: satishkumarreddy011@gmail.com

PROFESSIONAL SUMMARY

Data Engineer with over 3+ years of expertise in developing and implementing high-performance data solutions. Skilled in programming languages such as Python and SQL, with a comprehensive grasp of large-scale data processing technologies like Apache Spark. Experienced in leveraging cloud platforms, especially AWS and Azure. Proficient in architecting ETL processes, managing data repositories, and constructing data integration workflows. Familiar with platforms such as Databricks and Snowflake and experienced in employing data visualization tools. Expertise in creating and optimizing CI/CD workflows while adhering to DataOps best practices. Proven track record of problem-solving and effective communication. Adept at collaborating with cross functional teams to drive data-centric solutions and strategies.

SKILLS

- **Programming & Scripting:** Python, SQL, Scala, PySpark
- **Big Data Technologies:** HDFS, Hive, Pig, Sqoop, YARN, Apache Spark (Core, Spark SQL), Kafka
- **Machine Learning & NLP:** Linear/Logistic Regression, Decision Trees, Random Forest, SVM, KNN, K-Means, Gradient Boosting, AdaBoost, PCA, LDA, NLP (NLTK, SpaCy, Gensim)
- **Cloud Platforms & Services:** AWS (EC2, S3, Redshift, EMR, Lambda), Azure (Synapse, ML Studio), GCP (BigQuery, Dataflow, GCS)
- **ETL & Data Integration:** IBM InfoSphere (v8 – v9.1), SSIS, Autosys
- **Databases & Warehousing:** Snowflake, Teradata, Oracle 9i/10g, DB2, SQL Server, MySQL
- **IDEs & Development Tools:** PyCharm, Jupyter Notebook, Spyder
- **Reporting & Visualization:** Tableau, Power BI
- **Version Control:** Git, GitHub
- **Frameworks & Libraries:** TensorFlow, Keras
- **Other Tools & Utilities:** Google Analytics, MS Office Suite
- **Operating Systems:** Linux, Unix, Windows
- **Project Methodologies:** Agile, Data Security & Compliance, Customer Satisfaction

PROFESSIONAL EXPERIENCE

Globus Medical Inc

Jan 2024 – Present

Data Engineer

- Architected and deployed Spark-based pipelines on AWS (S3, Lambda, Glue) for processing large volumes of clinical data.
- Built ETL workflows in Databricks and Snowflake to ingest 5TB+ of patient, EHR, and claims data daily.
- Automated healthcare data workflows with Azure Data Factory and DBT, reducing manual interventions.
- Implemented data quality checks using Alation and custom validators, lowering clinical data inconsistencies by 60%.
- Collaborated with data science teams to prepare ML-ready datasets for patient outcome and risk models.
- Led integration of on-prem EMRs with cloud platforms, ensuring traceability and schema consistency.
- Drove migration of healthcare systems to cloud-native architecture for better scalability and performance.
- Promoted CI/CD in healthcare data projects via GitHub Actions, and mentored junior data engineers regularly.

Coforge

Jan 2021 -Nov 2022

Data Engineer

- Engaged in big data requirement analysis to develop and design solutions for ETL and Business Intelligence systems.
- Created 3NF data models for Operational Data Stores (ODS), Online Transaction Processing (OLTP) systems, and dimensional data models using Star and Snowflake Schemas.
- Worked within the Snowflake environment to eliminate data redundancy and integrate real-time data from various sources into HDFS via Kafka.
- Developed a comprehensive data warehouse model in Snowflake, incorporating over 100 datasets.
- Managed the ingestion and processing of structured and semi-structured data on AWS using S3 and Python, and facilitated the migration of on-premises big data workloads to AWS.
- Designed data aggregation processes in Hive for ETL tasks on Amazon EMR, ensuring alignment with business requirements.
- Supported the migration of data from traditional RDBMS to Hadoop using Sqoop, assessing the performance.
- Implemented data validation through MapReduce programs to filter out unnecessary records before transferring data into Hive tables.

- Designed and configured a Kafka cluster to handle high throughput of up to 1 million messages per second, utilizing Kafka producer API 0.6.3.
- Extracted real-time data feeds using Kafka and Spark Streaming, converting them into RDDs, processing them as DataFrames, and saving the results in Parquet format within HDFS.

EDUCATION

- **Master's degree in computer science**
Southeast Missouri State University, Cape Girardeau, MO
- **Bachelor's degree in computer science**
Karunya Institute of Technology and Sciences, Coimbatore, TN, IND

CERTIFICATION

- **DeepLearning.AI: Neural Networks and Deep Learning**
- **University of Michigan: Python Data Structures & Python for Everybody**