

SATHWIK REDDY KONUDULA

Seattle, WA | ☎: (518)-522-2514 | ✉: konudulasathwik@gmail.com | [linkedin.com/in/sathwik99](https://www.linkedin.com/in/sathwik99) | <https://github.com/Wiky123>

SUMMARY

- Software Engineer with 2+ years of experience in scalable distributed systems, AI infrastructure, and cloud-native platforms.
- Built fault-tolerant microservices and model-serving APIs using Golang, Python, Spring Boot, and Kubernetes.
- Designed and deployed inference pipelines optimized for LLM workloads, GPU utilization, and low-latency performance.
- Experienced with CI/CD, observability (Prometheus, Grafana, ELK), and multi-cloud deployments on AWS, GCP, Azure.

EDUCATION

Georgia Institute of Technology	Atlanta, GA
Master of Science: Computer Science (Online), Current GPA-4.00/4.00	January 2025
Stony Brook University	Stony Brook, NY
Bachelor of Science: Information Systems, GPA-3.52	February 2021 to May 2024

SKILLS

- **Programming Languages:** Golang, Python, Java, Scala, C#, Bash, SQL, JavaScript
- **Frameworks:** Spring Boot, .NET Core, gRPC, REST APIs
- **AI/ML:** TensorFlow, Hugging Face, Redis, Kafka, PyTorch (familiar), NLP, LLM Inference
- **Cloud & Infrastructure:** AWS (EC2, Lambda, S3, VPC, IAM, KMS), Azure, GCP, Terraform, Ansible, Puppet, ArgoCD, CloudFormation
- **Containers & Orchestration:** Docker, Kubernetes, Helm, Serverless, Microservices Architecture, Load Balancers
- **Observability & Monitoring:** Prometheus, Grafana, ELK Stack, Datadog, CloudWatch, Splunk
- **DevOps & CI/CD:** Jenkins, Screwdriver, Git, Postman, OAuth2, OpenID Connect, Agile/Scrum, Jira
- **Other Tools & Systems:** RabbitMQ, MQTT, Wireshark, Scalable Queues, Fault-Tolerant Systems, GPU Scheduling

WORK EXPERIENCES

Goldman Sachs (via Pyramid Consulting Inc)	Remote
Backend Software Engineer, Cache Infrastructure Systems	July 2024 to July 2025
• Developed scalable AI inference services using Java, and Spring Boot on AWS, integrated with Kubernetes and gRPC. • Designed high-throughput Redis-based model caching and GPU-optimized queueing for 99.99% uptime under peak loads. • Tuned performance for LLM-serving microservices with gRPC and REST APIs, optimized API response times below 50ms. • Integrated observability stack (Prometheus, Grafana, ELK) and auto-healing bots, cutting incident MTTR by over 50%. • Built cost-aware scheduling for large language models with Redis, Kafka, and load balancing logic.	
Tekgente Solutions Inc	South Plainfield, NJ
DevOps & Reliability Engineer – Hybrid Cloud & Edge Computing (Co-op)	February 2023 to May 2024
• Managed AI model deployment pipelines across AWS, Azure using Terraform, ArgoCD, and Ansible for consistent scaling. • Built ETL and inference pipelines on GPU-accelerated nodes with Kubernetes, Redis, and serverless model endpoints. • Maintained cloud-based LLM catalogs and secure inference queues with Kafka, gRPC, and S3-based checkpointing. • Reduced model load latency by 45% through real-time data streaming and container-based optimizations on edge clusters. • Implemented distributed offline transfer modules with 100% reliability across 50+ multi-cloud nodes using Puppet, Redis. • Applied SIEM tools (Splunk, ELK) to protect AI systems against DDoS, achieving 95%+ detection accuracy via ML models.	
Google	New York, NY
SPS Intern	May 2022 to August 2022
• Designed a Java-based web app using GCP services like App Engine and Datastore, improving deployment speed by 64%. • Automated CI/CD pipeline builds and Kubernetes cluster deployments via Screwdriver, reducing rollout time by 40%. • Implemented fault-tolerant design using Kafka and Redis, improving financial transaction reliability by over 40%. • Resolved complex component dependencies, reducing system errors by 90% and improving runtime stability significantly. • Contributed to open-source projects with distributed teams, performing code reviews and enhancing codebase quality.	

PROJECTS

Optimized Distributed Storage System AWS, Terraform, Python, Kubernetes
• Built a scalable distributed storage system, enabling multi-node replication and automatic failover. • Optimized data consistency and caching strategies using Redis & CloudFormation, reducing read latency by 45%. • Deployed infrastructure as code with Terraform and automated scaling policies, increasing system uptime and efficiency.
Automated Risk Management System for Financial Transactions Scala, gRPC, Kubernetes, Prometheus, Redis, Python
• Developed a high-frequency trading risk management system using Scala, Spring Boot, and PostgreSQL, enabling real-time detection of fraudulent transactions with sub-10ms response times. • Built an event-driven architecture leveraging Kafka and Redis, ensuring 99.9% uptime and instant anomaly detection alerts. • Designed API integrations for financial regulatory compliance (SEC, FINRA) using gRPC, OpenID, and OAuth, enabling real-time reporting and reducing compliance violations by 30%.