

Siva Mani Subrahmanya Hari Vamsi Pullipudi

+1(571) 723 2415 | spullipu@gmail.com | pssharivamsi.site | linkedin.com/in/pssharivamsi | github.com/PSSHARIVAMSI

Summary

Data Engineer / Data Analyst with over three years of professional experience designing and operating reliable data platforms. Skilled at developing and maintaining scalable pipelines that ingest millions of records, building unified data models, and collaborating with business teams to deliver dashboards and insights. Proficient in automating ETL workflows using Fivetran, Matillion and MuleSoft, tuning warehouse performance, and ensuring data quality. Adept at communicating findings to technical and non-technical stakeholders and supporting data-driven decision making.

Technical Skills

- **Data Engineering:** Snowflake, SQL, Fivetran, Matillion, MuleSoft, Snowpark Container Services
- **Analytics & Visualization:** Tableau, Power BI, Looker, Streamlit
- **Cloud & DevOps:** AWS (EC2, S3, Glue), Docker, Git, CI/CD, Azure DevOps
- **Programming Languages:** Python, R, C

Certifications

SnowPro Advanced Data Engineer | SnowPro Advanced Architect | SnowPro Core | Fivetran Certified | Coalesce Fundamentals | AWS Cloud Practitioner | MTA Python | CCNA | Juniper Networking

Professional Experience

Kipi.ai

Remote | Oct 2021 - Dec 2024 | 3 Years 2 Months

Lead Engineer / Senior Software Engineer / Software Engineer

- Led a team of four engineers to build the *LEAF Accelerator* Snowflake-Streamlit platform, delivering 25,000+ KPIs across multiple industries and eliminating more than 400 analyst hours each quarter. Established CI/CD standards for Snowflake native apps, raising deployment reliability to 99.9%.
- Architected an enterprise data consolidation platform for Oscar W. Larson Co., integrating seven source systems (SQL Server, Acumatica, Sage) into a unified Snowflake warehouse. Engineered 50+ MuleSoft pipelines with automated error handling and retry mechanisms, achieving 99.5% data reliability while processing over two million records daily and cutting report generation time from four hours to 15 minutes.
- Mentored six junior engineers on Snowflake optimisation and Matillion orchestration best practices, and collaborated with business analysts to design data models that improved accessibility and facilitated data-driven decision making.
- Executed an Exasol to Snowflake migration for Blue Yonder, converting 300+ legacy SQL templates, building Python CI pipelines for regression testing, and tuning warehouse configurations to reduce compute costs by 18%.

Projects

LEAF Accelerator – KPI and USECASES Platform

Sep 2024 – Dec 2024

Lead Engineer | Snowflake, Streamlit, AWS

- Designed and deployed a Snowflake-backed Streamlit native app surfacing 25 000+ KPIs and 10 000+ analytics use cases across manufacturing, banking, insurance, medical and logistics domains.
- Automated data ingestion and cleansing in Snowflake and containerised the solution on AWS.

PMO – Quality Gate Initiative

Jun 2024 – Aug 2024

Data Engineer | Salesforce, Excel

- Collaborated with project managers to implement a standardized quality gate process across 100+ active projects, designing and integrating best-practice checklists within Salesforce for enhanced compliance tracking.
- Optimized checklist workflows by leveraging Excel automation and Salesforce integration, improving consistency, audit preparedness, and adherence to organizational quality standards.

Financial Industrial Solutions – Geospatial Analytics

Mar 2024 – May 2024

Data Engineer | Python, Snowflake, Streamlit

- Built a Streamlit application for an insurance client to visualise geospatial data aggregations on interactive maps with dynamic filters, enabling rapid exploration of millions of spatial records.
- Managed data loading and cleansing in Snowflake, and containerised the app with Snowpark Container Services.

Oscar W. Larson Co. USA – Data Migration
Senior Consultant | Snowflake, MuleSoft, Acumatica, Sage

June 2023 – Feb 2024

- Architected a multi-source integration platform consolidating financial, sales and inventory data from seven enterprise systems into a unified Snowflake warehouse serving 100+ business users.
- Set up RBAC, databases, schemas, and tables in Snowflake while maintaining documentation of row counts and query performance metrics to ensure data integrity and operational transparency.
- Designed and implemented 50+ MuleSoft pipelines with advanced error handling and retry mechanisms, ensuring 99.5% data reliability across multiple business subsidiaries and integration layers.
- Optimised processing through Snowflake stored procedures and task scheduling, reducing report generation time from four hours to 15 minutes while processing millions of records daily.
- Established DevOps best practices using Git version control and automated deployment pipelines, improving code quality and reducing production incidents by 60%.

Query Analyzer – Snowflake Native App
Data Engineer | Snowflake, Streamlit

Feb 2023 – May 2023

- Designed the architecture and developed a Snowflake Native App, currently live on the Snowflake Marketplace that analyzes long-running queries and provides automated performance tuning recommendations.
- Implemented stored procedures and analytic queries to identify high-impact performance bottlenecks, increasing overall app performance by 60% through faster query execution, reduced latency, and optimized resource utilization.

Industrial Solutions – LLM Model
Solution Consultant | Python, AWS, Snowflake

Nov 2022 – Jan 2023

- Created a proof-of-concept large language model pipeline to answer questions from corporate 10-K and 10-Q reports; scraped and processed more than 50 corporate filings into PDFs and uploaded them to AWS S3.
- Developed Python routines to ingest documents into Snowflake and prototype question-answering workflows, demonstrating document intelligence capabilities.

Kipi Stone – Real-Time Stock Market Analytics
Senior Data Engineer | Snowflake, Kafka, Matillion, Tableau, AWS

Sep 2022 – Oct 2022

- Simulated and analyzed real-time NIFTY stock data by streaming live feeds from Yahoo Finance via Apache Kafka into Snowflake with sub-second latency, enabling timely insights for technical and fundamental analysis.
- Built robust data architecture with Snowflake (databases, schemas, transformations) to process 500,000+ daily transactions, orchestrated ETL pipelines with Matillion, and delivered interactive KPI dashboards in Tableau.

Snowflake Object Discovery
Data Engineer | Snowflake, Streamlit, Balsamiq

Jul 2022 – Aug 2022

- Developed a Streamlit web app to visualize enterprise-wide Snowflake accounts across 10 business units, surfacing datasets, users, and roles to support executive-level data governance decisions.
- Extracted metadata from Snowflake system views, designed wireframes using Balsamiq, and delivered a streamlined user interface with intuitive navigation in Streamlit.

Blue Yonder Germany – Exasol to Snowflake Migration
Consultant | Snowflake, Python, Docker, DBeaver, Jira

Jan 2022 – June 2022

- Migrated 300+ complex Exasol SQL templates to Snowflake, ensuring syntactic compatibility, improved query performance, and increased data accuracy across critical reporting workflows.
- Optimized Snowflake warehouse configurations by implementing right-sizing strategies and auto-suspend/resume policies, resulting in an 18% reduction in compute costs and improved resource utilization.
- Built a prototype using Python & Docker in client's codebase, resolved 50+ Jira issues with 98% on-time delivery rate.

Education

George Mason University
Master of Science in Data Analytics

Virginia, USA

Jawaharlal Nehru Technological University
Bachelor of Technology in Electronics and Communication Engineering

Andhra Pradesh, India