

Sravani Vakiti

Data Engineer & Analyst | Scalable ETL Pipelines, ML-Ready Data, and Automated BI Solutions | Python • Azure • Snowflake • Power BI • SQL

MA, USA | +1(508) 373-4157 | vakitisravani3@gmail.com | [LinkedIn](#)

Professional Summary

Detail-oriented and impact-driven Data Engineer & Analyst with around **4 years** of combined experience across industry and academia, delivering **scalable, end-to-end data solutions** across **finance, energy, and public sector** domains. Proven success in **building robust ETL pipelines**, orchestrating **real-time and batch workflows**, and deploying **ML models** for advanced analytics. Highly proficient in **Python (Pandas, PySpark)**, **SQL**, **Snowflake**, **Azure Data Factory**, **Power BI**, and **Tableau**, with strong command over **cloud platforms (AWS, Azure)** and **DevOps tools** like **Docker**, **Kubernetes**, and **GitHub Actions**. Known for driving measurable results—**increased pipeline efficiency by 40%**, **reduced churn by 18%**, and **enhanced forecast accuracy**. Adept at **Agile collaboration**, cross-functional teamwork, and translating complex data into **actionable business insights**.

TECHNICAL SKILLS

- **Programming Languages:** Python (Pandas, NumPy, Matplotlib, Seaborn, PySpark), R, SQL, SAS, Ruby, Java, HTML, CSS, JavaScript, MATLAB
- **Databases:** Redshift, DynamoDB, Snowflake, PostGreSQL, MySQL, SQL Server, MongoDB, InfluxDB, MySQL, PostgreSQL, Cassandra
- **Tools & Libraries:** Django, Flask, Splunk, MS Excel, MS Visio, PowerPoint, Word, Microsoft Power BI, Tableau, Google Sheets, UAT, Alteryx, JIRA, Confluence
- **Automation & ETL:** Power Automate, SQL Query Optimization, Data Cleaning
- **Machine Learning & NLP:** TensorFlow, Scikit-learn, PyTorch, NLTK, Pandas, NumPy, Keras, Matplotlib, Seaborn
- **Cloud and DevOps:** AWS (S3, EC2, RDS, EBS, IAM, Lambda), Azure Data Factory, Azure Synapse, Azure DevOps, Git, REST APIs, Docker, Kubernetes, Terraform, Git, GitHub Actions
- **Data Engineering & ETL:** Azure Data Lake, Snowflake, Hadoop, Hive, Spark (PySpark), Databricks Apache Airflow, Kafka, Spark, informatica
- **Operating Systems:** Windows, Linux, MacOS
- **Additional Skills:** Agile, Remote Desktop, Troubleshooting Issues, Process Design, Management Consulting, Information Technology and Analytical Thinking.

EDUCATION

Master of Science in Computer Science (Clark University)

Aug 2023 – May 2025

Bachelor of Engineering in Electrical and Electronics Engineering (CVR College of Engineering)

Aug 2018 – May 2022

EXPERIENCE

DataSage: Data Engineer

Sept 2024 – Present

- Designed and deployed **scalable, production-grade ETL pipelines** using **Python**, **Azure Data Factory**, and **Snowflake**, automating ingestion from **REST APIs** and **cloud databases** to support enterprise analytics.
- Engineered **high-performance PySpark workflows** in **Databricks**, improving large-scale data transformation efficiency by **30 %**.
- Orchestrated **real-time and batch data pipelines** using **Apache Airflow**, **Kafka**, and **Azure Data Lake**, enabling robust and reliable data movement across systems.
- Developed **interactive BI dashboards** in **Power BI** and **Tableau**, empowering business stakeholders with **real-time insights** and **data-driven decision-making**.
- Optimized **complex SQL queries** and integrated **Power Automate** to streamline repetitive tasks, cutting **processing time by 20%**.
- Partnered with ML engineers to build **ML-ready datasets** and support **model training and deployment** using **Scikit-learn**, **PyTorch**, and **TensorFlow**.
- Led **containerization and CI/CD** of data workflows via **Docker**, **Kubernetes**, and **GitHub Actions**, managing infrastructure using **Terraform** on **AWS** and **Azure**.
- Documented data architecture and business logic using **Confluence**, **Visio**, and **JIRA** within an **Agile delivery** environment.

Capgemini: Data Engineer

Aug 2021 – June 2023

- Engineered **high-performance ETL pipelines** using **Azure Data Factory**, **PySpark**, and **SQL**, streamlining ingestion and transformation of **structured and unstructured data** across large-scale systems.
- Integrated diverse data sources—including **Azure Data Lake**, **REST APIs**, **PostgreSQL**, and **MongoDB**—into unified pipelines, enhancing **data availability** and **cross-platform analytics**.
- Automated complex **data cleaning and preprocessing tasks** using **Python (Pandas, NumPy)** and **Power Automate**, reducing manual effort and increasing pipeline throughput by **30%**.
- Tuned and optimized **SQL queries** and implemented robust **data validation frameworks**, resulting in improved **reporting accuracy** and

dashboard responsiveness.

- Delivered **analytics-ready datasets** in **Snowflake**, **Databricks**, and **Hive**, supporting **real-time streaming** and **batch workloads** across cloud ecosystems.
- Designed and deployed **dynamic dashboards** in **Power BI** and **Tableau**, empowering stakeholders with **real-time KPIs**, trend insights, and strategic decision-making capabilities.
- Collaborated within **Agile teams** using **JIRA**, while maintaining detailed documentation of **data flows**, architecture, and logic in **Confluence** to ensure transparency and alignment.
- Deployed modular, cloud-native components via **Docker** on **AWS EC2** and **AWS Lambda**, ensuring scalable and **cost-effective pipeline execution**.
- Led **UAT support** and resolved production bottlenecks, ensuring seamless alignment of data solutions with **business logic**, **compliance**, and **security standards**.

Capgemini: Junior Data Engineer

Oct 2020 – Aug 2021

- Supported the development of **scalable ETL pipelines** using **Python (Pandas, NumPy)**, **SQL**, and **Azure Data Factory** to process large and complex datasets across business domains.
- Contributed to **data ingestion and transformation** from **Azure Data Lake**, **AWS S3**, and **Snowflake**, ensuring high **data availability** and **scalability** for enterprise analytics.
- Performed robust **data cleaning and transformation** using **PySpark** and **SQL**, increasing **data quality** and enabling more reliable insights for analytics teams.
- Created **summary reports** and **data visualizations** using **Power BI**, **Tableau**, and **Excel**, supporting business users and senior engineers in operational and strategic decisions.
- Automated routine data handling tasks with **Power Automate**, improving **workflow efficiency** and reducing manual intervention.
- Actively participated in **Agile ceremonies** via **Azure DevOps** and **JIRA**, contributing to **sprint planning**, **testing**, and **bug resolution**.
- Maintained clear and accessible **technical documentation** in **Confluence** and **Visio**, detailing pipeline architecture, logic, and data mapping for streamlined handovers.
- Provided **UAT support**, identifying and resolving data mismatches, and ensuring smooth transitions from development to **production-ready environments**.

Projects

Boston Crime Trends Analysis (2019–2023) Tools:

Tools: Python, SQL, Tableau, Excel

- Processed and analyzed **300,000+ crime records** using **Python (Pandas)** and **SQL**, uncovering key **spatial and temporal trends** in crime data.
- Performed **data cleaning and preprocessing**, enhancing model and dashboard accuracy by improving **data reliability**.
- Developed **interactive dashboards** in **Tableau** to visualize crime **hotspots**, **peak times**, and trend patterns, aiding **strategic public safety planning**.
- Discovered critical insight into **increased summer-night shootings** in high-risk neighborhoods, directly informing **targeted law enforcement interventions**.
- **Impact:**
Helped reduce crime rates by **12%** in key areas and improved decision-making through **real-time, data-driven insights**.

Automated Cloud Cost Optimization Dashboard

Tools: AWS CloudWatch, Azure Data Factory, Snowflake, Python, SQL, Power BI, AWS Lambda, GitHub Actions

- Engineered an automated ETL pipeline to aggregate daily multi-cloud cost and usage data from **AWS CloudWatch** and **Azure Cost Management APIs**, enabling unified billing visibility.
- Performed data transformation using **Python** and **SQL**, and loaded clean, analytics-ready data into **Snowflake** for scalable querying.
- Designed dynamic **Power BI dashboards** to track high-cost services, unused assets, and monthly usage patterns, improving cost transparency.
- Implemented **GitHub Actions** for workflow scheduling and integrated **AWS Lambda** to trigger real-time alerts for threshold breaches.
- **Impact:**
Reduced cloud cost reporting lag by **40%**, cut manual data processing by **30%**, and enabled teams to act on cost anomalies in real time—driving measurable budget control.

Certifications

- **Microsoft Azure Fundamentals**
Certification Number- 2Y3A9C-572B6D

March 2023