

SRI SAI LALITHA MALLIKA YETURI

551-229-8653 | Jersey city, New Jersey 07306 | yeturi.mallika@gmail.com | [Linkedin](#) | [Github](#) | [Eportfolio](#) |

EDUCATION

Yeshiva University, Katz School of Science and Health

New York, NY

Master of Science in Artificial Intelligence, GPA: 3.7

May 2025

Relevant Coursework: Data Analytics, Machine Learning, Financial Modeling, Statistical Analysis

Villa Marie College for Women

Hyderabad, India

Bachelor of Engineering in Computer Science, GPA: 8.95/10

May 2022

WORK EXPERIENCE

Full Stack Developer Intern

February 2025 - Present

TERN, New York, USA

- Deployed a fine-tuned “InsurGPT-Lite” (7 B params) with vLLM on four A10G GPUs; sustained 8 k req/s at 72 ms p95 for policy-document Q&A
- Applied INT8 GPTQ quantization & head pruning, shrinking VRAM 20 GB→9.5 GB and saving \$2.1 k / month.
- Instrumented CloudWatch + Prometheus; built Nginx→Triton fail-over routing to maintain 99.97 % uptime.
- Authored internal playbook “Latency Wars” (Nsight Sys, CUDA Graphs) adopted across engineering for GPU optimization.

Graduate Teaching Assistant

August 2023 – May 2025

Katz school of Science and Health, New York, USA

- Delivered weekly lectures and lab sessions for Advanced Statistics (35 students); live-coded analyses in R (tidyverse, dplyr, ggplot2) and graded all assignments.
- Designed automated grading pipeline that cut manual review time 40 %.
- Introduced a “Prompt Evaluation Lab,” teaching students to design ablation tests for legal-research LLM prompts.

TECHNICAL SKILLS

- **LLM Inference & Serving:** vLLM, TensorRT-LLM, Triton, ONNX Runtime, DeepSpeed MII
- **Streaming:** INT8 GPTQ, head pruning, KV-cache streaming
- **GPU & HPC:** Multi-GPU orchestration, NCCL, Nsight Sys, CUDA Graphs
- **Frameworks:** PyTorch, TensorFlow, Hugging Face, LangChain
- **DevOps:** Docker, GitHub Actions/Jenkins, CloudWatch, Prometheus
- **Programming:** Python, Java, C/C++, R

PROJECTS

Deep Learning-Based Assessment of Canine Cardiomegaly via Vertebral Heart Scale Prediction

Jan 2025 - Mar 2025

- Developed an EfficientNet-B7 model achieving 86% accuracy in detecting heart enlargement via Vertebral Heart Scale.
- Automated data preprocessing and enforced reproducibility through Git for version control. • Optimized inference speed by ~20% through model pruning and parallel data loading.

NeuralNexus — RAG-Powered NLP Textbook Tutor

Feb 2025 - April 2025

- Converted Jurafsky & Martin’s 1,000-page textbook into an interactive tutor via RAG pipeline (HTML/PDF → Chroma/FAISS 1.5 k-dim embeddings → GPT-4o).
- Tuned retrieval (k = 6) & prompting (T = 0.3) to hit 92 % citation accuracy and 87 % retrieval precision—2.5 × baseline.
- Implemented conversation memory & transparency UI in Gradio, earning 4.6 / 5 clarity and 78 % user preference in Grad-level evaluation. Custom chunking cut retrieval latency 34 %; open-sourced repo, merged 12 external PRs, and built an active contributor base.

AWARDS AND CERTIFICATIONS

- Microsoft certifications (Introduction to AI, AI on Azure)
- IBM Watson (Getting Started with AI)
- Speaker, “RAG for Legal QA,” NYC LLM Meetup (2025)