

Venkata Srihitha Guntuboina

480-847-8344 · gvs190101@gmail.com · linkedin.com/in/srihitha-guntuboina · github.com/gsrihitha

SUMMARY

M.S. in Information Technology (4.0 GPA) with strong experience in designing scalable data pipelines, building cloud-native microservices, and integrating generative AI solutions. Skilled in Python, FastAPI, Langchain, and modern DevOps practices (Docker, Kubernetes, AWS). Proven ability to develop robust backend APIs, streamline ETL workflows, and deliver AI-driven automation for enhanced business operations.

EDUCATION

Arizona State University, Tempe, AZ

May 2025

M.S. in Information Technology

GPA: 4.0

Relevant Coursework: Cloud Computing and Security, Big Data Analytics and Visualization, System Analysis and Design, NLP, Agile Methodologies

Gandhi Institute Of Management and Technology, Andhra Pradesh, India

April 2022

B.S. Computer Science Relevant Coursework: Data Structures and Algorithms, Computer Networks, Operating Systems

TECHNICAL SKILLS

Languages: Python, JavaScript, Objective C, C++, Bash, HTML, CSS3

Back-End & Tools: Flask, Docker, Git, Postman, VS Code, RESTful API Design

Databases & Big Data: MySQL, MongoDB, Hive, Hadoop, Spark, PySpark, Apache Kafka, DBT

Cloud & DevOps: AWS (EC2, S3, Lambda, RDS), CI/CD, GitHub Actions, Kubernetes

Analytics & ML: Pandas, NumPy, Scikit-learn, Tableau, Matplotlib, Seaborn

PROFESSIONAL EXPERIENCE

Software Engineer Intern (AI & Dev Ops), ZofAI, USA (Remote)

July 2025 - Present

- Built and deployed automated agents using Langchain and FastAPI to integrate generative AI capabilities into business workflows
- Used Python and modern APIs to create AI-driven solutions for task automation, focusing on leveraging generative models for enhanced user interaction

Data Engineering Associate, Accenture, Bengaluru, India

October 2022 – July 2023

- Gained practical experience with Hadoop ecosystem tools; implemented distributed data workflows using PySpark
- Optimized 3 ETL pipelines via data partitioning and caching, reducing job runtimes by 20%
- Resolved ETL failures, schema mismatches, and latency issues under tight deadlines through client-facing technical support and cross-team collaboration; reduced reporting errors by 15%
- Recognized by senior stakeholders for adaptability and proactive engagement in resolving platform-level data quality concerns in a fast-paced production environment

PROJECTS

Real-time Product Recommendation System

June 2025

- Architected a scalable real-time data pipeline to capture and process user interaction streams using Kafka and Spark Structured Streaming for personalized product recommendations
- Designed and deployed a modular microservices architecture on AWS EC2 with Docker Compose, delivering low-latency recommendations through a FastAPI-powered REST API
- Implemented a robust CI/CD workflow using GitHub Actions to automate testing, containerization, and cloud deployment; integrated S3 for efficient storage and retrieval of processed data

JobPulse - Job & Internship Portal

April 2025

- Engineered a scalable React frontend with reusable components, cutting load times by 60%
- Built and integrated Flask RESTful backend with MongoDB Atlas, supporting real-time updates
- Leveraged GCP Cloud Functions and Storage to build an automated web scraping pipeline, enabling continuous job data ingestion and recruiter activity tracking
- Automated hiring status tracking to reduce manual effort.

Capstone - Cloud Architecting

March 2025

- Deployed a secure, scalable PHP web application using EC2, RDS MySQL, and ALB across multiple AZs
- Configured Secrets Manager for secure credential management, enabled SSH access via private subnets and set up Autoscaling to maintain performance