

Sai Surendra Raju Dantuluri

surendra@myjobitmail.com | +1 (314) 575 7534 | St Louis, MO | [LinkedIn](#)

SUMMARY

Data Scientist with 4+ years of experience delivering AI/ML and data engineering solutions across healthcare, cloud, and enterprise platforms. Proven expertise in building scalable ML pipelines, GenAI chatbots, and predictive analytics using AWS, Azure, Python, and LLMs like GPT and BERT. Adept at deploying production-ready models, automating workflows with Docker/Kubernetes, and integrating data pipelines using Spark, SQL, and Redshift. Skilled in healthcare analytics, RAG-based systems, NLP, and regulatory-compliant data architectures. Holds an M.S. in MIS from Saint Louis University.

SKILLS

Methodologies: SDLC, Agile, Waterfall

Languages: Python, R, SQL, MATLAB, HTML, CSS

Frameworks & Libraries: TensorFlow, PyTorch, Keras, Scikit-Learn, NumPy, Pandas, Matplotlib, Flask, Transformers Architecture

Big Data & ETL Tools: PySpark, Apache Spark, Hadoop, MapReduce, HDFS, Spark Streaming

Databases: MySQL, PostgreSQL, SQL Server 2008, MongoDB

Data Visualization: Tableau, Power BI, Excel (Pivot Tables, VLOOKUP)

Machine Learning & Deep Learning: Linear/Logistic Regression, Clustering, SVM, PCA, Random Forest, Boosting, Lasso, Ridge, CNN, RNN, Finetuning & Transfer Learning

Generative AI & LLMs: Foundation Models (GPT, LLaMA, Claude), Prompt Engineering, Tokenization (BPE, WordPiece), Reinforcement Learning from Human Feedback (RLHF)

Data Management: S3, Glue

Tools & APIs: Postman, REST APIs, JSON Parsing

IDEs: Jupyter Notebook, PyCharm, Visual Studio Code

Operating Systems: Windows

Soft Skills: Problem-Solving, Communication, Collaboration

EXPERIENCE

Hugging Face, USA | Data Scientist

Jan 2025 – Present

- Integrated Electronic Health Records (EHR) with Retrieval-Augmented Generation (RAG) models and deployed GenAI chatbots using Amazon Bedrock and Transformers, improving digital engagement and self-service adoption by 20% in healthcare navigation tools.
- Automated end-to-end ML deployment pipelines using Docker, AWS CodePipeline, and GCP Kubernetes, ensuring HIPAA-compliant, scalable workflows compatible with CMS and NCQA regulatory standards.
- Developed secure data pipelines using AWS Glue, Redshift, and S3 for real-time claims analytics and utilization tracking, and built interactive dashboards in Tableau and SSRS, reducing analytics latency by 30–40%.
- Built ML-driven care intelligence applications using SageMaker, Bedrock, and LangChain for use cases such as hospital readmission prediction (KED), personalized care planning, and NLP-based medical note summarization, boosting clinical team efficiency.
- Designed optimized relational databases in PostgreSQL and Oracle for claims processing, member eligibility, and care coordination systems, ensuring HIPAA-compliant design and enhancing query performance across healthcare platforms.
- Developed and deployed machine learning models (Logistic Regression, Gradient Boosting, CNNs) under Agile frameworks for fraud detection, risk stratification, and population health analytics, reducing healthcare data processing time by 25%.

Tata Consultancy Services (TCS), India | Azure Data Engineer

Jul 2021 – Aug 2023

- Designed and deployed scalable data pipelines using Azure Data Factory, Databricks, and Synapse Analytics to migrate 12+ TB from legacy systems, reducing infra costs by 25% and improving data availability.
- Engineered robust ETL workflows with HDInsight, Spark SQL, T-SQL, U-SQL, and Azure SQL, enabling standardized data transformation and improving processing speed by 40%.
- Automated ingestion from RDBMS to Azure Data Lake using Spark, Sqoop, Oozie, and Airflow, reducing manual overhead and delivering near-real-time pipelines with 20% lower latency.
- Built 50+ ADF pipelines integrating Blob Storage, SQL Data Warehouse, and APIs (SQL/Mongo), maintaining 99.9% SLA and ensuring unified data ingestion with 99% integrity.
- Created parameterized JSON templates in Azure Data Factory with integrated SQL Activity, reducing pipeline development time by 25% and enhancing cross-service orchestration.
- Developed Power BI dashboards and analytics on 10M+ daily records, enabling actionable insights and accelerating business decisions through optimized visualization and reporting.

Rlogical Techsoft, India | AI/ML Engineer

Jan 2020 – Jun 2021

- Designed and implemented deep learning-based diagnostic solutions (CNN, RNN, LSTM, Transformers) for enterprise clients to detect anomalies and monitor system health, achieving 90%+ model accuracy and reducing MTTR across embedded and IoT platforms.
- Developed scalable ML pipelines for predictive analytics using AWS SageMaker, Lambda, S3, and Redshift, enabling real-time telemetry ingestion and forecasting (ARIMA, Prophet) for automotive and manufacturing clients.
- Delivered end-to-end AI solutions including real-time inference engines and auto-scaling workflows using AWS SageMaker, EC2, and Lambda, improving operational decision-making and analytics throughput for client platforms.
- Built production-ready NLP pipelines (BERT, spaCy) for use cases like sentiment analysis, feedback classification, and automated log tagging, increasing triage accuracy by 25% and enhancing customer experience for service-based clients.
- Engineered deep learning models in TensorFlow and PyTorch for visual quality inspection and multivariate anomaly detection across sectors like automotive, healthcare, and e-commerce, supporting smarter diagnostics and automation.
- Automated MLOps pipelines using Docker, Kubernetes, and Jenkins for cross-cloud deployments (AWS/GCP), reducing client deployment cycles by 40% and ensuring secure, scalable, and fault-tolerant AI service delivery.

EDUCATION

Master of Science in Management Information Systems | Saint Louis University, Missouri

May 2025