

PROFESSIONAL SUMMARY

Results-driven **Cloud & Data Engineer** with 4+ years of experience in **data engineering, analytics, and cloud-native development**. Expert in **SQL, Python, PySpark, and Databricks**, with strong hands-on experience in **ETL design, data modeling, and dashboard development**. Skilled in **AWS architecture, CI/CD automation, and API integrations**. Recently contributed to **Generative AI (GenAI) initiatives at Wells Fargo**, developing **LLM-based retrieval pipelines (RAG)** to enhance data-driven insights. Passionate about turning complex data into actionable intelligence and delivering reliable, high-performance cloud solutions.

EDUCATION

Master’s in data science, GPA: 3.62/4.0 | University at Albany, SUNY Albany

SKILLS

- Programming & Data Processing:** Python, SQL, PySpark, Java, Bash
- Big Data & Analytics:** Databricks, Apache Spark, Hadoop, Hive, Kafka
- Cloud & DevOps:** AWS (S3, Glue, Lambda, Redshift, Step Functions), Terraform, Docker, Jenkins, Git
- Data Management:** Data Modeling, Data Warehousing, ETL Frameworks, Data Quality Automation
- Visualization & BI:** Tableau, Power BI, Spotfire, AWS QuickSight
- APIs & Automation:** FastAPI, Flask, RESTful APIs, JSON/XML Integration
- AI & Machine Learning:** Generative AI (LLMs, RAG, LangChain, Pinecone, OpenSearch), Scikit-learn, TensorFlow
- Collaboration:** Agile/Scrum, JIRA, Confluence, Cross-Functional Stakeholder Management

WORK EXPERIENCE

Wells Fargo | Cloud Data Engineer Dec 2023 – present

- Architected and deployed AWS-based data pipelines integrating Databricks, Glue, and Redshift to support enterprise-scale data processing for compliance and risk analytics.
- Developed Generative AI-driven retrieval pipelines (RAG) using LangChain, Pinecone, and OpenSearch, enabling context-aware document search and summarization for regulatory reporting.
- Fine-tuned LLM models on domain-specific datasets to enhance natural language query performance within the internal data platform.
- Designed and optimized ETL frameworks in Databricks (PySpark, SparkSQL), reducing data latency by 40% across financial datasets.
- Built RESTful APIs using FastAPI and Flask to deliver secure, real-time data services to downstream applications.
- Automated CI/CD pipelines with GitHub Actions, Terraform, and Docker, improving deployment consistency and reducing manual errors by 50%.
- Created executive Tableau dashboards for AI model monitoring and data pipeline health, improving transparency for stakeholders.
- Collaborated with compliance and risk analytics teams to ensure AI systems met governance and audit requirements.

Accenture Pvt Ltd | Data Engineer Sep 2020 – Jul 2023

- Led design and implementation of large-scale ETL pipelines using Databricks and AWS Glue, processing multi-terabyte engineering and operational data for global clients.
- Designed star-schema data models in Redshift and Databricks to support analytical reporting, improving query performance by 45%.
- Created 20+ interactive Tableau dashboards visualizing asset reliability, outage patterns, and performance KPIs; dashboards adopted across engineering and finance teams for operational planning.
- Automated Tableau extract refreshes and publication via Python and Airflow, reducing reporting turnaround time by 60%.
- Developed REST APIs and microservices in Flask and FastAPI for seamless data exchange across distributed systems.
- Implemented CI/CD pipelines with Jenkins, Docker, and Terraform, streamlining deployments across environments.
- Built data validation frameworks and monitoring scripts to maintain 99% data accuracy and lineage traceability.
- Partnered with global teams across 4 business units to translate operational KPIs into analytics-ready datasets used for decision support.

PROJECTS

Customer Churn Forecasting (Python, Scikit-learn, Tableau)

- Applied time-series and classification models to predict customer churn with ~87% accuracy.
- Delivered insights via Tableau dashboards, driving targeted retention campaigns.

Retail Demand Forecasting (PySpark, AWS, ARIMA/Prophet)

- Built forecasting models for demand prediction using transaction and seasonal sales data.
- Enabled supply chain optimization by reducing stock-outs and overstock by 15%.

Data Analysis on E-commerce Sales (Pandas, NumPy, Matplotlib, Seaborn, SQL, AWS Athena)

- Processed and cleaned 500,000+ transaction records from an e-commerce dataset.
- Used AWS Athena to run SQL queries on structured data stored in S3, identified top-selling products, seasonal trends, and customer purchasing behavior using Pandas and SQL queries.
- Built interactive visualizations in Matplotlib and Seaborn to present sales trends and customer segmentation and recommended business strategies based on customer retention analysis and product demand forecasting.

CERTIFICATIONS

Google Data Analytics Certificate from Coursera | AWS Certified Data Engineer – Associate