

Tejaswini. T

tejaswini.dataeng@gmail.com | 9402053292 | [linkedin](#) | United States



[LINK](#)



[LINK](#)



[LINK](#)



[LINK](#)



[LINK](#)



[LINK](#)



[LINK](#)

Professional Summary

- 5+ years of experience as a Data Engineer specializing in building scalable analytics platforms, real-time data pipelines, and enterprise reporting solutions across Azure and AWS ecosystems.
- Deep expertise in Big Data processing using Apache Spark, PySpark, and Hadoop, supporting large-volume analytical workloads.
- Strong proficiency in end-to-end ETL/ELT design, integrating ingestion, transformation, orchestration, and optimization for structured and unstructured data.
- Hands-on experience with Azure services including Cosmos DB, Synapse Analytics, Synapse Link, Data Factory, and Azure SQL for real-time and batch analytics.
- Advanced skills in Microsoft Fabric, including Lakehouses, Direct Lake models, data engineering workloads, and Power BI integration.
- Experienced in building serverless analytics using Azure Synapse Serverless SQL Pools with optimized T-SQL querying.
- Skilled in developing analytics-ready data models, performance tuning, and implementing enterprise BI solutions using Power BI and DAX.
- Proficient with AWS services such as S3, Glue, Lambda, Athena, Kinesis, and CloudWatch for cloud-native data engineering.
- Experienced in IaC automation using Terraform, Azure ARM, and CloudFormation to streamline environment creation and governance.
- Strong programming skills in Python, SQL, PySpark, R, and Java, with additional experience in COBOL, JCL, and HTML.
- Expertise working with relational and NoSQL databases including Snowflake, Azure SQL, PostgreSQL, MySQL, DB2, and DynamoDB.
- Hands-on experience with real-time streaming technologies such as Kafka, Kinesis, and Azure Event Hub for event-driven pipelines.
- Delivered CI/CD automation using Azure DevOps, GitHub Actions, Jenkins, and Terraform modules to support robust data operations.
- Strong background in data modeling, governance, security, and compliance frameworks across enterprise analytics platforms.
- Highly skilled with visualization and reporting tools including Power BI, Tableau, and QuickSight for insight-driven decision-making.
- Knowledgeable in LLMs, Azure AI, and AI-driven automation for analytics enhancement, anomaly detection, and workflow optimization.
- Strong collaborator in Agile/Scrum environments, delivering high-quality, scalable cloud data solutions aligned with business goals.

Technical Skills

- **Programming & Scripting:** Python, R, SQL, Java, PySpark, HTML, JCL, COBOL
- **Data Engineering:** Data Pipeline Architecture, ETL/ELT Development, Synapse link cosmos DB, Real-Time Analytics, Serverless SQL pools, Streaming Data, Metadata Management, Dependency & Workload Management, Data Modeling, Performance optimization
- **Cloud Platforms:** AWS (S3, Glue, Lambda, Athena, CloudFormation, CloudWatch, SNS, SQS, Kinesis), Azure (Cosmos DB, Synapse Analytics, Synapse Link, Data Factory, Azure SQL, Microsoft Fabric, AI Services, Resource Manager, DevOps), Snowflake, Databricks
- **Infrastructure & Automation:** Terraform, CloudFormation, Infrastructure as Code (IaC), Docker, Kubernetes
- **Databases:** SQL Server, MySQL, PostgreSQL, DB2, DynamoDB, Azure Cosmos DB (NoSQL)
- **Analytics & BI Tools:** DBT, Power BI, Tableau, QuickSight, Excel
- **AI/ML & Advanced Tech:** AWS SageMaker, Azure ML, Databricks, NLP, OCR, Document Intelligence, LLMs, Computer Vision, Flask API Deployment, Regression, Decision Trees, Clustering, SVMs, Neural Networks, Time Series Forecasting
- **Streaming & Messaging:** Kafka, Kinesis, Event Hub
- **DevOps & Version Control:** Azure DevOps, GitHub, Jenkins, CI/CD Pipelines
- **Data Governance & Quality:** Data Quality, Lineage, Schema Drift, Security & Compliance (HIPAA), Governance Frameworks
- **Network Analytics:** 4G/5G Performance Metrics, WiFi Analytics, RAN KPIs, Network Monitoring, Capacity Planning
- **Big Data:** Hadoop, Hive, Spark (Advanced), Streaming Telemetry Analytics
- **Visualization:** Tableau (Network Performance Dashboards), Alteryx (workflow automation)
- **Math & Statistics:** Probability, Statistical Modeling, Hypothesis Testing, Optimization Techniques
- **Other:** Agile & Scrum Methodologies, Automation, Real-Time Processing, Cost Optimization, Performance Tuning.

Professional Experience:

BigCommerce - Senior Data Engineer

03/2025 – Present

- Developed and optimized **real-time and batch data pipelines** using **PySpark**, **Python** and **SQL** improving data processing efficiency by **35%** and supporting large-scale BI reporting.
- Integrated **AWS S3**, **Glue**, **Lambda**, and **Athena** to automate data ingestion, transformation, and querying of structured and unstructured datasets, enabling **near real-time insights**.
- Built **scalable data lake architecture** managing **2 TB+** of streaming data with schema drift handling, data quality checks, and lineage tracking for reliable analytics.
- Designed and implemented **ETL/ELT workflows** with **Azure Data Factory** and **Snowflake**, improving pipeline reliability and downstream data availability.
- Automated infrastructure provisioning with **Terraform** and **Azure ARM templates**, reducing environment setup time by **40%**.
- Integrated **Kafka** and **Azure Event Hub** for real-time event streaming, enabling faster operational decision-making.
- Built **CI/CD pipelines** using **Azure DevOps** and **GitHub Actions**, streamlining testing and deployment of data applications.
- Developed **interactive dashboards** with **Power BI**, **Tableau**, and **QuickSight**, empowering data-driven decision-making across business units.

- Implemented **data governance and security controls** with **Azure Key Vault**, **AWS IAM**, and compliance frameworks, ensuring secure handling of sensitive business data.
- Partnered with AI/ML teams to support **predictive analytics** and automation using **SageMaker** and **Azure AI Services**.
- Optimized **SQL** queries and **PySpark** transformations to reduce data latency by 30%, enhancing performance for real-time reporting and analytics.
- Collaborated with cross-functional teams to design scalable data architectures and implement best practices, improving system reliability and maintainability.

Johnson and Johnson - Data Engineer

09/2024 – 03/2025

- Streamed **real-time web logs and operational events** using **Kafka** into **PySpark Streaming on Azure Databricks**, enabling sub-second processing of **1M+ events/day**.
- Implemented **data quality validation** in PySpark, classifying data as valid, bad, or passable to ensure reliability of analytics and reporting.
- Performed **business transformations** and aggregations on streaming data in PySpark, reducing processing time by **40%**.
- Built **ETL pipelines** using **DBT**, **Azure Data Factory**, **Snowflake**, and **Databricks**, integrating structured, semi-structured, and unstructured datasets.
- Developed **Python** and **R** scripts for **historical data preprocessing**, flattening JSON documents, and optimizing file formats (CSV, Avro, TXT, Parquet) for analytics.
- Configured **AWS and Azure infrastructure** with **Terraform**, provisioning storage (S3, Blob), EMR clusters, and networking resources integrated with CI/CD pipelines.
- Optimized queries using **SQL**, **PySpark DataFrames**, and **Spark SQL** to handle large-scale data aggregation and machine learning preparation.
- Designed and deployed **data models** for customer data in **Snowflake**, **Databricks**, and **Azure SQL**, ensuring high data integrity.
- Orchestrated **real-time and batch workflows** in Databricks, integrating Kafka event streams and storing results in cloud storage for downstream analytics.
- Developed **interactive dashboards** with **Power BI**, **Tableau**, and **QuickSight**, delivering actionable insights and supporting predictive and operational analytics.

American Airlines – Data Engineer

02/2024 – 08/2024

- Developed and optimized **Python**, **PySpark**, and **R** scripts for data manipulation, integrating **Pandas**, **NumPy**, **SQL**, **JCL**, and **COBOL** to streamline flight operations, customer data, and maintenance logs, improving reporting efficiency.
- Built and maintained **ETL/ELT pipelines** using **Azure Data Factory**, **Snowflake**, and **Databricks**, ensuring reliable ingestion, transformation, and storage of structured, semi-structured, and unstructured datasets.
- Implemented **Infrastructure as Code (IaC)** using **Terraform** and **Azure ARM templates**, automating cloud resource provisioning, improving environment consistency, and reducing setup time by 40%.
- Leveraged **Azure Databricks** and **Apache Spark** for large-scale streaming and batch data processing, enabling **real-time analytics** and supporting operational decision-making.
- Collaborated with data scientists to deploy **Machine Learning models** using **Azure AI Services**, **SageMaker**, and **Databricks**, enabling predictive analytics for **flight delays**, **customer demand forecasting**, and **maintenance scheduling**.

- Designed **CI/CD pipelines** using **Azure DevOps, GitHub Actions, and GitLab**, automating testing, deployment, and version control of data applications to improve release cycles and reliability.
- Containerized applications with **Docker** and orchestrated deployments using **Azure Kubernetes Service (AKS)**, ensuring **scalability, fault tolerance, and high availability** of data solutions.
- Created interactive dashboards using **Power BI, Tableau, and QuickSight**, providing **real-time operational insights** to airline operations, management, and business units.
- Implemented **data governance, security, and compliance** measures using **Azure Key Vault, AWS IAM, and FAA regulations**, ensuring protection of sensitive operational and customer data.
- Developed **anomaly detection workflows and automated alerts**, optimized SQL queries in **PostgreSQL, MySQL, and Snowflake**, and provided **mentorship and documentation** for junior engineers to ensure high-quality data operations and smooth knowledge transfer.
- Conducted large-scale network performance analytics leveraging streaming telemetry to optimize wireless service availability and end-user connectivity reliability.
- Designed and automated **KPI** dashboards (latency, throughput, QoS failures, signal performance) using Python and Tableau to support capacity planning and incident reduction.
- Applied clustering, regression, and anomaly detection ML techniques to predict network degradations, reducing **mean-time-to-detection (MTTD)** for performance incidents.
- Led data engineering initiatives including architecture definition, **ETL** automation, and cross-team collaboration with **Cloud Security, SRE, and Network Engineering** groups.

DXC Technology – Associate Data Engineer

06/2021 – 07/2023

- Designed and implemented **scalable data engineering solutions** in **AWS** using **Glue and EMR**, processing and analyzing large datasets to reduce data processing time by **25%**.
- Developed and applied **data modeling best practices** for structuring data in **Data Warehouses and Data Lakes**, optimizing storage and analytics performance.
- Built and maintained **ETL pipelines** using **Python, PySpark, and Spark SQL**, converting various file formats to ensure **data quality and consistency**.
- Configured **Databricks Spark environments** on **Google Cloud Platform**, enabling seamless integration with cloud-based data lakes and improving performance for analytics and machine learning workloads.
- Collaborated with data scientists and engineers to integrate **Spark clusters with cloud storage solutions** for **real-time data processing**, enhancing system reliability and throughput.
- Re-designed data infrastructure to **improve scalability and performance** for analytics and reporting, leveraging **AWS, Azure, DBT and cloud-native technologies**.
- Integrated **Power BI dashboards** with multiple data sources, including **SQL databases and cloud storage**, delivering **real-time insights** for business stakeholders.
- Implemented **Data Warehouse solutions** using **Snowflake and Redshift**, enabling scalable data storage and enhanced analytical capabilities.
- Led **performance optimization efforts** on existing ETL pipelines, identifying bottlenecks and implementing improvements to achieve a **40% increase in efficiency**.
- Participated in the **full development lifecycle**, including CI/CD testing, documentation of data pipelines, models, and processes, and stakeholder-driven data modeling sessions to support complex analytical queries.

DXC Technology – Intern Data Engineer-Hyderabad, INDIA

07/2020-05/2021

- Performed data analysis by using tools like **SQL** for data transformation and **Microsoft excel** for visualization
- Performed data ingestion using **SSIS** and created and maintained database tables for future audits
- Created **Tableau** dashboards to visualize employee data across multiple departments
- Authored and maintained **SQL** scripts for different engagements
- Collaborated with stakeholders during project meetings to collect and document data requirements for engineering solutions.
- Worked on various file formats like **Parquet, Json, Avro, CSV**.
- Used **Spark-SQL** to process the data and to run on Spark engine.
- Involved in complete **SDLC** – Requirement Analysis, Development, Testing and Deployments.
- Developed end to end **ETL pipeline** using **Spark-SQL** on **Spark engine**.
- Worked with **Hadoop 2.x** version and **Spark 1.0**.
- Responsible for **Data Ingestion, Data Cleansing** and **Data Transformation**.
- Authored **Python** Scripts to load data in **Microsoft Excel** with specialized formatting.

Certification

- Azure Fundamentals (AZ-900) – Microsoft
- Azure Data Fundamentals (DP-900) - Microsoft
- Fabric Data Engineer (DP-700) - Microsoft
- Cloud Practitioner - AWS
- Snowflake Associate Platform – Snowflake University
- Snowflake Data Warehousing - Snowflake University
- Machine Learning Foundations – AWS
- Generative AI Fundamentals - Databricks