

Vikas Rao Gudipudi

Data Engineer | Dallas, Texas | +1 (945)-274-8862 | vikasraog6@gmail.com | [LinkedIn](#)

Summary

Data Engineer with around 5 years of experience building scalable data platforms across compliance-driven industries. I specialize in turning raw, high-volume data into reliable pipelines and analytics systems that teams can trust. My work typically spans real-time streaming with Kafka and Spark, cloud-native processing on AWS and Databricks, and warehouse design on Snowflake and Redshift. I focus heavily on automation, performance tuning, and data quality, ensuring that every pipeline I build reduces manual effort and accelerates decision-making. Known for working closely with engineering and analytics teams, I translate business needs into robust, production-ready solutions that improve speed, consistency, and visibility across the organization.

Skills

Programming & Scripting: Python, SQL, Shell Scripting, Java, Scala
Cloud & Infrastructure: AWS (S3, Glue, Lambda, Redshift, Athena, EC2), Databricks, Terraform
Streaming & Messaging: Kafka, Event-Driven Processing
Orchestration & Integration: Airflow (DAGs, Sensors), AWS Step Functions, dbt, Talend, REST APIs, Salesforce API
Data Warehousing & Databases: Snowflake, Redshift, PostgreSQL, MySQL, SQL/PLSQL, Dimensional Modeling (Star & Snowflake Schema)
Data Quality, Governance & Modeling: Automated Data Validation, Metadata Management, Schema Evolution, Compliance-Grade ETL, Medallion Architecture
CI/CD & Version Control: GitHub, GitHub Actions, Docker
Visualization & Reporting: Tableau, Power BI

Experience

Truist, USA | Data Engineer Jan 2024 - Present

- Built and maintained cloud data pipelines that moved critical business and compliance data across AWS and Databricks, ensuring the data arrived on time and in a usable, high-quality form.
- Designed reliable ingestion workflows using Spark and Delta Lake that handled daily streaming and batch loads without failures, even when volumes spiked unexpectedly.
- Improved pipeline speed by simplifying transformations, optimizing Spark jobs, and fine-tuning cluster resources so teams received fresh data faster.
- Created clean, well-structured Snowflake models that made it easy for analysts and business teams to explore data and build dashboards without depending on engineering.
- Introduced automated checks and validation steps that caught data issues early, reducing manual fixes and preventing broken dashboards.
- Used Terraform to standardize infrastructure and simplify deployments, helping the team maintain consistent environments across dev, test, and production.
- Worked closely with product, analytics, and engineering teams to understand what data they actually needed and delivered pipelines tailored to their use cases.

Tech Stack: Databricks, PySpark, AWS, Kafka, Snowflake, Terraform, SQL, Python

OneAmerica Financial partners, India | Data Engineer Feb 2022 – July 2023

- Developed end-to-end data pipelines that supported reporting, customer analytics, and operational insights across cloud and on-prem systems.
- Rebuilt legacy workflows in Spark and AWS to make them faster, more stable, and easier to maintain for the wider data team.
- Created streaming processes using Kafka to handle high-volume event data and make it available for near real-time dashboards.
- Migrated data from old warehouse systems to Redshift, improving query speed and reducing the cost of storing and processing large datasets.
- Set up monitoring, alerts, and Airflow schedules so the pipelines ran predictably and engineering teams were notified early if something failed.
- Improved data accuracy by adding validation rules and organizing data into clear layers, which made downstream consumption more reliable.

Tech Stack: Spark, Kafka, Redshift, S3, Athena, Airflow, Python, SQL

Juspay, India | Data Engineer Intern July 2021 – Jan 2022

- Helped automate recurring ETL tasks using Python and Talend to reduce manual work and deliver consistent data for analytics teams.
- Connected multiple data sources—APIs, databases, files—into a single unified pipeline that fed key dashboards and daily reports.
- Used Spark to process and clean large datasets so they were ready for reporting and business review meetings.
- Provided data extracts and visualizations that supported product, marketing, and operations teams in making decisions faster.

Tech Stack: Python, Talend, Spark, SQL, APIs, Tableau

Projects

Real-Time Data Streaming Pipeline

- Built a real-time pipeline that streamed events from Kafka into Databricks, processed them with Spark, and stored curated tables in Delta Lake.
- Implemented checkpoints, schema handling, and incremental processing to keep the pipeline reliable during traffic spikes.
- Delivered low-latency datasets used by analytics and monitoring teams for near-real-time insights.

Cloud Data Warehouse & Analytics Layer

- Designed a Snowflake-based warehouse with staging, transformation, and analytics layers built using modular dbt models.
- Scheduled automated refreshes with Airflow and documented lineage to improve transparency and self-service analytics.
- Built clean data marts that enabled faster reporting and reduced dependency on engineering teams.

Creation of a Presurgical Epilepsy Evaluation Platform

- Created a validation framework to detect schema drift, missing values, duplicates, and unusual patterns before loading data into production tables.
- Logged validation results, triggered alerts, and produced summary reports for engineering teams.
- Improved trust in downstream dashboards by ensuring clean, consistent, and audit-ready data.

Certification

- AWS Certified Data Engineer – Associate
- Python for Data Science

Education

Elmhurst University	Chicago, Illinois
Master of Information Technology	Aug 2023 -May 2025