

Vishal Chhagan Khairnar

vishalk.khairnar@gmail.com | (862) 849-8052 | [linkedin.com/in/vishalkhairnar](https://www.linkedin.com/in/vishalkhairnar) | Newark, NJ-07103

EDUCATION

Master of Science in Computer Science, (GPA: 3.9/4)

Jan 2024 - Dec 2025

New Jersey Institute of Technology, *Newark, NJ*

TECHNICAL SKILLS

Programming Languages : Python, SQL, R, Java, Scala, React, TypeScript, Shell

Databases : MongoDB, Snowflake, Redshift, Oracle, Cassandra and Teradata.

Cloud & Big Data: AWS (Glue, S3, EMR, Lambda, Athena, SageMaker), Databricks, Spark, PySpark, Hive, EMR, Kafka, CloudWatch

ETL & Orchestration: Airflow, Amazon BDT tools, Informatica, IICS, ODI 12C, Kubernetes, Docker, AWS CDK, Terraform, CI/CD

Visualization & Project Management: Tableau, Power BI, Agile, Scrum, Jira, Confluence.

Pharma Vendor's : IQVIA(XPO,FIA,APLD), LAAD. Veeva CRM, MMIT, KOMODO, SHS(Sales/APLD), GA, Specialty Pharma.

EXPERIENCE

Amazon, Bellevue, USA

May 2025 – Aug 2025

Data Engineer Intern

- Built reusable AWS Glue ETL components to scan S3 folders with schema metadata driving field-level extraction from Salesforce objects; leveraged the Salesforce-Glue connector to replace AppFlow, staged data in S3, and loaded into Amazon Redshift, deployed via AWS CDK.
- Architected a secure VPC peering tunnel from Native AWS to Amazon Prod to enable OAuth-authenticated API calls to Big Data Tools, utilizing IAM, VPC endpoints, Route 53, Lambda, and security groups; all infrastructure was provisioned using AWS CDK & CI/CD.
- Developed an Airflow framework to onboard Glue Crawlers, Glue jobs, and Amazon's BDT tools (Cradle, Datanet) into a unified DAG with centralized logging; implemented CloudWatch metric filters to send automated alerts, improving incident response time by ~30%.
- Collaborated with analytics teams to build data models in Redshift for reporting dashboards tracking KPI performance across business units, enabling sales managers to make data-driven decisions in real time.

Axtria, Pune, India

May 2023 – Dec 2023

Lead Data Engineer

- Led a team of 4 members; conducted weekly stakeholder meetings for sprint planning and requirements gathering; collaborated closely with the Scrum Master and led UAT cycles to secure client sign-off, ensuring alignment and on-time project delivery.
- Built and Led efficient ETL pipelines with Snowflake, IICS, Redshift, and Python, processing millions of records weekly, improving performance by 30%, and saving 20+ hours monthly.
- Developed a Market Access stencil framework integrating Claims, Sales, and Lives datasets from common pharma vendors, supported by automated ingestion, dimensional modeling, and a reusable KPI catalog.
- Automated ingestion of 500+ non-standardized one-time shipment files using Python and Snowflake; partnered with stakeholders to standardize schemas, resolve data gaps, define operational workflows, and apply business rules for analytics and dashboard consumption.
- Maintained RAID logs, WSR dashboards, BRDs, TDDs, and warehouse architecture documentation to support Agile delivery, SDLC compliance, and transparent stakeholder communication.
- Led the development of a data pipeline for shipment & sales data using AWS Glue and IICS, integrating with 6 Power BI screens and finalizing 25+ KPIs for enhanced insights.

ZS Associates, Pune India

Aug 2019 - May 2023

Associate Data Engineer

- Designed and implemented a scalable data warehouse for 6+ pharmaceutical clients, managing over 1 billion records weekly and leveraging AWS EMR, EC2, Spark, Redshift, S3, Hive, IICS, Glue and Tableau/Power BI to achieve a 30% reduction in processing times.
- Led a team of 6 members in delivering an end-to-end data warehouse integrating Sales, Marketing, and Veeva CRM data to power FR and HQ dashboards. Developed dedicated screens for product, geography, market, and customer views to monitor key indicators such as territory growth trends, brand penetration, call engagement rates, and account coverage.
- Developed Market Access dashboard, integrating MMIT Lives, IQVIA XPO, and APLD claims to track payer trends, formulary status, market share, and patient access across geographies and different channel, enabling cross-functional teams to identify access barriers and high-value payers.
- Analyzed patient-level data for oncology pharma clients, finalizing 60+ metrics and designing 4 dashboards with 28 screens in Tableau.
- Developed a Data Quality Framework using PySpark, Hive, Glue, and Redshift orchestrated with Airflow to detect anomalies (outlier identification, delta checks), reducing data discrepancies by 25% and improving SLA compliance.
- Implemented Master Data Management (MDM) across 5+ vendors, enhancing data quality and consistency by 20% and reducing discrepancies in consolidated warehouse datasets
- Set up end-to-end operations for a pharma client by orchestrating data warehouse loads using Airflow and Oozie, establishing credential management and critical checkpoint validations to halt pipelines on failure, supporting daily and weekly reporting.
- Served as the primary point of contact for 4 teams, aligning business rules and data requirements. Delivered 50+ data contracts with 100% on-time accuracy, enhancing collaboration

Accenture, Pune India

Jan 2017 - July 2019

Software Engineer

- Developed an automated Python script to transfer data from AWS S3 to Redshift, reducing manual intervention, improving logging for efficient debugging, and sending alerts for timely failure resolution.
- Onboarded cascading JSON pricing files into internal DE framework built on SQS, Lambda, Spark, and Glue, applying multi-stage business transformations and exposing processed outputs to Amazon Redshift and S3.
- Developed an automated PL/SQL and Python process that analyzes weekly production logs, generates Excel reports, updates Oracle log tables, and sends error alerts via email, saving about 20 hours per week by eliminating manual intervention.