Vishnu Varma
PH: 863-904-7060 | vishnuvarma757@gmail.com | www.linkedin.com/in/vishnu-varma-3537a820a

Professional Summary

- 5+ years of experience in designing, developing, and managing complex data pipelines and infrastructure.
- Skilled in ETL/ELT, data modeling, data wrangling, and enrichment using Hadoop, Spark, PySpark, SQL, Scala, Python, Airflow, Azure, AWS, Snowflake, Databricks, Tableau, and Power BI. Proven ability in real-time processing, CI/CD automation, and dashboard/reporting delivery for enterprise-scale solutions.
- Worked on Spark apps in Azure Databricks using Spark SQL for ETL from diverse formats; optimized queries.
- Developed PySpark pipelines in Spark Streaming for large-scale processing from Kafka, S3, and Kinesis.
- Tuned Snowflake queries using profiling, execution plans, and rewrites for faster performance.
- Integrated Hadoop, PySpark, HBase, MongoDB, and Hive for big data analytics.
- Automated ETL using Python with SQL integration and data quality validation.
- Experience in Star/Snowflake schema design, SCDs, fact/dim modeling using Erwin.
- Migrated SQL workloads to Azure SQL DB, Synapse, and Data Lake via ADF; also migrated on-prem SQL to AWS Redshift.
- Worked on CI/CD workflows using Jenkins on Kubernetes; leveraged Git, Docker, and Ansible.
- Delivered interactive dashboards in Tableau, Power BI, and QuickSight using structured/semi-structured data.
- Implemented data governance and security best practices including RBAC, encryption, and compliance with HIPAA/GDPR.

Education Details

- Masters in Data Analytics at Indiana Wesleyan University, Marion, IN

Skills Summary

- **Languages:** Python, SQL, Scala, PySpark
- **Big Data:** Hadoop, Spark, HBase, Kafka, Pig, Cassandra, MongoDB, Snowflake, Airflow
- **Cloud Platforms:** AWS (S3, Redshift, Lambda, Glue, Kinesis, CloudFormation, DynamoDB,Cloudtrail), Azure (ADF, Synapse, Databricks, SQL DB, Data Lake, DevOps), GCP (Big Query, Dataflow)
- **ETL & Streaming:** Glue, Airflow, DataStage, Kafka, Spark Streaming, Azure Stream Analytics, dbt, Fivetran, Informatica, Talend
- **DevOps & CI/CD:** Docker, Kubernetes, Jenkins, Git, GitHub Actions, Terraform, Ansible, IaC practices
- **Visualization:** Tableau, Power BI, QuickSight
- **Data Modeling & Governance:** Dimensional Modeling, Star/Snowflake Schema, Erwin, Data Validation (Great Expectations), RBAC, Encryption, HIPAA, GDPR
- **Methodologies:** Agile, Scrum, CI/CD

Professional Experience

BNY Mellon, US **Oct 2023 - Present**
Data Engineer

- Designing and managing enterprise-scale data lake and real-time analytics systems for global financial data.
- Developed Python-Spark apps to process data from RDBMS and streaming sources (Kafka, Kinesis)
- Configured Snow pipe to load real-time data from S3 into Snowflake with <5 min latency
- Built scalable Glue ETL pipelines to ingest, transform, and load high-volume structured/unstructured data into Redshift
- Automated AWS Glue ETL jobs ingesting 50 M+ records/hour into Redshift, slashing manual intervention by 80%
- Used Spark Streaming APIs for real-time actions; stored in DynamoDB and Snowflake
- Integrated CodeStar + CodeCommit for version control; automated deployment with Jenkins and Ansible
- Enabled micro-batching to ingest millions of files from S3 staging to Snowflake cloud
- Built Python scripts to process CSV, JSON, and Parquet from S3 and store in DynamoDB/Snowflake
- Used CloudWatch/CloudTrail for monitoring ETL jobs and user activity
- Built Alteryx workflows and Tableau dashboards for regulatory and risk reporting
- Tuned Redshift performance with sort/dist keys and WLM queues, boosting query efficiency by 40%
- Ensured audit-ready pipelines by building custom Python-based data quality checks

Environment: AWS (S3, Redshift, EMR, Lambda, Glue, DynamoDB, Kinesis, CloudFormation, SageMaker), Snowflake, Python, PySpark, Kafka, Jenkins, Tableau, Alteryx, SQL, HDFS, Hive, Pig, RDBMS, Teradata

Vishnu Varma
PH: 863-904-7060 | vishnuvarma757@gmail.com | www.linkedin.com/in/vishnu-varma-3537a820a

Kogentix, India **Sept 2021 – Dec 2022**
Data Engineer

- Worked on building scalable, cloud-native data pipelines and analytics solutions using Microsoft Azure services for enterprise clients.
- Built data pipelines using Azure Data Factory, Spark SQL, and Data Lake Analytics with minimal impact on production systems
- Orchestrated Azure Data Factory & Databricks workflows processing 1.5 TB of batch data and 200 M streaming events/day
- Migrated on-prem SQL Server data to Azure Synapse & Azure SQL DB, applying transformations with PySpark
- Used Kafka & Cassandra for distributed data processing and streaming integration
- Created REST APIs in ADF for seamless integration between systems
- Stored structured/semi-structured data efficiently in Parquet/Avro formats to improve query performance
- Automated workflows with ADF scheduling & triggers; integrated Git & CI/CD for DevOps practices
- Delivered actionable insights using Power BI integrated with ADF pipelines
- Developed reusable data products for schema validation, complex transformations, and multi-port outputs (ADLS/SQL)
- Migrated 5 TB of on-prem SQL Server data to Azure Synapse and Azure SQL DB, completing ahead of schedule by 30%
- Imported data into Synapse from ADLS using PolyBase and Spark connector

Environment: Azure Data Factory, Databricks, Synapse, Azure SQL DB, Azure DevOps, Spark SQL, Kafka, Cassandra, PySpark, Python, Git, Power BI, U-SQL, Kubernetes, Jenkins

Aurobindo Pharma, India **July 2019 – Aug 2021**
Data Engineer / Production Support

- Developed and supported data infrastructure using AWS, Spark, SQL Server, and ETL tools for pharma analytics.
- Migrated legacy media data to Wide Orbit using AWS Redshift and custom SQL mappings
- Designed dimensional models (Kimball) with facts, dimensions, and referential constraints
- Built SSIS ETL flows to move data from FTP and flat files to S3 and Redshift
- Converted Informatica ETL to SSIS with dynamic control/script tasks
- Created Tableau dashboards with parameters/actions and SSAS cubes with MDX calculations
- Delivered Power BI reports using Power Pivot & Power View
- Migrated 10 million+ legacy media records to AWS Redshift via SSIS, cutting nightly ETL runtimes by 60%
- Slashed report generation time from 5 h to 30 m with Python scripts
- Cleaned, transformed, and validated over 75 million transaction and EHR records using SQL and Python, cutting downstream data-error rates from 4% to 0.5%

Environment: AWS Redshift, S3, EMR, Kafka, SQL Server, SSIS, SSAS, Tableau, Power BI, T-SQL, Visual Studio